

The partition, not the count: a count-matched measurement of step-size granularity in online meta-gradient optimisation — and the mechanism we could not find

M. Ahmaditeshnizi*

S. Salehkaleybar*

Abstract

MetaOptimize (Sharifnassab et al., 2025) meta-learns one step size per parameter group and reports that finer partitions help inconsistently. That inconsistency has never been separated from the group *count*, so the partition’s contribution remains unmeasured. We hold the count fixed and swap an architecture-aligned partition (one step size per output channel) for uniform chunks, over 2,173 runs on CIFAR-10 and CIFAR-100 with ResNet-18, ResNet-34 and ResNet-50, every contrast taken within one submission, so batch effects cancel.

The uniform partition wins in every count-matched cell we measured. Pooled over the eight SGDM cells the effect is $+0.556 \pm 0.045$ pp and homogeneous there (Cochran Q 4.21 on 7 df), while cells that differ in base optimiser are strongly heterogeneous. Alignment itself is a bounded null: at fixed count and fixed size multiset, permuting which weights share a group is worth -0.009 ± 0.157 pp, 95% CI $[-55\%, +51\%]$ of it. Of nine candidate mechanisms, none survives as a general carrier, and we report all nine.

Scope, stated here rather than deferred. The corpus is CIFAR-resolution vision, and a Lion meta-optimiser carries every count-matched cell but one. Every accuracy is a test-set quantity with no held-out validation split. The effect is ≈ 0.6 pp inside a method trailing tuned cosine by 1.8–4.2 pp. This should be read as a constraint on partition design, not as support for practitioners.

1 Introduction

Adaptive-step-size methods that *learn* their step sizes online — IDBD (Sutton, 1992), Autostep (Mahmood et al., 2012), hypergradient descent (Baydin et al., 2018), SwiftTD (Javed et al., 2024), MetaOptimize (Sharifnassab et al., 2025) — all face a design choice that is usually made silently: how many weights share one step size. Call this the **granularity** of the partition. The choice ranges from one global step size ($m = 1$) to one per weight ($m = P$), and every method in the family exposes it, but almost none of them study it.

MetaOptimize exposes the choice and flags it as unresolved. Its §7.1 CIFAR-10 experiment uses two granularities, scalar ($m = 1$) and a “blockwise” partition of six blocks — “*one for each linear layer and four blocks for the ResNet modules*”; §7.2’s non-stationary experiment uses a “blockwise” of **two**. So the paper runs exactly two granularity values anywhere, and the value of its coarse arm differs between experiments. Its §7.3 ImageNet aside reports “*Unlike CIFAR10, here the blockwise versions of MetaOptimize showed no improvement over the scalar versions*”, and its Limitations section asks for future work on “*the layer and weight levels*”.

That is also how the paper closes on the question: “*while increasing the number of step sizes is anticipated to enhance performance, our experimental findings in Section 7 reveal that this improvement is not consistent across the MetaOptimize approximations evaluated.*” We take that question up with 2,173 runs ($\approx 1,625$ GPU-hours; 1,724 admissible) on CIFAR-10 and CIFAR-100 — ResNet-10, -18, -34 and -50 across the corpus, and ResNet-18, -34 and -50 in every count-matched cell — and report one robust measurement, one bounded null, and a mechanism we could not find.

*LIACS, Leiden University, the Netherlands. Correspondence: mohammadrezaahmaditeshnizi@gmail.com.

That is an unusually clean open question, and the obvious way to answer it is to build the ladder and read off the shape. We did that, and the shape is not the answer — because “the number of step sizes” is not one experimental variable. Along the way we found that:

- moving the shared meta-stepsize η by one decade moves the measured layerwise-minus-scalar gap from $+3.291 \pm 0.146$ pp ($\eta = 10^{-3}$, the paper’s own default) to $+0.655 \pm 0.176$ pp ($\eta = 10^{-4}$), within one batch at fixed seeds — so four fifths of the reported granularity benefit at the default η is tolerance to an over-large meta-step, not accuracy (§4.1);
- the initial step size α_0 sets the *sign* of the ordering at the parent’s own configuration (§4.1);
- the group count, holding the partition *family* fixed, is a smooth function of $\log m$ whose local slope runs from $+0.350$ to $+0.828$ pp per decade inside a single batch, so no single “count slope” exists to import (§4.2);
- and — this is the paper — **at fixed group count the partition still matters**, by an amount larger than the count effect it is usually confused with (§4.3).

The last of these is the measurement we can defend. The rest of the paper is about how far it generalises (§4.4–§4.8), what it is *not* (§5), and what we could not measure (§6–§7).

The measurement. Replacing an architecture-aligned partition (one step size per output channel, *nodewise*) by a uniform one of the *same group count* (fixed-size chunks of K weights) is worth a positive amount of final accuracy in **every within-batch, count-matched cell we measured** — Table 4 gives the cells and the count — spanning three networks (ResNet-18, ResNet-34, ResNet-50), 2 datasets, 4 base optimisers, 2 meta-optimisers, 2 meta-stepsizes and 2 budgets. On ResNet-18/CIFAR-10 with an SGDm base the effect $D = +0.456 \pm 0.195$ to $+0.727 \pm 0.200$ pp across **six independent batches**; on ResNet-34 $D = +0.666 \pm 0.094$ (9 v 9); on ResNet-50 $D = +0.881 \pm 0.261$; on CIFAR-100 $D = +1.640 \pm 0.245$ and $+1.485 \pm 0.238$. Over the fourteen cells that run the same ResNet-18 partition contrast, D varies genuinely across configurations ($Q = 102.47$ on 13 df, $p = 5.5 \times 10^{-16}$, $\tau = 0.295$ pp against 0.143 pp rms measurement error), and splitting those fourteen by base optimiser leaves them homogeneous inside every level (within $Q = 7.36$ on 10 df, $p = 0.69$) while accounting for **92.8% of that Q** (between-base $Q = 95.12$ on 3 df). **We do not claim the base optimiser is identified**: the share is a property of the grouping and not of the label, and the obvious rival — batch identity — accounts for 95.8% of the same Q . §4.4 states exactly what that analysis can and cannot support. At a fixed base optimiser D does not vary: over eight SGDm cells spanning seven separate submissions, two meta-stepsizes, two budgets and three step-size clip boxes, $D = +0.556 \pm 0.045$ pp with $Q = 4.21$ on 7 df ($p = 0.76$, $\tau = 0.000$). We report that as an estimate and not as a threshold clearance: this campaign’s pre-registered minimum interesting effect, 0.30 pp, was fixed before mm1 ran and carried unchanged into eight registered scorers, but it is registered for a single within-batch cell and never for a pool (§3.3).

The bounded null. Architecture *alignment* is not a large carrier, and we can put a number on how large it could still be. Permuting **which** weights share a group while holding the group count **and the exact per-tensor group-size multiset** fixed is worth -0.009 ± 0.157 pp ($t = -0.06$, 3 v 3, one batch), scored NULL against a symmetric band registered in advance. The 95% interval is $[-0.317, +0.298]$ pp = $[-55\%, +51\%]$ of the same batch’s D , and the effect this design could have detected at 80% power is **0.440 pp** = 76% of D . So an alignment effect accounting for most of D is excluded; one accounting for half of it is **not** — power against $A = D/2$ is 0.46. We also record a defect in our own registration: the NULL band’s half-width (0.15) is **narrower than the standard error the batch achieved** (0.157), so a genuinely zero effect would have scored NULL only 66% of the time. What is left as the leading carrier is the **group-size distribution**, which takes $+0.590 \pm 0.107$ ($t = 5.53$) of the same decomposition.

What we could not find. **We do not have a mechanism**, and that is worth saying in the first section rather than the last. We examined nine candidate mechanisms and report the verdict on each in §5 — which are refuted, which is narrowed to a base-meta pairing, which is not separable from the axes it is aliased with, and which cannot be decided by this instrument or this design — because those verdicts are half the contribution. One refutation is fully pre-registered: at the one corner where the base *and* the

meta-optimiser both carry a second-moment normaliser — the corner that predicted the smallest gap in the corpus — $D = +0.889 \pm 0.228$ pp (t 3.89), the largest in the AdamW family, against a bar of $D \geq +0.55$ committed before the runs existed. Chief among the rest: the obvious carrier — degenerate size-1 groups — is not it (removing 100% of a network’s singletons buys $+0.115 \pm 0.133$ pp, t 0.87); the tail story is scoped to a base-meta pairing rather than a base (pooled $D - G = +0.514 \pm 0.056$ under SGDm + Lion, -0.061 ± 0.085 under AdamW + Lion, and $+0.629 \pm 0.242$ under AdamW + RMSProp, where G itself does not move at all: -0.001 ± 0.104); no summary statistic of the size distribution is *identifiable* from this corpus, because at fixed count the design contains exactly one contrast type and every candidate collapses to an indicator for the aligned arm; and **no measurable property of a configuration predicts D out of sample** better than the corpus mean by a margin that survives the power bound ($|r| \geq 0.632$ needed at 10 design points). A cross-validated *null* at $n = 10$ is defensible in a way a cross-validated success at $n = 10$ never is; we report both, and the null is the one that survives.

Scope, stated here and not deferred to a threats section. (i) Everything is CIFAR-resolution vision with ResNets and one meta-learning framework; ImageNet is out of reach on our data allocation (489 of 1000 train classes present, validation set unlabelled). (ii) MetaOptimize trails a tuned SGD + cosine baseline at 95.124 ± 0.047 ($n = 5$, the interior maximum of a bracketed grid), a **1.807 pp deficit**, and the deficit widens to 2.56 pp on ResNet-34 and 4.21 pp on ResNet-50. We make no competitiveness claim. (iii) The base optimiser has been varied four ways and the meta-optimiser once: of the 420 admissible runs in the partition families, 408 carry a **Lion** meta-optimiser and 12 — one cell, the second-moment corner that §6.1’s void batch failed to test — carry RMSProp. One cell is not an axis, so no general statement about the meta-optimiser is available from this corpus. (iv) The effect is ≈ 0.6 pp inside a method that is 1.8 to 4.2 pp behind a cosine schedule, and in our implementation the uniform arm also costs $\approx 17\%$ more wallclock — so this is a **designer-facing** result, addressed to whoever chooses where the groups go in a step-size adapter, and not a recommendation to any practitioner to adopt this method or this partition (§4.7). (v) Every accuracy here is a test-set quantity and no validation split was held out anywhere (§3.3, §7 T12). (vi) We inherit, and partly overlap with, Choi et al. (2019) on tuning-protocol sensitivity, Zheng and Kwok (2019) on blockwise adaptivity, and CAM-HD (Jie et al., 2020) on the granularity ladder and its interior optimum; §2 states exactly what is left.

1.1 Contributions

1. **A count-matched isolation of the group-size distribution from the group count** in meta-learned step sizes, replicated in 20 within-batch cells across three networks, 2 datasets, 4 base optimisers and 2 meta-optimisers (§4.3–§4.5, Table 4, Figure 1).
2. **A pre-registered permutation null, reported with its resolution:** at fixed count *and* fixed per-tensor size multiset, membership is worth -0.009 ± 0.157 pp, 95% CI $[-0.317, +0.298] = [-55\%, +51\%]$ of D , MDE 0.440 pp. This **bounds** alignment’s contribution rather than eliminating it, in one batch at one design point, and we report the registration defect that limits it (§4.6).
3. **A candidate moderator for the effect’s heterogeneity, with its confound measured rather than asserted:** over fourteen same-contrast ResNet-18 cells the base optimiser is the coarsest partition that leaves them internally homogeneous (within Q 7.36 on 10 df, p 0.69) and it accounts for 92.8% of the between-cell Cochran Q (95.12 of 102.47 on 3 df); inside the SGDm level D is homogeneous ($\tau = 0.000$, 95% upper limit 0.109 pp) across seven submissions, two meta-stepsizes, two budgets and three clip boxes. **All four levels now rest on two or more independent submissions**, and each is internally homogeneous (within-level Q 0.17/1, 1.37/1, 1.61/1 and 4.21/7); a registered replication (bm2) bought the two thinnest levels, and it is what makes the decomposition a testable restriction rather than an arithmetic identity. **We report, in the same place, that this is not yet identification:** the share is invariant to any relabelling that induces the same grouping, batch identity accounts for 95.8%, and base survives conditioning on batch only at ΔQ 4.11 on 2 df, p 0.13 (§4.4, Figure 2).

4. **A ranking of the four variables** that “the number of step sizes” conflates, with the within-batch magnitude of each (§4.1–§4.2).
5. **Nine candidate mechanisms — four refuted, one narrowed to a base–meta pairing, one not separable, three undecidable by this design — and two nulls**, reported as a section rather than an appendix (§5), including the identifiability limit that makes a tenth unanswerable from this corpus. One of the four refutations is fully pre-registered: the bar was committed in the batch script before the runs existed and the batch cleared it in the direction the mechanism forbade (§5.5, M9).
6. **A designer-facing prescription with its scope attached, and the scope is not where we first put it**: merging each one-dimensional tensor into a single group is worth $+0.328 \pm 0.084$ to $+1.363 \pm 0.151$ pp across twelve within-batch cells under SGDM, SGD and RMSProp bases, and costs nothing in learned parameters. Under an AdamW base with a **Lion** meta-optimiser it is worth nothing measurable ($+0.007 \pm 0.056$ over two batches); under the same AdamW base with an **RMSProp** meta it is worth $+0.988 \pm 0.231$. The exception is a base–meta pairing, not a base (§4.7, §5.4, §5.5).
7. **A measurement-discipline appendix** documenting a batch that failed silently on its own axis, a metric column that produced two withdrawn headlines, an argument-line defect that voided two batches, and an internal variance claim that did not reproduce (§6, Appendix A).

2 Related work, and what is left

2.1 MetaOptimize (Sharifnassab et al., 2025)

The parent method learns $\alpha = \exp(\beta)$ per group from a discounted meta-gradient carried by an eligibility trace h . Its configuration, extracted from the released text rather than recalled: ResNet-18, CIFAR-10, batch size 100; (base, meta) $\in \{(\text{AdamW}, \text{Adam}), (\text{Lion}, \text{Lion}), (\text{RMSProp}, \text{Adam}), (\text{SGDM}, \text{Adam})\}$; $\eta = 10^{-3}$ for every MetaOptimize row; $\alpha_0 = 10^{-6}$; $\gamma = 1$; blockwise = six blocks. Its §7.5 states that η “*works universally well... with no sweeping required*”.

Three points matter for what follows and are easy to get wrong.

- **§7.1 makes no blockwise-beats-scalar claim in its own text.** Its stated CIFAR-10 finding is “*In every tested combination, MetaOptimize outperforms its corresponding fixed-step-size baseline*” — the method against a tuned fixed LR. The positive granularity reading rests on the §7.3 aside and one figure.
- **§7.1 states no seed count and shows no error bars.** `grep -icE "error bar|shaded|standard deviation|confidence interval"` over the full text returns 0; the only “averaged over 5 random seeds” sentence belongs to §7.2. So the curves supporting the granularity inference are of unstated replication.
- **The §9 sentence is indexed by approximation**, verbatim: “*this improvement is not consistent across the MetaOptimize approximations evaluated.*” It is not a claim that an accuracy-versus- m curve turns down; no such curve was measured. Any paper claiming to *explain a reported non-monotonicity of the granularity curve* is answering a sentence that does not exist. We do not make that claim.

We also record that six blocks on a ResNet-18 (He et al., 2016) have sizes 1856, 147968, 525568, 2099712, 8393728 and 5130 weights: the parent’s coarse arm **cannot contain a size-1 group**, so no degenerate-group story can be about the parent’s own experiment.

2.2 Choi et al. (2019), *On Empirical Comparisons of Optimizers*

Choi et al. (2019) show that optimiser rankings are set by the hyperparameter tuning protocol, and that where an **inclusion relationship** holds — optimiser A’s hyperparameter space contains a setting that

reproduces B — the more general optimiser never underperforms once tuned. This pre-empts a large part of our §4.1: our layerwise-minus-scalar gap moving from +3.291 to +0.655 pp when both arms are tuned is exactly their phenomenon, and we say so.

It does **not** pre-empt §4.3, and the reason is structural rather than empirical. Granularity is not an inclusion relationship in Choi’s sense. The finer partition’s extra degrees of freedom are **state** (the per-group β), not hyperparameters. In the released reduction, the coarse partition’s meta-gradient is a *raw sum* over its members’ meta-gradients with no per-group normalisation (`block_product, patches/HF_patched.py:161--172: scalar` returns `[sum((u*v).sum() for all tensors)]`; `layerwise` returns the per-tensor sums), so a coarse group is not obtainable from a fine partition by tying anything the search space exposes. The single hyperparameter setting at which every granularity coincides is $\eta = 0$, where β never moves and all arms degenerate to a fixed step size. So the inclusion holds only at a point where the method is switched off. That is why a granularity comparison can be non-monotone *at each arm’s own optimum*, and §4.5 shows that it is: tuning each arm to its own argmax moves D by -0.090 ± 0.178 pp.

2.3 Zheng and Kwok (2019), *Blockwise Adaptivity*

Zheng and Kwok (2019) partition network parameters into blocks and make the **second-moment normaliser** blockwise rather than coordinate-wise, arguing theoretically for comparable convergence and better uniform stability, and reporting lower test error than Adam (Kingma and Ba, 2015) on CIFAR-10 ResNet-56/110 for all four of their block constructions. Those constructions are: per parameter tensor/matrix/vector; per output dimension/output channel; per convolution kernel; and per input dimension with a **single step-size per parameter vector**. The last of these is, structurally, the prescription we arrive at independently in §4.7 — a whole 1-D tensor gets one group.

Two things separate their work from ours, and one caveat is ours to declare. They adapt a second-moment *preconditioner*, not a meta-learned step size, so the object being grouped is different. And their four schemes vary group **count** and group-size **distribution** simultaneously, so their ranking of block constructions cannot separate the two — which is the confound the present paper exists to remove. The caveat: the markdown rendering of the paper we hold locally collapses every scheme label to the same token, so **we cannot identify from our copy which of their four schemes they report as best**, and we make no claim about their ranking beyond “all four beat Adam”.

2.4 CAM-HD (Jie et al., 2020), *Adaptive Multi-level Hyper-gradient Descent*

This is the closest prior work and it must be cited in the first paragraph of any granularity claim. Working in the hypergradient-descent framework (Baydin et al., 2018), CAM-HD extends learning-rate adaptation to **layer-wise, unit-wise and parameter-wise** levels, and:

- **states the small-sample mechanism, in words, four years before us:** “*For the model involving a large number of learning rates for different groups of parameters, the updating for each learning rate only depends on the average of a small number of examples. Therefore, when the batch size is also not large, over-parameterization is an issue to be concerned.*”
- **reports the interior optimum explicitly:** “*usually the optimal performance is neither at full global level nor full layer/filter level, but a weighted combination of two levels of adaptive learning rates.*”
- **fixes it with the classically correct remedy:** a learned weighted combination of levels, i.e. hierarchical partial pooling / shrinkage, with the combination weights themselves trained.

What is not in CAM-HD, and is therefore ours: degenerate and size-1 groups, and one-dimensional tensors as a distinguished object; the decomposition of the granularity contrast into a **size-distribution** part and an **alignment** part, with the alignment null; and **count matching** — CAM-HD’s levels differ in group count and in size distribution simultaneously, so its interior optimum is exactly the confound we remove.

CAM-HD also frames the referee question we owe an answer to: *the classical remedy for noisy small groups is shrinkage, and CAM-HD already published partial pooling for this. Why a hard rule rather than the soft, classically optimal version?* We answer it with data in §5.9: three pooling operators were implemented and all three failed their own controls. In particular, full pooling in meta-gradient space (**zpool**, $r = 0$) lands at 87.880 ± 0.070 ($n = 5$), which is 87.879 ± 0.046 ($n = 12$) — plain scalar — to one part in a thousand, confirming the operator does what it says; and at the good partition it is monotone harmful, $r = 1 \rightarrow 91.037 \pm 0.079$ versus $r = 0 \rightarrow 87.880$ (-3.157 pp). The best interior cell we can see on **layerwise** is $r = 0.99$ at 91.112 ($n = 1$), $+0.075$ pp over no pooling and well inside seed noise. Pooling *does* rescue an over-fine partition — **weightwise** goes 79.235 ± 0.154 ($r = 1$) to 89.272 ± 0.162 ($r = 0.99$), a $+10.04$ pp rescue — but the rescued arm is still **2.18 pp below plain six-block** (91.448 ± 0.049) and 2.47 pp below plain **nodewise**. Shrinkage repairs a bad partition; it does not beat a good one.

2.5 Methods that ship a tensor-level rule and do not justify it

Table 1: Methods that impose a tensor-level rule on one-dimensional parameters, and the reason each gives for it. No method in this survey imposes a **minimum group size** on step sizes; two impose a maximum.

method	what it does to 1-D tensors	stated reason
Adam-mini (Zhang et al., 2024)	Algorithm 3, <i>Partition for non-Transformers</i> , is verbatim for name, param in parameters: param_blocks[name] = param — one block per tensor . On a ResNet, Adam-mini is layerwise and already ships our prescription.	Its stated principle is Hessian sub-block alignment . Our permutation null (§4.6) bounds what that principle can be buying in this regime — at fixed count and fixed per-tensor size multiset, membership is worth -0.009 ± 0.157 pp, 95% CI $[-0.317, +0.298]$ — which excludes a large alignment effect within one layer at one design point and does not exclude a moderate one.
Adalayer (Zhao et al., 2024)	LayerNorm parameters form whole-tensor blocks.	Concludes that adaptivity on the last layer and LayerNorm is <i>necessary</i> . Note the direction: this is about merging across norm tensors versus shattering within one, and about a second-moment preconditioner, not a meta-learned step size. It is not evidence for or against our contrast; we flag it because a careless reading makes it look like both.
Muon (Jordan et al., 2024; Liu et al., 2025)	1-D parameters, embeddings and the head are routed to AdamW (Loshchilov and Hutter, 2019).	Tensor rank , from the geometry of Newton–Schulz orthogonalisation.
LARS / LAMB (You et al., 2017, 2020)	BatchNorm (Ioffe and Szegedy, 2015) and bias parameters conventionally exempted from the layerwise trust ratio.	Folklore: tiny norms, unstable ratio.
bitsandbytes (Dettmers et al., 2022)	Tensors with < 4096 elements kept at 32-bit.	Precision, not grouping.
Shampoo (Shi et al., 2023)	max_preconditioner_dim is a cap from <i>above</i> ; 1-D parameters get diagonal Adam.	Cubic preconditioner cost — the opposite direction.
Adafactor (Shazeer and Stern, 2018)	min_dim_size_to_factor falls back to full per-coordinate second moments below a threshold.	Memory — again the opposite direction.

Table 1 surveys the field. No method in it imposes a **minimum group size** on step sizes; two impose

a maximum. Our §4.7 supplies the first count-matched measurement of what such a floor buys, and §5.4 supplies its scope limit.

An attribution we withdraw. An earlier version of this project attributed a $\sqrt{N}$ noise-averaging argument for coarse grouping to Adam-mini (Zhang et al., 2024), Adalayer (Zhao et al., 2024) and SGG (Li et al., 2025). Reading all three in full, that attribution is wrong and is withdrawn: Adam-mini’s non-Transformer partition is per-tensor and its stated principle is Hessian block structure, and Adalayer’s conclusion runs the other way. We record the withdrawal because the misattribution was ours.

3 Method and experimental setup

3.1 The partitions

All partitions are defined over the network’s `named_parameters()` list and are constructed, not assumed; every count in Table 2 is re-derived analytically and cross-checked against the optimiser’s own allocated `beta` on the cluster.

Table 2: The partitions, on ResNet-18 (62 tensors, 11,173,962 weights).

arm	rule	ResNet-18: m	min group	size-1 groups	size CV
scalar	one group	1	11,173,962	0	—
resnet18_blocks	the parent’s six blocks	6	1,856	0	—
layerwise	one per tensor	62	10	0	—
nodewise	one per output channel / row	14,420	1	9,610 (66.64%)	1.916
permnode	nodewise’s exact per-tensor size multiset, membership randomised within each tensor	14,420	1	9,610	1.916
chunk777	uniform chunks of $K = 777$ weights	14,421	10	0	0.045
nodewise1d	one per output channel, but each 1-D tensor is one group	4,851	10	0	0.756
chunk2325	uniform chunks of $K = 2325$	4,851	10	0	0.088
weightwise	one per weight	11,173,962	1	all	0

ResNet-18 has 62 parameter tensors and 11,173,962 weights, of which **41 tensors are one-dimensional** — every BatchNorm scale and shift plus the classifier bias — totalling **9,610 weights = 0.0860% of the network**. **nodewise** gives each of those 9,610 weights its own step size; that is where two thirds of its groups live, over less than a thousandth of its parameters.

3.2 The contrasts

Write $\mathcal{P}$ for a partition of the network’s weights into m step-size groups and $\mu(\mathcal{P})$ for the mean `plateau5` of the arm that trains under it, taken over the seeds of one batch. All four differences below are **within batch** and **count-matched by construction**.

The primary contrast replaces an architecture-aligned partition by a uniform one at the same group count:

$$D = \mu(\text{chunk777}) - \mu(\text{nodewise}), \quad m = 14,421 \text{ vs } 14,420 \quad (1)$$

— one group apart, 3×10^{-5} decades of count. The same contrast with the degenerate size-1 tail **already removed from both arms**, count-matched exactly, is

$$G = \mu(\text{chunk2325}) - \mu(\text{nodewise1d}), \quad m = 4,851 \text{ exactly} \quad (2)$$

so that $D - G$ is the tail’s contribution and, being a difference of two within-batch differences taken in the same batch, carries no cross-batch floor. The alignment leg holds the count *and* the exact per-tensor

group-size multiset and randomises only which weights share a group, within each tensor:

$$A = \mu(\text{permnode}) - \mu(\text{nodewise}), \quad m = 14,420, \quad \text{size multiset identical.} \quad (3)$$

The pure count axis holds the partition *family* fixed and moves only m :

$$U = \mu(\text{chunk2325}) - \mu(\text{chunk777}), \quad \log_{10}(14,421/4,851) = 0.4732 \text{ decades}, \quad (4)$$

and the practitioner’s move — merge each one-dimensional tensor into a single group — is

$$T = \mu(\text{nodewise1d}) - \mu(\text{nodewise}). \quad (5)$$

T is count-confounded by construction, and the confound is exact rather than approximate:

$$T = (D - G) + U. \quad (6)$$

Two further identities are used and are exact by construction, not by fitting. With $B = \mu(\text{chunk777}) - \mu(\text{permnode})$,

$$A + B = D, \quad (7)$$

which is what licenses reporting alignment and size-distribution as *shares* of one effect in §4.6.

Every difference in (1)–(5) is reported with a Welch standard error (Welch, 1947) formed from the two arms’ own variances, never from a pooled or an assumed one:

$$\widehat{\text{se}}(\mu_1 - \mu_2) = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}, \quad t = \frac{\mu_1 - \mu_2}{\widehat{\text{se}}}. \quad (8)$$

Because a percentage point is not comparable across error budgets, the commensurable companion to D is the share of the aligned arm’s remaining error that it removes:

$$\rho = \frac{D}{100 - \mu(\text{nodewise})}. \quad (9)$$

Figure 1(b) plots ρ ; §4.3 states, and this paper obeys, the rule that CIFAR-10 and CIFAR-100 values of D are never averaged and never placed on one percentage-point axis.

On ResNet-34, ResNet-50 and ResNet-18/CIFAR-100 the count matching in (1) is close but not exact: **chunk835** 25,562 vs **nodewise** 25,556 (R34); **chunk295** 79,796 vs 79,700 and **chunk884** 26,715 vs 26,677 (R50); **chunk771** 14,595 vs 14,600 and **chunk2293** 4,943 vs 4,941 (C100). The largest mismatch, 0.14% of the count, is $685\times$ below the resolution floor at the measured count slope.

A structural caveat specific to ResNet-50. Its largest 1-D tensor has 2,048 elements, which exceeds $K = 295$, so **chunk295 splits** some 1-D tensors across groups instead of giving each one a single group. It still contains no size-1 group (min group size 10). ResNet-18, ResNet-34 and ResNet-18/CIFAR-100 all have largest 1-D tensor $512 < K$, so their chunk arms give each 1-D tensor exactly one group. The ResNet-50 D is therefore a slightly different object and is reported as such.

3.3 Metric, admissibility, multiplicity, and units of replication

Metric. For a run r that completed E_r test epochs with accuracies $a_r(1), \dots, a_r(E_r)$, the primary metric throughout is

$$\text{plateau5}(r) = \frac{1}{5} \sum_{e=E_r-4}^{E_r} a_r(e). \quad (10)$$

The 20-epoch analogue $\frac{1}{20} \sum_{e=E_r-19}^{E_r} a_r(e)$ is present in the corpus as the CSV’s **plateau** column and is **not used**: two of this project’s headlines were withdrawn for quoting it (Appendix A.2), and it is banned as a primary.

Admissibility. A run enters an analysis iff

$$\text{adm}(r) = [\text{window_ok}(r) = 1] \wedge [\text{complete}(r) = 1] \wedge [\text{plateau5}(r) \text{ readable}], \quad (11)$$

where $\text{window_ok}(r) = [E_r > 20]$ — the tail in (10) must be a genuine plateau and not most of a short probe — and $\text{complete}(r) = [E_r \geq 0.95 E_r^{\text{req}}]$. The second condition is not redundant: **24 of 2,173 runs pass window_ok while having completed under 95% of their requested epochs** (23 of the 24 finished under 90%; the twenty-fourth stopped at 94 of 100). The worst of them, **rs-blk6-1e4-s2**, ran 29 of 100 epochs and still reports plateau5 85.228; it sits in a **resnet18_blocks** arm of the **rs** meta-stepsizes sweep, where including it moves that arm’s mean by 1.22 pp and inflates its sem 19-fold. **None of the 24 is in a count-matched arm of any cell reported in this paper; seven are truncated snapshots of the in-flight rp1 batch (§3.5), which contributes no number here** — see the attrition ledger in §8, where attrition inside the primary contrasts is zero. Of 2,173 rows, **1,724 are admissible**.

Selection on the test set, stated in full. No validation split was held out anywhere in this project. CIFAR-10 and CIFAR-100 (Krizhevsky, 2009) ship a 50,000/10,000 train/test split; we trained on the 50,000 and evaluated on the 10,000, and plateau5 is a mean of five of those test evaluations. Every tuning decision in this paper was therefore made on the same 10,000 images the paper reports accuracy on. Six places where that matters:

1. **The metric is the test set.** Every accuracy in this paper, including both numbers in §7 T4, is a test accuracy with no held-out estimate behind it.
2. **The operating point.** $\eta = 10^{-4}$ and $\alpha_0 = 10^{-3}$ were chosen because both arms score better there (§4.1). D depends on η at -0.189 pp per decade within batch (§4.5), so D is reported at a test-selected operating point, not at a random one.
3. **§4.5’s per-arm optima** are argmaxes over a two-rung η ladder taken on the same test accuracies that define D .
4. **§7 T4’s absolute numbers are all argmaxes:** the MetaOptimize cell is the maximum of a seven-rung α_0 ladder and the baseline the maximum of a four-rung learning-rate grid, both scored on the test set.
5. **The β -box** was widened between batches in response to observed clipping, which is an analysis-affecting choice made after seeing runs (*Box occupancy* below, §6).
6. **The metric window.** plateau5 was made primary after the 20-epoch column produced two withdrawn headlines (Appendix A.2). The choice has a mechanical justification — a trailing window longer than the plateau contains the mid-training trough of §4.8 at short budgets and not at long ones — and it has been applied uniformly since, but it was not fixed before the first analysis.

Over the campaign, **418 distinct admissible configuration cells** were scored on that test set. Two consequences, and we separate them because they are not the same size. **For absolute accuracies the effect is the usual one and it is not small:** every level in this paper should be read as an optimistic, selection-inflated estimate, and we make no claim that any of them would survive on a genuinely held-out split. **For the contrasts it is much smaller, and for a structural reason:** D , G , A , T and U are differences between two arms trained and evaluated identically inside one batch, and **no arm and no cell was ever selected** — §4.3 shows that every count-matched contrast the corpus contains is reported, including the one that is excluded. What is selected is the operating point at which the contrast is measured, not the contrast. A reader should read D as “the effect at a tuned operating point”, which is the setting a practitioner is in, and should not read any absolute accuracy in this paper as a benchmark result.

Box occupancy. β is clipped to a box, and a run at a guard is uninterpretable. Occupancy is read **per coordinate and per seed**, never from a per-tensor summary — reading the 62-element summary instead reported 0.000000 occupancy on 24 runs that were clipped in every record and inverted an arm ranking, which voided one batch (**ar1**, excluded from every primary here). The scorers report **rec_lo/rec_hi** per run; all **cc1**, **mm1**, **pp1** and **bn1** arms read 0.0000 on both guards. The box is **not** constant across cells: the fourteen-cell pool of §4.4 spans three boxes, and one cell (**hz3**) contains a single seed whose two arms sit in *different* boxes (§7 T9, §4.8).

Unit of replication. Every primary contrast is taken **within one batch** (one contiguous submission), so any offset shared by the two arms cancels identically. §6.3 reports what we can and cannot say about the size of such offsets. Two cells are exceptions to “one contiguous submission” and are marked as such in Table 4: the two **r13** rungs and the two **n11** bases each share a batch, and **hz3** is *two* submissions 32,288 job ids apart.

Heterogeneity. Across k cells with estimates y_i and standard errors se_i , writing $w_i = se_i^{-2}$ and $\bar{y} = \sum w_i y_i / \sum w_i$, we report Cochran’s (Cochran, 1954)

$$Q = \sum_{i=1}^k w_i (y_i - \bar{y})^2 \quad \text{on } k - 1 \text{ df}, \quad \hat{\tau}^2 = \max \left(0, \frac{Q - (k - 1)}{\sum w_i - \frac{\sum w_i^2}{\sum w_i}} \right) \quad (12)$$

(DerSimonian–Laird, DerSimonian and Laird, 1986; I^2 follows Higgins and Thompson, 2002), and, where a moderator is claimed, the exact partition of Q into its within-level and between-level parts, $Q = Q_{\text{within}} + Q_{\text{between}}$ with degrees of freedom adding likewise. Figure 2(b) is that partition.

The precision every pooled quantity is computed at. Every pooled estimate, Cochran Q (Cochran, 1954) and DerSimonian–Laird τ (DerSimonian and Laird, 1986) in this paper is computed from full-precision arm means in the run table, which is what the deposited code computes and what **make reproduce** re-derives. Computing the same quantities from the three-decimal (D, se) pairs *as printed in Table 4* gives values that differ in the second decimal ($43.19 \rightarrow 43.01$ on twelve cells; $36.40 \rightarrow 36.29$ on eleven); both are given in Appendix A.4 so that a referee who recomputes from the printed table alone finds no surprise.

Multiplicity. One family of tests in this paper is large enough that an uncorrected nominal α would be misleading: the secondary contrast G , which is measured once in every count-matched cell. We declare the family — the G leg of every cell in Table 4 that has one, twelve tests as the table stood when the family was fixed — and report Holm–Bonferroni step-down adjusted p -values (Holm, 1979) for all twelve in §5.4, together with the enlarged fourteen-test family that the four new Table 4 cells produce, so that the effect of enlarging it is visible rather than absorbed. The primary D is **not** corrected and is not presented as a family of hypothesis tests: it is one quantity estimated in twenty cells, reported as twenty intervals and pooled once, and the claim made of it is “positive in every cell”, not “significant in k cells”. Where a p -value is quoted anywhere in this paper it is two-sided; at $n = 3$ v 3 we give the Welch–Satterthwaite value as primary and the normal approximation alongside it, because the two differ materially at 2–4 degrees of freedom and the paper’s t columns are the normal-approximation quantity.

The minimum interesting effect, and where it was fixed. Every registered scorer in this campaign judges a within-batch difference against the same three-way band: $|\Delta| \leq 0.15$ pp is a null, $|\Delta| > 0.30$ pp is an effect, and $(0.15, 0.30]$ is registered **in advance** as undecided so that a middling result cannot be argued into either camp afterwards. The band was fixed in `analysis/c76_mm1_score.py` before **mm1** ran and has been transcribed unchanged into seven further scorers — **c77_pp1**, **c78_bn1**, **c79_ar1**, **c81_cc1**, **c84_gn1**, **c88_scorers** (**aw1**) and **c97_bm2** — the last of which records the provenance in its own source (DECIDE = 0.30 # the REPLICATES line, unchanged from **cc1/gn1/aw1**). **0.30 pp is therefore this paper’s**

pre-registered minimum interesting effect, at the level of a single within-batch cell. Two forms of the rule are in use and we say which applies where: most scorers test the point estimate with a *t* or *se* side condition, while *aw1*’s tests the 95% interval.

It is registered for a cell, not for a pool, and we do not extend it. The headline $+0.556 \pm 0.045$ *pp* is an inverse-variance pool over eight cells; no scorer registered a bar for a pooled quantity, and inventing one after the pooling is exactly the failure mode §3.4 exists to prevent. **We therefore report the pooled headline as an estimate with its interval, not as a threshold clearance.** For a reader who wants the comparison anyway, it is stated once and not relied on: the pool clears the 0.30 line on both forms of the rule (point estimate $+0.556$; 95% lower limit $+0.468$), and cell by cell all eight clear it on the point estimate (minimum $+0.428$) while five of eight clear it on the interval — *mm1* ($+0.169$), *hz3* ($+0.258$) and *gn1*’s BatchNorm arm ($+0.287$) do not. That spread is the reason we pool, and it is also the reason the pool’s clearance is not evidence the individual cells were each registered-significant.

Duplicate runs. Two differently-named runs are *duplicates*, not replicates, when they resolve to the same experiment: the same effective argument line after *argparse*’s last-wins resolution, the same ENV line, and the same *--seed*. The rule such runs carry is: **any *n*, *se* or *t* computed over rows sharing a *dup_group* first drops rows marked superseded, then averages within the group, and counts *n* as the number of distinct groups.** The run table’s *dup_group* column carries twenty-one groups — the eighteen collapsed pairs (twelve in *m12*, four in *h2*, two in the CIFAR-100 smoke tests) and the three *a0* reruns already on record, where one member of each carries *superseded* = 1. Twelve of the eighteen are *m12*’s, which is why *m12* is a 3 v 3 cell with *se* 0.195 and not a 6 v 6 cell with *se* 0.142 (§6.1).

Two facts about that column a reader should have before relying on it. **Within a batch it is exhaustive, and we checked rather than assumed:** applying the signature above to every run that logs an ENV: line returns exactly those eighteen pairs and no others. (One hundred and nine runs predate the ENV: line; for those the hierarchical arm axis was never logged, and six apparent collapses in *ha*, *hs*, *det* and *hv* are that blind spot rather than duplicates.) **Across batches the same signature finds far more:** 154 groups covering 364 runs are the same experiment at the same seed in two or more submissions. None of those inflates the *n* of a within-batch primary, which is one of the things “within batch” buys — but they are why §4.4 no longer calls its four repeated cells independent, and one of them is §4.1’s parent cell, where nine *pp*_/PP_ pairs are two submissions of one nine-run design at the same three seeds. **The column does not yet carry those nine.** §4.1 applies the rule to them by hand and reports that cell at *n* = 3; a reader who reads the cell off the run table alone will recover *n* = 6. That is the one bookkeeping gap this paper ships open, we name it rather than let the column imply a completeness it does not have, and *analysis/args_repair.py* is the one-line place it would be closed.

3.4 Registration discipline

Where a batch has a scorer committed before its runs existed, that scorer is run **unedited** and its printed verdict is quoted rather than paraphrased. This rule exists because it was broken: two cycle-91/92 headlines were withdrawn after being produced by reductions hand-written at read time (Appendix A.2). The scorers used here are *analysis/c76_mm1_score.py*, *c77_pp1_score.py*, *c78_bn1_score.py*, *c81_cc1_score.py*, *c82_fa1_score.py*, *c83_gc1_score.py*, *c87_rl3_score.py*, *c87_hz3_score.py*, *c88_scorers.py* (committed at 5129e74, before any *ub9/aw1* run existed; md5 0363bcccb4d3bbad50beb19c9281b9be), and *c84_gn1_score.py* (md5 82c515d1ad490228ecb20f83288a510f), **which halted *gn1* at its commensurability gate T0.6 — a 2.862 *pp* difference in level against a 2.0 *pp* bar — and issued no verdict; we report that outcome rather than the contrast it declined to compute (§5.6, §7 T7). Run instead against the deposit, which excludes the probe files, the same scorer halts one gate earlier and for a different reason (§7 T7, §8).**

The rule has exactly three exceptions in this paper, and they are listed here rather than left to be discovered. In each case we depart from a registered scorer’s own instruction, we say so at the point of use, and we say so again here so that a reader is not required to reconstruct the set.

1. **§4.4 pools across step-size clip boxes, and analysis/c87_rl3_score.py’s registered header forbids it** — “NO POOLING ACROSS BOXES”, on the registered ground that a box change moves the optimiser and not merely the instrument. The base-optimiser decomposition pools fourteen cells spanning three boxes ($-15:-2.3026$, ten cells; $-30:9.0$, three; $-25:-2.3026$, one). Our justification is a measurement rather than an assertion, and we give it with its own history rather than only in the state that flatters it. On the eleven-cell pool of the previous draft, partitioning Q by box left between-box Q 1.08 on 2 df (p 0.58) against within-box Q 35.32 on 8 df; on the thirteen-cell intermediate pool that contained **bm2** but not **sm3** it rose to **5.14 on 2 df (p 0.077)**, because both **bm2** cells sit in $-15:-2.3026$; on the fourteen-cell pool it is **0.69 on 2 df (p 0.71)**. That figure moves because the box axis is partly confounded with the base axis — all four non-SGDm cells sit in one box — so we do not rest the pooling on it. The uncontaminated test is the one taken inside a single base: over the eight SGDm cells, which between them span all three boxes, between-box Q is **0.76 on 2 df, p 0.68** (pools $+0.582 \pm 0.080$, $+0.629 \pm 0.123$, $+0.523 \pm 0.060$). That is the number the pooling rests on. A β -box column is carried in Table 4 and in Appendix B so that a reader who does not accept the argument can redo the split. It remains an exception.
2. **Table 4 row 4 is a quantity we re-derived past the point where its own scorer halts.** analysis/c84_gn1_score.py evaluates its T0 validity gates first and per run; T0.5 is box occupancy measured on **gn1**’s own probe records, which are not in the deposit, so the scorer halts at T0.5 and never reaches T1, the positive-control gate at which D_{BN} would have been computed. Row 4’s $+0.587 \pm 0.153$ is therefore ours, not the scorer’s. The GroupNorm arm of the same batch, which is the contrast the scorer was written to judge, is excluded from every pool and quoted nowhere (§4.3, §7 T7).
3. **§4.8 reports the budget contrast on plateau5 and declines the scorer’s registered primary window.** analysis/c87_hz3_score.py fixes a 50-epoch mean as its primary reading and returns **GROWS** on it, with $D(300) - D(100) = +0.229$; plateau5, this paper’s primary metric everywhere else, returns -0.149 ± 0.105 ($t = 1.42$). We keep plateau5 for consistency and print the registered verdict alongside it rather than instead of it, and §4.8 explains the mechanism of the disagreement — a 50-epoch trailing window at $B = 100$ contains the mid-training trough and at $B = 300$ does not.

There is no fourth. Every other scorer named in this section was run unedited and its printed verdict quoted, including where the verdict cost us a subsection (§5.6, §7 T7) and including where it refuted a mechanism we had proposed (**sm4**, §5.5). Supplying a scorer’s own documented arguments — `--root`, `--runs`, `--probes`, `--csv` — is not editing it.

This list is a list of the scorers we *have*, not a claim of coverage. **Five batches behind material reported in this paper have no scorer that was registered before their runs existed** — **m12**, **n11**, **r50**, **gm2** and **sm3** — and the per-batch provenance table in §8 marks each one. **n11** is the serious case: it supplies two of Table 4’s rows and both of the single-batch levels of the base-optimiser moderator in §4.4, so the newest headline’s weakest flank is also its unregistered one. Their run-table values were reconciled field-by-field against their own runs’ **ARGS**: lines with zero mismatches; what is missing is not the data but the commitment-before-the-fact. The replication that repairs **n11**’s half of it has since run: **bm2** (§3.5, R3) was registered before any of its runs could exist and supplies a second, independent batch at both of **n11**’s levels (§4.4).

Every figure in this paper is generated by analysis/c98_figures.py directly from results/all_runs.csv and the raw .out series; `--numbers` prints each plotted value. **No figure contains a value that was typed.**

analysis/c98_reproduce.py re-derives and asserts the numbers that carry a claim, and we state its scope rather than let “the audit passes” be read as coverage. It asserts: the corpus and admissibility counts (§8); every cell of Table 4 with its standard error and the count of positive cells; the commensurable ratio ρ of (9) for all twenty cells **and the rank of each cell we name** (§4.3); the fourteen-cell pool, τ , the within/between-base Q partition and the conditional test against batch identity (§4.4); the alignment legs A and B and the additivity identity (§4.6); the prescription table T (§4.7); the tail decomposition $D = G + (D - G)$ and its pools (§5.4); the count axis U and the **ck1** ladder (§4.2); the meta-stepsize pair (§4.1); the budget ladder read within run from the raw series (§4.8); the competitiveness deficit (§7

T4); the **gn1** commensurability gate’s levels, budgets and gated difference (§7 T7); Appendix A.4’s printed-table pools; the partition-family meta-optimiser census (§7 T1); and the deposit reachability register of §8. It does **not** assert: prose-only quantities, group counts, the attrition ledger’s upstream cluster-side rows, wallclock and byte counts, Appendix B’s arm means, or any value that exists only inside a registered scorer’s own printed output — those are quoted from the scorer, not re-derived. Measured by `python3 analysis/c98_reproduce.py --census`, the audit executes **283 assertions covering 216 of the 725 distinct quantity-numerals** in this manuscript. We print the fraction rather than a superlative because a superlative is exactly the kind of claim this audit exists to catch: the first version of it checked no ρ , and a false ρ superlative survived two review cycles in §4.3 as a result.

3.5 Four pre-registered batches: two scored, two still in flight

Three weaknesses this paper states about itself, and one experiment it declares was never run, are addressed by four batches submitted while this draft was being written. Two of the four — **bm2** and **sm4** — have since completed and been scored by running their registered scorers **unedited**, and their verdicts are folded into §4.4, §4.7, §5.4, §5.5 and §7. Two — **rp1** and the **hz3** seed-5 trio — are not scored, and **no number from either enters any claim in this paper**. The registrations are stated here in full regardless of outcome, so that the decision rules are on the record ahead of the numbers, and so that a reader can check that the two verdicts we did read are the ones we said we would read.

Every one of the four was submitted under the two rules that the failures in §6.1 paid for: **STANDING RULE 20** — a batch’s science is what the runs’ own **ARGS:** line says, never what the submission script’s header claims — enforced by `analysis/argsline_guard.py` and `bin/_lib_guards.sh` before and after launch; and **STANDING RULE 21** — no batch is submitted without a scorer that already exists, whose selftest already passes, and whose sha256 is recorded.

Table 3: The four pre-registered batches. Two have been scored by running their registered scorers unedited; the other two contribute no number to this paper.

tag	batch	jobs	what it repairs	registered scorer	status
R1	rp1	24	the alignment null’s power and its permutation-seed confound (§4.6)	<code>c97_rp1_score.py</code>	complete on disk, not scored
R2	hz3 seed-5 trio	3	the budget cell’s clip-box and GPU-class mismatch (§4.8, §7 T9)	<code>c87_hz3_score.py</code> , reused unedited	queued
R3	bm2	12	the base-optimiser moderator’s single-batch SGD and RMSProp levels (§4.4)	<code>c97_bm2_score.py</code>	SCORED — both levels REPLICATE
R4	sm4	12	the second-moment corner the void batch of §6.1 failed to test	<code>c97_sm4_score.py</code>	SCORED — mechanism REFUTED

R1 — rp1, the alignment replication and the two-way variance decomposition. `nodewise` $\times$ 6 fresh run seeds plus `permmode{101,202,303}` $\times$ the same 6 seeds, fully crossed: 24 jobs at **pp1**’s own cell (ResNet-18/CIFAR-10, SGDm + Lion (Chen et al., 2023), $\eta = 10^{-4}$, $\alpha_0 = 10^{-3}$, 100 epochs, box $-15: -2.3026$). The permutation seed is an explicit constant **decoupled from the run seed**, which **pp1** could not do (§4.6, limit 2), and the guard asserts on the live tree that each draw reproduces `nodewise`’s per-tensor size multiset exactly, is a true non-identity permutation, differs from the other two draws, and coincides with none of **pp1**’s. Registered in advance: the design takes `se(A)` from 0.157 to ≈ 0.09 and the minimum detectable effect from 0.440 pp to ≈ 0.25 pp — below half of D and inside the pre-registered band for the first time — and yields a permutation-draw $\times$ run-seed variance decomposition. **Also registered in advance, and this is the part that binds us: if *rp1* returns an interval that still spans half of D , the alignment leg is to be reported as *underdetermined*, not as a null.** The scorer additionally prints **pp1**’s registration defect (band half-width $0.15 < \text{realised se } 0.157$) whichever way the new data land.

R2 — the hz3 seed-5 trio, re-run box- and hardware-matched. hz3-ch-s5, hz3-c23-s5 and hz3-n1d-s5 re-run at 300 epochs in the batch’s own box $-30:9.0$ and pinned to the batch’s own GPU class for seed 5 (RTX 2080 Ti), restoring a matched 6 v 6 at every budget. Nothing else in hz3 is touched; its seed-5 nodewise partner is already in the right box. The composed argument line was verified byte-identical to the original hz3-ch-s5 run’s own ARGs: line modulo the partition flag, and the box arithmetic was pre-registered: at $\text{ms} \cdot T = 10^{-4} \times 300 \times 500 = 15.0$ nats of travel from $\beta_0 = -6.907755$, β is confined to $[-21.908, +8.092]$, so **both rails of $-30:9.0$ are provably unreachable** while the superseded box’s floor of -15 is reachable from epoch 162 — which is exactly what the probe records show happened (§7 T9). This batch changes the data the registered reader reads; it does not change the reader. analysis/c87_hz3_score.py is reused **unedited** (sha256 0be1f5201d..., selftest 144/144 PASS) rather than a second scorer being written, because writing one would create two registrations for one question.

R3 — bm2, the base-moderator replication. One submission, four arms (SGD/RMSProp $\times$ nodewise/chunk777), three **fresh** seeds (3, 4, 5 against n11’s 0, 1, 2), 12 jobs at n11’s own cell taken from n11’s own ARGs: line. Independence is bought twice — a separate submission, so the batch random effect is resampled, and fresh seeds, so the m12 failure mode (two “independent” halves that shared their seeds) cannot recur. Registered bands, committed before the runs existed: **REPLICATES** if $D' \geq 0.30$ and $t \geq 2$; **FAILS TO REPLICATE** if $D' \leq 0.15$ and $\text{se} \leq 0.15$; **UNDECIDED** otherwise; agreement with n11 if $|D' - D_{\text{n11}}| \leq 0.5$, the unmodelled cross-batch offset. The registration also states, in advance, that **the powered-null branch is unreachable on the planning noise** — $\text{se}(D') = 0.194$ exceeds the 0.15 the null branch requires — so a small D' must read **UNDECIDED** and may not be written up as a refutation. That sentence exists because this project has already published an underpowered null as a refutation once (§4.6).

R4 — sm4, the second-moment corner, properly composed. AdamW base $\times$ RMSProp meta (Tieleman and Hinton, 2012) at four granularities $\times$ 3 seeds, 12 jobs: the cell that sm3 declared and did not run (§6.1). The payload is built once in a single shell array and handed to sbatch unmodified, so the \$BFLAGS-then-append shape that voided both m12 and sm3 is absent by construction; the first launched job’s own ARGs: line was read off the cluster and carries **exactly one --alg-meta**, reading RMSProp, and the runtime prints neither of the attribute-mismatch warnings that sm3 printed. Registered in advance against anchors re-derived at write time (aw1 $D + 0.279 \pm 0.087$, sm3 $D + 0.141 \pm 0.064$, pooled AdamW + Lion $D + 0.210 \pm 0.056$): **REFUTED** if $D \geq +0.55$; **CONSISTENT** if $D \leq +0.279$ and the 95% upper bound is below $+0.55$; **UNDECIDED** otherwise, in which case the registration forbids adding seeds to this batch and requires a fresh registered replicate instead. A blocking operating-point gate fires first: an RMSProp meta-optimiser has never run anywhere in this corpus and its meta-stepsize was matched to aw1 rather than tuned, so if the 12-run mean plateau5 lands more than 3.0 pp from the 93.161 anchor **every contrast in the batch is demoted to descriptive and labelled off-anchor** — reported, never discarded.

Status, stated because “in flight” is not a uniform state. R3 and R4 are done and are in the paper. bm2’s twelve jobs and sm4’s twelve jobs completed at 100/100 epochs each. Each was scored exactly as registered: the scorer file was run **unedited**, with only its own documented --root / --runs / --probes arguments supplied, and its verdict is quoted rather than paraphrased — bm2’s in §4.4, sm4’s in §5.5. Both scorers’ ARGs gates (RULE 20) returned CLEAN on the runs’ own ARGs: lines, and both box gates were measured on the batches’ own probe records rather than inherited: sm4’s four arms and bm2’s twelve probe directories all report rec_lo and rec_hi of exactly 0.0000 over $T = 10,000$ recorded steps, so neither batch touched a clip rail. Both verdicts are among the six that a reader holding the deposit alone cannot regenerate, because both gates read the excluded probe records; §8 says so and names them.

R1 is complete on disk and deliberately unscored. All 24 rp1 runs reach 100 epochs in their own .out series. Seven of the twenty-four were ingested into results/all_runs.csv mid-flight and their rows still carry the partial epoch counts of that snapshot, so the run table is behind the disk for this batch. We have not run analysis/c97_rp1_score.py, and **no rp1 number appears anywhere in this paper**, including in §4.6 where it would matter most. The registration stands as written, including the clause that binds us: if rp1 returns an interval that still spans half of D , the alignment leg is to be reported as

underdetermined, not as a null.

R2 is queued. The three seed-5 jobs sit PENDING on the same congested partition on which the original hz3 seed-3 and seed-4 runs waited 16 and 23 hours, so a wait is expected rather than anomalous. STANDING RULE 20's post-launch ARGS check is owed on them the moment they start.

4 Results

4.1 The meta-stepsize and the initial step size each move the granularity gap by more than the gap

Meta-stepsize. Within one batch (ms), at fixed seeds, ResNet-18/CIFAR-10, SGDm base + Lion meta, $\alpha_0 = 10^{-3}$:

η	layerwise	scalar	layerwise – scalar
10^{-3} (the parent's default)	91.259 ± 0.076 ($n=3$)	87.967 ± 0.124 ($n=3$)	+3.291 $\pm$ 0.146 , t 22.6
10^{-4} (both arms' bracketed optimum)	92.960 ± 0.024 ($n=3$)	92.305 ± 0.174 ($n=3$)	+0.655 $\pm$ 0.176 , t 3.72

Both arms are **better** at $\eta = 10^{-4}$ than at $\eta = 10^{-3}$ — scalar by 4.338 pp, layerwise by 1.701 pp. The granularity gap at the default η is therefore mostly the scalar arm's greater sensitivity to an over-large meta-step. This is Choi et al.'s tuning-protocol effect, in the specific place where the parent's protocol fixes η and states that no sweeping is needed.

Initial step size. At the parent's own (AdamW base, Adam meta, $\eta = 10^{-3}$) configuration:

α_0	scalar	six blocks	layerwise
10^{-6}	91.630 ($n=3$)	91.681 ($n=6$)	91.781 ($n=3$)
10^{-3}	92.631 ($n=6$)	92.237 ($n=9$)	91.382 ($n=16$)

The ordering is monotone **increasing** in granularity at $\alpha_0 = 10^{-6}$ and monotone **decreasing** at $\alpha_0 = 10^{-3}$. (These cells pool across batches with unequal composition and are descriptive; the in-batch statement is the η pair above.)

The parent's own cell, checked for a clipping artefact — and it is three seeds, not six. Eighteen unaugmented AdamW/Adam runs at $\eta = 10^{-3}$, $\alpha_0 = 10^{-6}$ in the parent's own clip box are the only runs in the corpus that touch the parent's actual experiment. They read scalar 73.830, six blocks 74.352, layerwise 73.449, so the parent's blockwise-beats-scalar step reproduces at **+0.522** and the *next* rung reverses it at **−0.903**. The eighteen are not eighteen independent runs. They are two submissions of one nine-run design — pp_* (job ids 4683697–4683705) and PP_* (4700548–4700556) — and for all nine same-arm pairs the effective argument line, the ENV line and the --seed are identical. Under this paper's own duplicate rule (§3.3) the cell is **3 v 3**, and the two steps become

six blocks – scalar = $+0.522 \pm 0.293$, t **1.78** — **UNRESOLVED** (Welch p 0.17 on 3.05 df; normal-approximation p 0.075), where reading the eighteen as six seeds gave ± 0.208 and t 2.52; and **layerwise – six blocks** = -0.903 ± 0.278 , t **3.25**, which the Welch statistic §3.3 makes primary at 3 v 3 also leaves unresolved (p 0.058 on 2.61 df) though the normal approximation does not (p 0.0012).

We state that plainly: at the parent's own configuration and in the parent's own clip box, this corpus does not resolve the parent's blockwise-beats-scalar step. Per seed the step reads

+0.481, +0.336, +0.750, and neither submission resolves it alone (**pp** $+0.477 \pm 0.294$, **PP** $+0.567 \pm 0.358$). That is a finding about how thin the parent’s own cell is, not a defect in the replication, and it is the reason the reproduction question had to be bought with fresh runs rather than settled on the eighteen.

It also had to be bought for a second reason. $\beta_0 = \ln(10^{-6}) = -13.8155$ sits 1.1845 nats above a -15 floor while the travel budget is 50 nats, so this cell is arithmetically capable of binding on both guards, and **runs/PP** holds no **probe.jsonl** — occupancy is unmeasured in both directions. We re-ran the cell (**ub9**, 9 jobs, three arms $\times$ seeds 0–2, one submission) in a wide box with clip counting on. The registered scorer’s verdict, verbatim: “*REPRODUCTION STANDS. The parent’s blockwise>scalar step survives a box in which clipping is measured rather than assumed, AND the next rung still reverses it.*” That scorer’s rule is a threshold on arm means fixed before the runs existed, so the duplicate correction above cannot move it. In the wide box, and at a genuine 3 v 3 with no duplicate pair in it, the step is $+2.375 \pm 0.557$ (t 4.26) and the reversal at the next rung is -0.806 ± 0.273 (t 2.95).

So the reproduction claim rests on **ub9 and not on the eighteen.** In the parent’s own box the step is +0.522 and unresolved; in a box where clipping is measured it is +2.375 and resolved, and the non-monotonicity survives both. The eighteen are still worth their place — nine same-config, same-seed pairs across two submissions give mean $|\Delta\text{plateau5}|$ 0.306 **pp** and max 0.510 **pp**, one of the two direct re-measurement bounds in this corpus (§6.1 gives the other) — but they are not six seeds and this paper no longer reports them as six.

4.2 The group count is smooth within a family, but shallow and sign-unstable across cells

Holding the partition family fixed (uniform chunks) inside **one** batch (**ck1**), ResNet-18, SGDm + Lion, $\eta = 10^{-4}$:

arm	m	plateau5 ($n=3$)
chunk1 (per weight)	11,173,962	90.979 ± 0.131
chunk2	5,586,981	91.095 ± 0.024
chunk16	698,373	91.411 ± 0.277
chunk128	87,297	92.159 ± 0.063
chunk1024	10,913	92.526 ± 0.077

The five-rung ascent is $+1.547 \pm 0.152$ **pp** (t 10.2) over 3.01 decades. Its *local* secants, in ladder order from finest to coarsest, run $+0.383$, $+0.350$, $+0.828$, $+0.407$ **pp** per decade — a factor of 2.4 between adjacent rungs of one curve, which is why mutually incompatible “count slopes” can all be fitted to one corpus.

Measured *directly*, at fixed partition family and inside a batch, the count axis is small and sign-unstable. Over fourteen in-batch measurements of $U = \text{chunk2325} - \text{chunk777}$ (0.4732 decades):

cell	U	cell	U
cc1	$+0.101 \pm 0.190$	m12	$+0.337 \pm 0.112$
r13 @10⁻⁴	$+0.064 \pm 0.154$	g3m (R34)	$+0.263 \pm 0.074$
r13 @3$\times$10⁻⁴	$+0.017 \pm 0.098$	r50 (R50)	$+0.321 \pm 0.259$
fa1	$+0.019 \pm 0.094$	gm2 (C100)	-0.054 ± 0.196
hz3 (300 ep)	-0.148 ± 0.080	n11/SGD	$+0.182 \pm 0.283$
aw1 (AdamW)	$+0.045 \pm 0.097$	n11/RMSProp	-0.037 ± 0.085
sm3 (AdamW)	$+0.071 \pm 0.082$	sm4 (AdamW, RMSProp meta)	$+0.359 \pm 0.074$

U changes sign across cells — it is negative in three of the fourteen — and its magnitude never exceeds +0.36 **pp**, which is smaller than D in every cell where both are measured. The maximum is **sm4**’s $+0.359 \pm$

0.074, the corpus’s one non-Lion cell, where it is 40% of that cell’s own D ; on the thirteen Lion cells the maximum is `m12`’s $+0.337$ and the previous draft’s $+0.34$ bound still holds. (`m12`’s entry is computed on its three seed groups, not its six runs, per §3.3: on six runs it would read ± 0.088 rather than ± 0.112 .) **No single count slope exists to import**, and any count correction must be measured in the batch and at the budget being corrected.

4.3 The primary: at fixed group count, the partition matters

Table 4: D (Eq. 1) = uniform chunk – architecture-aligned **nodewise**, count-matched, within batch, **plateau5**. Welch difference of arm means (Eq. 8); se from the two arm variances. Every cell is a separate contiguous submission except the two `r13` rungs and the two `n11` bases, which share a batch, and **hz3 (row 9), which is two submissions 32,288 job ids apart, in two clip boxes and on two GPU classes** — see §7 T9. The β -box column matters in §4.4: the pool spans three boxes, and box is shown there not to be a moderator.

#	batch	network	data	base	η	epochs	β -box	n	D (pp)	se	t
1	<code>cc1</code>	ResNet-18	C10	SGDm	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.727	0.200	3.63
2	<code>mm1</code>	ResNet-18	C10	SGDm	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.485	0.161	3.01
3	<code>pp1</code>	ResNet-18	C10	SGDm	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.581	0.141	4.11
4	<code>gn1</code>	ResNet-18	C10	SGDm	10^{-4}	100	$-15:-2.3026$	4 v 4	+0.587	0.153	3.83
5	<code>m12</code>	ResNet-18	C10	SGDm	10^{-4}	100	$-15:-2.3026$	3 v 3 ($\times 2$ reruns)	+0.456	0.195	2.34
6	<code>r13</code>	ResNet-18	C10	SGDm	10^{-4}	100	$-30:9.0$		+0.681	0.173	3.93
7	<code>r13</code>	ResNet-18	C10	SGDm	3×10^{-4}	100	$-30:9.0$	3 v 3	+0.591	0.096	6.18
8	<code>fa1</code>	ResNet-18	C10	SGDm	3×10^{-4}	100	$-25:-2.3026$	6 v 6	+0.629	0.123	5.11
9 [†]	<code>hz3</code>	ResNet-18	C10	SGDm	10^{-4}	300	$-30:9.0^{\dagger}$	6 v 6	+0.428	0.086	4.94
10	<code>aw1</code>	ResNet-18	C10	AdamW	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.279	0.087	3.19
11	<code>n11</code>	ResNet-18	C10	SGD	10^{-4}	100	$-15:-2.3026$	3 v 3	+1.035	0.109	9.54
12	<code>n11</code>	ResNet-18	C10	RMSProp	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.973	0.251	3.87
13	<code>g3m</code>	ResNet-34	C10	SGDm	10^{-4}	100	$-15:-2.3026$	9 v 9	+0.666	0.094	7.08
14	<code>r50</code>	ResNet-50	C10	SGDm	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.881	0.261	3.37
15	<code>gc1</code>	ResNet-18	C100	SGDm	10^{-4}	100	$-15:-2.3026$	4 v 4	+1.640	0.245	6.71
16	<code>gm2</code>	ResNet-18	C100	SGDm	10^{-4}	100	$-15:-2.3026$	3 v 3	+1.485	0.238	6.24
17	<code>bm2</code>	ResNet-18	C10	SGD	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.978	0.086	11.40
18	<code>bm2</code>	ResNet-18	C10	RMSProp	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.631	0.149	4.23
19	<code>sm3</code>	ResNet-18	C10	AdamW	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.141	0.064	2.22
20 [§]	<code>sm4</code>	ResNet-18	C10	AdamW	10^{-4}	100	$-15:-2.3026$	3 v 3	+0.889	0.228	3.89

[†] `hz3`’s seed-5 `chunk777` run sits in a narrower clip box ($-15:-2.3026$) and on a different GPU class from its **nodewise** partner (§7 T9). The matched 5 v 5 reading is **$+0.455 \pm 0.096$** , t 4.75; the cell’s sign, magnitude and resolution are unchanged, the difference being 0.028 pp at full precision (0.027 from the rounded table entries) against a 0.086 pp se. Run R2 (§3.5) restores a matched 6 v 6.

[§] **Row 20 is the only cell in this table whose meta-optimiser is not Lion.** `sm4` runs an **RMSProp** meta-optimiser (§5.5, §7 T1); every other row in the table, and every other count-matched contrast in the corpus, runs Lion. Its own registered scorer forbids pooling it with `aw1` or `sm3`, so it enters no pool in §4.4 and appears here as a listed cell only. Rows 17 and 18 are `bm2`, the registered replication of §4.4’s two formerly single-batch levels at fresh seeds 3, 4, 5; row 19 is `sm3`, which declared an RMSProp meta and ran Lion (§6.1) and is therefore an independent replicate of `aw1` rather than the corner it was built for.

Row 5’s n is three, not six. `m12`’s two nominal halves are the same effective command line *including* `--seed`, so the batch is three seeds run twice rather than six seeds (§6.1). The point estimate is unaffected; the standard error is 0.195 rather than 0.142 and the t is 2.34 rather than 3.20.

D is positive in every cell. Eighteen of the twenty are resolved at $t \geq 3.0$; the two exceptions are `m12` at t 2.34, whose three independent seeds were run twice under two names, and `sm3` at t 2.22, which is the smallest D in the corpus ($+0.141 \pm 0.064$) and sits at the same AdamW cell as `aw1`. A further cell (`ar1`, $D = +0.697 \pm 0.118$) is **excluded** as box-void — it bound on the guards asymmetrically, in the direction that inflates D — and is reported here only so that its exclusion is visible. A **twenty-first** contrast exists and is also excluded: `gn1`’s ResNet-18/GroupNorm arms (Wu and He, 2018). Its registered scorer’s commensurability gate fired: the two halves sit 2.862 pp apart in level (BatchNorm 92.293, error budget

7.707 pp; GroupNorm 89.431, budget 10.569 pp) against a registered bar of 2.0 pp on that difference, and the scorer printed “**NO TRANSFER VERDICT IS ISSUED ... THIS IS NOT A NULL**” before reaching the contrast at all. We therefore quote no D for it here or anywhere else; its arm means are in Appendix B and its role in the design is discussed under T7 in §7. Consequently **every cell in this paper uses BatchNorm**.

Table 4 together with the excluded ar1 cell is the complete set of count-matched nodewise-versus-uniform-chunk contrasts in this corpus. None is omitted, and the excluded one is also positive. We verified this by enumeration rather than by recollection: of the 2,173 runs, every batch that ever ran a **nodewise** arm alongside a count-matched uniform-chunk arm is one of these twenty-two (the twenty above, **gn1**-GroupNorm, and **ar1**), and every other batch carrying a **nodewise** arm has no arm to match it against. The enumeration also shows that no such cell could have been lost to the admissibility gate: all 238 uniform-chunk, **nodewise1d** and **permnode** runs in the corpus outside the in-flight **rp1** batch are admissible, and the sixteen batches involved contribute 272 runs of which 272 are admissible (§8, Table 9). Two further count-matched contrasts exist in the corpus and are reported elsewhere in this paper rather than in Table 4, because neither yields a D : **bn1** ran **nodewise1d** and **chunk2325** at $m = 4,851$ without a **chunk777** arm, giving $G = +0.295 \pm 0.048$ and no D (§5.4), and **pp1**’s **permnode** arm is the alignment leg A at $m = 14,420$ (§4.6). There is no third.

A commensurability warning we obey. A percentage point is not comparable across error budgets. ResNet-18/CIFAR-10 sits on a ≈ 7 –8 pp budget and CIFAR-100 on ≈ 29 –30 pp. On **relative** error reduction (Eq. 9) the two CIFAR-100 cells carry the two **largest** D in the corpus, +1.640 pp (**gc1**) and +1.485 pp (**gm2**), and read $\rho = 0.055$ and $\rho = 0.050$ — the **fourth and third smallest of the twenty cells**, above only **sm3**’s 0.020 and **aw1**’s 0.040, and below every CIFAR-10 cell except those two, whose median ρ is 0.080. The largest is **n11**/RMSProp at 0.1241. The pp ordering and the commensurable ordering are therefore close to inverted, which is why we never average CIFAR-10 and CIFAR-100 D ’s, and never plot them on one axis. The rule is worth something measurable rather than being a stylistic preference: adding a single CIFAR-100 cell to the percentage-point $D - G$ pool of §5.4 moves the estimate by +0.044 pp and multiplies its Cochran Q from 4.52 on 7 df to 17.17 on 8 df.

4.4 The heterogeneity has a candidate moderator, the base optimiser — and a rival we can measure but not exclude

Restrict Table 4 to the fourteen cells that run the same ResNet-18 partition contrast (**nodewise** $\rightarrow$ **chunk777**, count-matched to +1 group on 14,420; rows 1–4, 6–12 and 17–19) so that the contrast itself is held fixed. **sm4** (row 20) is **excluded from every pool in this section** by its own registered scope note, which forbids pooling it with **aw1** or **sm3**; it runs a different meta-optimiser and is treated in §5.5. Those fourteen are heterogeneous: the fixed-effect pool is $+0.530 \pm 0.029$ with

$Q = 102.47$ on **13 df**, $p = 5.5 \times 10^{-16}$; DerSimonian–Laird $\tau = 0.295$ **pp** against an rms measurement se of **0.143 pp**, i.e. $I^2 = 87\%$.

Three of the fourteen were not available to the previous draft. Two are **bm2**, the registered replication of §3.5’s R3, and they are quoted from its scorer, **analysis/c97_bm2_score.py**, run unedited on the cluster where its probe records live (RULE 16; the box gate R0.5 is measured on **bm2**’s own twelve probe dirs, $T = 10,000$ records, all rails 0.0000):

```
D'_SGD base SGD +0.9780 +- 0.0858 t 11.40 3v3 -> REPLICATES
D'_RMS base RMSProp +0.6313 +- 0.1492 t 4.23 3v3 -> REPLICATES
SGD n11 +1.0353+-0.1085 | bm2 +0.9780+-0.0858 -> pool +1.0000 +- 0.0673 Q 0.172 on 1 df ->
REPLICATED
RMSProp n11 +0.9733+-0.2515 | bm2 +0.6313+-0.1492 -> pool +0.7204 +- 0.1283 Q 1.367 on 1 df ->
REPLICATED
```

bm2 is one submission of four arms at three **fresh** seeds (3, 4, 5 against **n11**’s 0, 1, 2), at **n11**’s own cell read off **n11**’s own **ARGS**: line. Independence is bought twice — a separate submission, so the batch




Figure 1: D in every count-matched cell, and the same effect made commensurable. (a) D = uniform chunk – architecture-aligned nodewise, plateau5, taken within batch, with 95% intervals from the Welch standard error of Table 4. Colour is the base optimiser, marker shape the dataset, marker size the network depth. D is **positive in every one of the twenty cells**. The horizontal rule separates CIFAR-10 from CIFAR-100: the two sit on $\approx 7-8$ pp and $\approx 29-30$ pp error budgets and a percentage point does not mean the same thing across it, so panel (a) must not be read across the rule. (b) The same twenty contrasts as a share of the aligned arm’s remaining error, $\rho = D/(100 - \text{aligned})$ (Eq. 9), which *is* commensurable. The ordering changes: CIFAR-100’s $+1.640$ pp, the **largest** D in panel (a), is $\rho = 0.055$ — the **fourth smallest** of the twenty, behind **sm3** (0.020), **aw1** (0.040) and the other CIFAR-100 cell **gm2** (0.050). The inset gives the fixed-effect pool over the fourteen cells that run the same ResNet-18 partition contrast; **sm4**, the one cell with a non-Lion meta-optimiser, is plotted but is not in that pool.

random effect is resampled, and fresh seeds, so the m12 failure mode cannot recur. The third is sm3, whose twelve rows are now in the deposited run table and which — under RULE 20, on its own ARGS: line — ran AdamW + Lion and is an independent replicate of aw1 (§6.1). n11 is **combined with bm2** and is not overwritten, and the two bm2 bases are **never pooled with each other**, both of which the scorer’s R5 forbids.

Split the fourteen by the base optimiser, the axis §5.5 was already looking at:

base optimiser	base state, from the runs’ own ARGS line	cells	batches	pooled D (pp)	within-level Q
SGD	no momentum, no second moment	2	2	$+1.000 \pm 0.067$	0.17 / 1
RMSProp	second moment 0.999, no momentum	2	2	$+0.720 \pm 0.128$	1.37 / 1
SGDm	momentum 0.99, no second moment	8	7	$+0.556 \pm 0.045$	4.21 / 7
AdamW	momentum 0.9 + second moment 0.999	2	2	$+0.189 \pm 0.052$	1.61 / 1

Between base optimisers, $Q = 95.12$ on 3 df ($p = 1.7 \times 10^{-20}$) — **92.8% of the total; within levels, $Q = 7.36$ on 10 df ($p = 0.69$).** Every level of this moderator now rests on at least two independent submissions. That is new, and it is the single most important thing bm2 and the sm3 ingest bought. Before stating what the decomposition does and does not license, we state the arithmetic that limits it, because a referee will find it otherwise.

The share is a property of the grouping, not of the label. A level with one cell contributes $Q = 0$ by construction. In the eleven-cell pool of the previous draft, *three* of the four levels were singletons, so Q_{between} was identically Q_{total} minus the eight-cell SGDm $Q = 36.4048 - 4.2060 = 32.1988$ — for **any** partition isolating those three cells, whatever it called them, and “88% is one identified moderator” therefore claimed more than the number could carry. On fourteen cells no level is a singleton: 3.15 of the present 7.36 within-level Q sits, on 3 df, in levels that previously had none, and those 3 df were a test the label could have failed — the scorer’s registered line is $Q > 3.841 \Rightarrow$ the level is not a stable quantity — and did not. The base partition is now a **restriction the data could have rejected**, and it did not reject it: all four levels are homogeneous, p 0.20–0.76. Over all 45,045 partitions of the fourteen cells with the realised shape $\{8, 2, 2, 2\}$, the base-optimiser partition ranks 14th (permutation $p = 0.00031$, against a median share of 0.333 and a maximum of 0.957). Leave-one-cell-out over the fourteen gives 89.6% to 95.7%.

The rival label, measured. The variable most likely to induce this grouping by accident is batch identity, which this paper’s own methods section makes a large random effect ($F(62, 85) = 5.47$, p 6.9×10^{-13}). Partitioning the fourteen cells by submission gives eleven levels, between $Q = 98.16$ on 10 df — **95.8%, more than the base optimiser’s 92.8%**. Nesting both inside their common refinement (base $\times$ batch, thirteen levels, within $Q = 0.21$ on 1 df) separates them:

Adding **batch given base** costs $\Delta Q = 7.15$ on 9 df, $p = 0.62$ — once you know the base optimiser, which submission a cell came from explains nothing further. Adding **base given batch** costs $\Delta Q = 4.11$ on 2 df, $p = 0.13$ — knowing the base optimiser does add something beyond submission identity, but not at $p < 0.05$.

On the eleven-cell pool the second of those two numbers was $\Delta Q = 0.05$ on 1 df, $p = 0.82$: base added nothing whatsoever beyond batch. That is what bm2 bought, and it is a smaller purchase than “the moderator is now replicated at every level” would suggest — the conditional test is **unchanged** by the sm3 ingest, because sm3 is its own batch as well as an AdamW cell. Nor is the base optimiser the best-fitting partition: cutting the same fourteen cells post-hoc into four contiguous groups in D order removes **97.5%** of Q , and the best partition of the realised shape removes 95.7%. Neither is a competitor, because both are chosen on the outcome; they are quoted so that “92.8%” is read as *this axis removes almost all of the heterogeneity*, not as *no other partition could*. Two axes that are not chosen on the outcome remove essentially none of it: β -box 0.7%, meta-stepsize 1.2%, budget 1.5%. **The base optimiser is the parsimonious description of these fourteen cells — four levels lose nothing against eleven — and it is not yet separated**

from the nuisance variable it co-varies with. We therefore write “candidate moderator”, not “identified moderator”, and we do not write that the base optimiser *carries* or *explains* the heterogeneity.

The one out-of-sample test that exists. `bm2` is the first occasion on which the moderator made a prediction before the data existed. `n11`’s levels predicted +1.0353 (SGD) and +0.9733 (RMSProp); the corpus mean predicted +0.5707 for both. Observed: +0.9780 and +0.6313, prediction errors -0.057 ($z -0.41$) and -0.342 ($z -1.17$). The base-level predictor’s RMSE is 0.245 pp against the grand mean’s 0.291 pp — **29% of the squared prediction error removed, on two points** — and on RMSProp alone the grand mean was the better predictor. This is offered as the direction the evidence points, not as a validated rule; §5.8’s out-of-sample null, whose unit is the design point and whose predictors are continuous, is unaffected.

The remaining within-level residual is not resolvable at all:

Over the eight SGDm cells — **seven separate submissions, two meta-stepsizes** (10^{-4} , 3×10^{-4}), **two budgets (100 and 300 epochs), three β -boxes and two clusters** — the pool is $+0.556 \pm 0.045$ with $Q = 4.21$ on 7 df, $p = 0.76$, $I^2 = 0\%$, $\tau = 0.000$ (one-sided 95% Q -profile upper limit **0.109 pp**, below the 0.146 pp rms measurement se of the cells themselves).

Four of those eight are submissions of the *identical* configuration — `cc1`, `mm1`, `pp1` and the BatchNorm arm of `gn1` — and they land at +0.727, +0.485, +0.581 and +0.587, $Q = 0.88$ on 3 df, $p = 0.83$. **They are not four independent replications and we no longer describe them as such.** Between them the four draw thirteen cell-seeds from **six** distinct seeds: `mm1`, `pp1` and `gn1` all use seeds 0–2, `gn1` adds seed 3, and `cc1` uses 3–5. The overlap is exact rather than nominal — `gn1-bn-ch-s0`, `mm1-ch-s0` and `pp1-ch-s0` are the same experiment, the same effective argument line after last-wins resolution with the same ENV and the same `--seed`, as are their *nodewise* partners (which `bn1` and `tw0` also re-run), and `cc1-*-s3` is the same experiment as `gn1-bn-*-s3`. The largest mutually seed-disjoint subset of the four is **two** cells: `cc1` against either `mm1` or `pp1`, which agree at $Q = 0.88$ and $Q = 0.36$, each on 1 df. Restricted to the three seeds they share, `mm1`, `pp1` and `gn1` read +0.485, +0.581 and +0.408 — a 0.173 pp spread over three measurements of one six-run experiment.

What their agreement measures is therefore re-measurement of a fixed seed set across batches, not replication across seeds, and §6.3 is where its worth is settled: with the batch variance component withdrawn ($sd_{\text{batch}} = 0.000$ pp, $F(39, 172) = 0.71$), a second submission at the same seeds is a rerun, not a replicate. $Q = 0.88$ on 3 df says that D does not move between reruns, which is worth knowing and is not nothing; it does not say that D survives a change of seed. The one comparison in this pool that does say that is `cc1` against `mm1` or `pp1`, and it agrees. The within-SGDm residual inherits the same caveat and we flag it rather than let the reader find it: four of the eight cells re-measure one seed set, which can only lower Q , so $Q = 4.21$ **on 7 df is an upper bound on homogeneity’s evidence, not a neutral test.** Leave-one-cell-out over the eight never resolves heterogeneity (Q 1.21–4.17 on 6 df, pool +0.544 to +0.603; $\tau = 0.000$ in all eight subsets), and no other axis available inside the subset competes with the base: β -box $Q = 0.76/2$ (p 0.68), meta-stepsize $Q = 0.68/1$ (p 0.41), budget $Q = 2.99/1$ (p 0.084), cluster account $Q = 0.76/1$ (p 0.38).

These fourteen cells are not one configuration, and we do not call them one. They span three step-size clip boxes — $-15:-2.3026$ (10 cells), $-30:9.0$ (3), $-25:-2.3026$ (1) — and `analysis/c87_r13_score.py`’s registered header forbids pooling across boxes, on the ground that a box change moves the optimiser and not merely the instrument. We pool anyway; §3.4 lists this as one of the paper’s three departures from a registered scorer. **The measurement that justifies it moves with the cell set, and we report that rather than only the state that flatters us.** On the eleven-cell pool, partitioning Q by box gave between-box $Q = 1.08$ on 2 df ($p = 0.58$) against within-box $Q = 35.32$ on 8 df; on the thirteen-cell intermediate pool that contained `bm2` but not `sm3` it rose to **5.14 on 2 df** ($p = 0.077$), because both `bm2` cells sit in $-15:-2.3026$; on the fourteen cells above it is **0.69 on 2 df** ($p = 0.71$). That figure is unstable because the box axis has become partly confounded with the axis under study — all four non-SGDm cells sit in one box — so we do not rest the pooling on it. The test that is not confounded is the one taken inside a single base: over the eight SGDm cells, which between them span all three boxes, between-box Q is **0.76 on 2 df**, $p = 0.68$, with pools $+0.582 \pm 0.080$ (-15), $+0.629 \pm 0.123$ (-25) and

$+0.523 \pm 0.060$ (-30). **That** is what the pooling rests on. A β -box column is carried in Table 4 and in Appendix B so the reader can redo either split.

So the finding is not “ D varies for reasons we cannot attribute”, and it is not “one identified moderator” either. It is:

Inside a base optimiser, D does not move. Under SGDm it is $+0.556 \pm 0.045$ pp across seven submissions, two meta-stepsizes, two budgets, three β -boxes and two clusters (Q 4.21 / 7, $\tau = 0.000$); under SGD it is $+1.000 \pm 0.067$, under RMSProp $+0.720 \pm 0.128$ and under AdamW with a Lion meta $+0.189 \pm 0.052$, each across two independent submissions (Q 0.17, 1.37 and 1.61, each on 1 df). **Between base optimisers the pooled level moves by a factor of 5.3**, and a four-way split on that axis leaves the fourteen cells internally homogeneous while accounting for 92.8% of their Cochran Q . **What we cannot say is that the base optimiser is the cause.** The same 92.8% would be earned by any covariate inducing the same grouping; the grouping’s chief rival, batch identity, accounts for 95.8%; and conditioning one on the other leaves base adding ΔQ 4.11 on 2 df, p 0.13 — real, and not resolved.

Say which denominator. The between-base Q of 95.12 is 92.8% of the fourteen-cell Q of 102.47, which is the pool the decomposition is computed on. The figure is not sensitive to the two cells whose inclusion is arguable: adding `m12` at its corrected se gives 92.6% (pool $+0.528 \pm 0.029$, Q 102.61 / 14, between 95.01 / 3), and restoring the withdrawn GroupNorm arm as a fifth level gives 93.2% (pool $+0.515 \pm 0.029$, Q 107.97 / 14, between 100.61 / 4). The corresponding figure on the previous draft’s eleven cells was 88.4% of a Q of 36.40, and against the legacy twelve-cell pool that contained the GroupNorm cell it was 74.6% — the “ $\approx 75\%$ ” that earlier versions of this work quoted. All are true of different denominators; none may be quoted without naming its own.

What we may not conclude from the direction. The four levels arrange themselves as a 2×2 in the base optimiser’s own state: the two levels whose base carries **no momentum term** (SGD, RMSProp) sit at $+1.000$ and $+0.720$, and the two that carry one (SGDm, AdamW) at $+0.556$ and $+0.189$. As contrasts on the four level pools, the momentum main effect is -0.488 ± 0.080 ($z = -6.09$) and the second-moment main effect -0.323 ± 0.080 ($z = -4.03$), with an interaction of -0.087 ± 0.160 ($z = -0.54$) that the design cannot resolve.

Both main effects now resolve, and that is a change from the eleven-cell reading, where the second-moment effect was -0.169 ± 0.145 ($z = -1.16$) and we drew a contrast between a resolved momentum effect and an unresolved normalisation one. With `bm2` and `sm3` in the pools that contrast is gone, and §5.5 withdraws the sentence that rested on it. What survives is that neither component *alone* is the moderator: collapsing the four levels to momentum-present / momentum-absent removes 61.1% of Q against the four-level partition’s 92.8%, and leaves a residual Q of **34.64 on 9 df** ($p = 6.9 \times 10^{-5}$) inside the momentum-present group — the SGDm-to-AdamW gap survives the collapse. **The candidate moderator is the base optimiser, not any single component of it.**

We still register the 2×2 as a prediction, not a result, for four reasons, and a referee should hold us to all four:

1. **The base optimiser is not separated from batch identity.** $\Delta Q = 4.11$ on 2 df, $p = 0.13$. `bm2` broke the `n11` co-dependence — SGD and RMSProp are no longer two halves of a single submission — which is why this number is 4.11 and not the 0.05 it was; it did not finish the job, and the `sm3` ingest does not move it, because `sm3` brings a new batch as well as a second AdamW cell.
2. **Each of the three non-SGDm levels rests on two batches, not more.** Two is enough to make a level’s homogeneity testable — it passes at all three (Q 0.17, 1.37, 1.61 on 1 df) — and not enough to estimate a between-batch variance component at that level.
3. **The momentum coefficients are not matched** (0.99 under SGDm, 0.9 under AdamW), so “momentum present/absent” is a two-point contrast in a continuous parameter.

4. **bm2 may not be asked about AdamW.** Its own registered scorer’s R5 forbids **bm2** any statement about the AdamW, SGDm or GroupNorm bases, about ResNet-34/50, about CIFAR-100, or about any other budget. Everything above that involves AdamW is carried by **aw1** and **sm3**.

The within-batch base contrast, and a phrase we may not use. **bm2** also reads the SGD-minus-RMSProp difference inside its own batch, which **n11** can also do. The scorer’s R3 verdict: $dD' +0.3467 \pm 0.1721$, $t\ 2.01 \rightarrow$ AGREE WITHIN RESOLUTION, against **n11**’s reading of the same contrast, $+0.0620 \pm 0.2739$. The scorer attaches the constraint in its own words — the agreement band of 0.40 is $1.46\times$ the planning standard error, so “AGREE WITHIN RESOLUTION is NOT an equivalence claim and may not be written as ‘identical.’” We do not write it as identical. The two momentum-free bases are not resolved apart by either batch, and neither batch is powered to say they are the same.

bm2 (R3, §3.5) was exactly the experiment this weakness named, and it has now been scored. It delivered two of the three repairs it was designed for: it resampled the batch random effect at the SGD and RMSProp levels, and it broke the **n11** co-dependence under which a batch $\times$ partition interaction peculiar to one submission would have moved both levels together with nothing in the corpus able to see it — such an interaction would now have to be reproduced, at both bases and in the same direction, by an independently submitted batch, and the between-batch spreads are -0.057 (SGD) and -0.342 (RMSProp) against a registered agreement band of 0.5. It did not deliver the third: the base optimiser is still not separated from batch identity ($\Delta Q\ 4.11$ on 2 df, $p\ 0.13$). So “the base optimiser is a moderator of D ” remains a within-corpus decomposition of fourteen measurements with one two-point out-of-sample check attached, **not an out-of-sample prediction rule** — which is why it does not contradict §5.8, where the predictors under test are continuous properties of a configuration and the unit is the design point. It should also be read against §3.4: **n11** is one of the five batches with no scorer registered before its runs existed, and it still supplies half of the SGD and RMSProp levels; **bm2**, which supplies the other half, was registered before any of its runs could exist.

The one cell we removed from this pool is the GroupNorm arm of **gn1** ($+0.202 \pm 0.137$): its own registered scorer refuses to issue a verdict on a 2.862 pp difference in level against a 2.0 pp bar, so under our own rule (§3.4) it may not be scored, and it is not reported here as a normalisation-scheme result. For the record, keeping it as a fifth level would raise the explained share to 93.2% and lower the pooled D to $+0.515 \pm 0.029$ — i.e. the decomposition does not depend on the exclusion; only our right to quote the cell does. Adding **m12** as a fifteenth cell at its corrected se gives pool $+0.528 \pm 0.029$, $Q = 102.61$ on 14 df, share 92.6%.

The second AdamW measurement is **sm3**, and it is now **in** the pool. **sm3** is void as designed (§6.1) — it declared an RMSProp meta-optimiser and, after argparse’s last-wins, ran Lion — and is salvageable only as an independent replicate of **aw1**, which is exactly what it is used for here. Its twelve rows are in the deposited run table, so a reader can re-derive them: $D = +0.141 \pm 0.064$ against **aw1**’s $+0.279 \pm 0.087$, a replicate spread of $+0.137 \pm 0.108$ ($z\ 1.27$), pooling to $+0.189 \pm 0.052$ with $Q = 1.61$ on 1 df. The previous draft kept **sm3** out of this decomposition on the ground that a number not in the deposited run table cannot be re-derived by a reader running **make reproduce**; that reason expired at the ingest and we do not replace it with another, because the alternative ground — that **sm3** is one of the five batches §3.4 names as having no scorer registered before its runs existed — does not distinguish it from **n11**, which supplies two cells of this same pool. The **sm4** scorer computes the same two-batch anchor its own way, as a 6 v 6 concatenation of the seeds — $D +0.2100 \pm 0.0559$ — and labels it, correctly, as a between-batch comparison and **not a pooled n for inference**. We use the inverse-variance pool in the table above because that is the estimator every other level in the table uses, and we report both so the difference is visible rather than buried.

A fifteenth same-contrast cell also exists and is excluded on a different ground: **sm4** (AdamW base, **RMSProp meta**, $D = +0.8893 \pm 0.2285$) is the corpus’s only partition cell with a non-Lion meta-optimiser. This pool is defined by a byte-identical contrast object at a Lion meta, and pooling **sm4** into it would silently convert a base-optimiser decomposition into a base $\times$ meta one; **analysis/c97_sm4_score.py**’s own scope note forbids pooling it with **aw1** or **sm3**. It is reported in §4.3, §4.7 and §5.5, not here.




Figure 2: The heterogeneity in D has a base-optimiser structure, every level of it is now replicated, and there is a rival we can measure. (a) The fourteen same-contrast ResNet-18 cells, grouped by base optimiser; each group’s band is its own inverse-variance pool ± 1.96 se. Every level carries **two or more independent submissions** — SGD **n11** + **bm2**, RMSProp **n11** + **bm2**, AdamW **aw1** + **sm3**, SGDm seven — and every level is homogeneous: Q 0.17 / 1, 1.37 / 1, 1.61 / 1 and 4.21 / 7 respectively, the eight SGDm cells spanning seven separate submissions, two meta-stepsizes, two budgets and three clip boxes at $\tau = 0.000$. (b) Partitioning the fourteen-cell Cochran Q (Eq. 12): **95.12 of 102.47 (92.8%) is between base optimisers**, on 3 df, $p = 1.7 \times 10^{-20}$; the within-level remainder is 7.36 on 10 df, $p = 0.69$. On the previous draft’s eleven cells the same partition read 32.20 of 36.40 (88.4%), and against the legacy twelve-cell pool containing the withdrawn GroupNorm cell, 74.6% — the “ $\approx 75\%$ ” figure. Every denominator must be named. The share is a property of the grouping and not of the label — batch identity, on eleven levels, accounts for 95.8% — so the panel is captioned as a decomposition and not as an attribution; the conditional test that distinguishes them (ΔQ 4.11 on 2 df, $p = 0.13$) is in the text. **sm4** is **not** in this figure: it is the only count-matched cell with a non-Lion meta-optimiser and its registered scorer forbids pooling it here.

4.5 Tuning each arm to its own optimum does not remove D

The referee objection ranked most likely to sink this result is that D compares two arms at one arm’s favourable meta-stepsize. We closed it inside one batch (r13, 24 runs, four arms $\times \eta \in \{10^{-4}, 3 \times 10^{-4}\} \times 3$ seeds). The registered scorer’s H3, verbatim:

```
D(ms=1e-4) = +0.681 se 0.126
D_own      = +0.591 se 0.126      (each arm at its OWN argmax)
dD         = -0.090 se 0.178 t -0.51 (16 df)
VERDICT: RULE 11 CLOSED ON R18/CIFAR-10 -- tuning each arm to its
own argmax does not resolvably move D.
```

The reportable sentence is “ D is unchanged at matched optima to within 0.356 pp”. This closes the objection **on ResNet-18/CIFAR-10 and nowhere else**: no meta-stepsize ladder exists on ResNet-34, ResNet-50 or CIFAR-100, and no ladder exists under AdamW.

dD is an upper bound in magnitude, and here is why. The two rungs are $\eta \in \{10^{-4}, 3 \times 10^{-4}\}$ and the argmax was taken on the same test accuracies that define D (§3.3). Both arms’ argmax landed on the same rung, $\eta = 3 \times 10^{-4}$, so D_{own} is itself a clean within-batch count-matched contrast and carries no cross-arm selection. What it does carry is rung selection over two rungs at $n = 3$, and a test-selected argmax rewards the noisier arm more: at the selected rung the **nodewise** arm has sd 0.164 against the **chunk777** arm’s 0.021 (plateau5, $n = 3$ each). Selection therefore pushes D_{own} down and pushes the measured shrinkage dD away from zero, so -0.090 ± 0.178 **is the largest shrinkage this two-rung ladder can produce on this metric, not an unbiased estimate of it.** The conclusion — that re-tuning does not remove D — is conservative in the direction that matters, and we state it as a bound rather than as a null.

The same batch measures D ’s sensitivity to η directly: $+0.681 \rightarrow +0.591$ across 0.477 decades, i.e. **-0.189 pp per decade of η** , within one batch. That is the number to use; the cross-batch cross-box value the record previously carried is superseded.

4.6 Alignment: a bounded null, and what it does not settle

permnode holds nodewise’s group **count** and its **exact per-tensor size multiset** and randomises only *which* weights share a group, within each tensor. Both properties are measured from the allocated **beta** on the built network by the submission script’s own guard, not inherited from a comment. The contrast was registered five-way with a symmetric band before the runs existed. The scorer’s P2, run unedited (**--selftest** 139/139 PASS), verbatim:

```
permnode<S> (m=14420, nodewise's EXACT size multiset, membership randomised)
92.003 +-0.090 (n=3)
nodewise     (m=14420, output channels)
92.012 +-0.129 (n=3)
A = permnode - nodewise = -0.009 pp   (se 0.157, t -0.06)
registered: A > +0.30 HURTS | (+0.15,+0.30] UND | [-0.15,+0.15] NULL |
[-0.30,-0.15] UND | A <= -0.30 HELPS
-> **NULL**
```

In the same batch $D = +0.581 \pm 0.141$ (t 4.11) and $B = \text{chunk777} - \text{permnode} = +0.590 \pm 0.107$ (t 5.53), and $A + B - D = 0$ exactly (Eq. 7) — a receipt on the reduction, not a finding.

What the verdict licenses, printed to the same precision as the verdict. NULL is a decision about a band; it is not an estimate. The estimate is $A = -0.009$ with a 95% interval of **$[-0.317, +0.298]$ pp**, which against this batch’s own $D = +0.581$ is **$[-55\%, +51\%]$ of D** . The smallest alignment effect this design could have detected at 80% power is **0.440 pp, i.e. 76% of D** ; its power against an alignment effect equal to *half* of D is **0.46**. The honest reading is therefore two-sided and asymmetric in its usefulness:

- **Excluded.** Alignment as the *principal* carrier of D . An effect of $+0.58$ — alignment explaining all of D — would have been detected here with probability 0.96, and was not.
- **Not excluded.** Alignment as a *contributing* carrier at up to roughly half of D , in either direction. The interval covers $+0.29$ and -0.32 .

The point estimate’s share of D , -1.6% , is quoted in the project record and should be read with the same interval attached: the share is -1.6% , 95% CI $[-55\%, +51\%]$. A ratio whose numerator is a null is not a precise quantity and we do not present it as one.

A defect in our own registration, since we require it of others. The NULL band’s half-width (0.15) is **narrower than the standard error the batch achieved** (0.157). The band was set from `mm1`’s effect-size scale — deliberately, so that the alignment and size-distribution legs would be read on one axis — and no power calculation was performed at the planned n . The consequence is computable: had the true effect been exactly zero, this batch would have returned NULL on only **66%** of realisations and an UNDECIDED wing or worse on the other **34%**. The verdict is therefore correctly *scored* and weakly *powered*, and we report it as such. The standard error is itself estimated on 3.6 degrees of freedom, with 95% interval $[0.092, 0.496]$; on a Welch t rather than a normal approximation the interval on A widens to $[-0.467, +0.448] = [-80\%, +77\%]$ of D . We print the normal-approximation interval as the headline because it is the one registered in the project record, and the t interval here so that a reader recomputing it finds no surprise.

The manipulation was not inert. The natural objection to any null is that the knob never turned. It did. Over the same nine runs, permuting membership moves the effective-sample-size field N_{eff}/m from 0.0547 ± 0.0010 to 0.0414 ± 0.0003 — **-0.0133 ± 0.0010** , $t = 12.7$, which is 70% of the whole **nodewise** $\rightarrow$ **chunk777** move in that statistic — while moving `plateau5` by -0.009 ± 0.157 . This is descriptive and carries no registered direction, and §5.2 has already shown that N_{eff}/m does **not** track accuracy. It establishes one thing only, and that thing matters here: the permutation demonstrably changed the optimiser’s internal state, so $A \approx 0$ is a statement about accuracy’s insensitivity to alignment and not about a manipulation that failed to apply.

Three scope limits, printed with the result rather than below it.

1. **Within-layer only.** `permnode` permutes *inside* each tensor, so this asks whether an output channel is special among the same-sized subsets **of its own layer**. It does not test whether *layer* boundaries matter. The registered scorer refuses to report this verdict as “architecture is irrelevant”, and neither do we.
2. **Permutation variance is confounded with seed variance.** `permnode<S>` takes S = the run seed (`bin/c77_permuted_partition.sh:447`), so seeds 0–2 are three different permutations *and* three different initialisations. $\text{se}(A) = 0.157$ is a compound of the two and this batch **cannot** decompose it. The direction is conservative — a second variance source can only widen the interval — but it means we cannot say how much of $[-0.317, +0.298]$ is draw-to-draw variation in the permutation itself, which is precisely the quantity the claim is about.
3. **One point in the design.** A is measured in **one batch, at one cell**: ResNet-18 / CIFAR-10 / SGDm base / Lion meta / $\eta = 10^{-4}$ / 100 epochs, with **nodewise** read at $\eta = 10^{-4}$ rather than at its own argmax of 3×10^{-4} . D is measured in sixteen count-matched within-batch cells across three networks, two datasets and four base optimisers. The two results are not on the same evidential footing and we do not present them as though they were.

What would settle it, and it is running. A single batch closes all three of the statistical gaps above: **R1** — **rp1**, `permnode` at three permutation draws $\times$ six seeds with the permutation seed decoupled from the run seed, plus a six-seed **nodewise** arm; **24 jobs (§3.5)**. No patch is needed — `permnode<S>`

already accepts an explicit S ; c77 simply passed the run seed. That design (i) takes $\text{se}(A)$ from 0.157 to ≈ 0.09 , which brings the MDE to ≈ 0.25 pp, below half of D and inside the pre-registered band for the first time; (ii) yields a **two-way variance decomposition** that separates permutation variance from seed variance, so limit 2 above becomes a measured number instead of a caveat; and (iii) supplies the replication A has never had. Until it reports, the sentence this paper is entitled to is the bounded one above and not a stronger one. **If rp1 returns an interval that still spans half of D , the alignment leg is to be reported as underdetermined rather than as a null**, and we register that in advance here and in that batch’s own scorer.

4.7 The prescription, and its exact scope

The practitioner’s move implied by §4.3 and §4.6 is: **give each one-dimensional tensor a single step size instead of one per element**. Measured as $T = \text{nodewise1d} - \text{nodewise}$ (Eq. 5), within batch:

Table 5: $T = \text{nodewise1d} - \text{nodewise}$, within batch, plateau5. Twelve cells under SGDm, SGD and RMSProp bases, every one resolved at $t \geq 3.0$; under AdamW the move is worth nothing measurable.

batch	network / setting	n	T (pp)	se	t
bn1	R18 / C10 / SGDm	3 v 3	+0.427	0.041	10.43
hz3 [‡]	R18 / C10 / SGDm, 300 ep	6 v 6	+0.337	0.069	4.85
r13	R18 / C10 / SGDm, $\eta \ 3 \times 10^{-4}$	3 v 3	+0.391	0.127	3.09
m12	R18 / C10 / SGDm	3 v 3 ($\times 2$ reruns)	+0.619	0.176	3.52
fa1	R18 / C10 / SGDm, $\eta \ 3 \times 10^{-4}$	6 v 6	+0.649	0.110	5.88
n11	R18 / C10 / SGD	3 v 3	+0.692	0.135	5.12
r13	R18 / C10 / SGDm, $\eta \ 10^{-4}$	3 v 3	+0.756	0.117	6.45
g3m	R34 / C10 / SGDm	9 v 9	+0.758	0.093	8.13
cc1	R18 / C10 / SGDm	3 v 3	+0.816	0.159	5.13
n11	R18 / C10 / RMSProp	3 v 3	+0.916	0.242	3.79
r50	R50 / C10 / SGDm	3 v 3	+1.049	0.317	3.31
gm2	R18 / C100 / SGDm	3 v 3	+1.363	0.151	9.00
aw1	R18 / C10 / AdamW , Lion meta	3 v 3	+ 0.091	0.078	1.16
sm3	R18 / C10 / AdamW , Lion meta	3 v 3	− 0.083	0.081	−1.03
sm4	R18 / C10 / AdamW , RMSProp meta	3 v 3	+ 0.988	0.231	4.28

[‡] hz3 matched at 5 v 5 (seed-5 trio excluded for the clip-box and hardware mismatch of §7 T9): **+0.328 ± 0.084**, t 3.89. The two r13 rungs share one batch and one clip box, as Table 4 rows 6 and 7 already note; they are two operating points, not two submissions. bn1 and m12 are the same experiment at the same seeds up to one flag Lion does not read: all twelve of their run pairs differ only by m12’s last-wins residual `--normalizer-param-meta 0.999`, whose runtime message is `args_meta includes unnecessary attributes {'normalizer_param'}`. Their T ’s differ by 0.192 ± 0.181 (z 1.06). Treat them as one measurement made twice, not two.

The move costs nothing in learned parameters — it *reduces* the number of learned quantities from 14,420 to 4,851 — and under an SGDm, SGD or RMSProp base it is worth **+0.328 to +1.363 pp** across **twelve** within-batch cells, **every one of which is resolved, the weakest at t 3.09** (r13 at $\eta = 3 \times 10^{-4}$; Welch p 0.037 on 3.94 df) and the strongest at t 10.4. Twelve cells is not twelve independent measurements, and the footnotes say why: bn1 and m12 are one experiment measured twice, so the count of distinct experiments behind the range is **eleven**, and two further pairs — r13’s two rungs and n11’s two bases — share a batch with each other rather than sitting in separate submissions.

The scope line is not where the previous draft put it. Under an AdamW base with a **Lion** meta-optimiser the move is worth nothing measurable, and that now rests on two independent batches rather than one: aw1 $+0.091 \pm 0.078$ and sm3 -0.083 ± 0.081 , pooling to **+0.007 ± 0.056** (their between-batch Q is 2.39 on 1 df, so the two do not agree closely, but neither is resolved and the pool is not). Under the **same AdamW base with an RMSProp meta-optimiser** the move is worth **+0.988 ± 0.231** (t 4.28) — the largest T in the CIFAR-10 corpus. So the exception is not “AdamW”. It is “AdamW **with a Lion**”

meta-optimiser”, and one cell is all the evidence there is on the other side of that line. §5.5 develops what this does to the mechanism story.

Two warnings travel with the **sm4** row, and both come from its own registered scorer, which classes T as descriptive only. First, T is **count-confounded by construction**: it moves 0.4731 decades of group count as well as removing the size-1 tail, so it is not a tail-removal measurement. Second, that confound is not a constant share. Decomposing T by Eq. 6 on **r13** at $\eta = 10^{-4}$ gives $T = +0.756 = 0.692 + 0.064$, where the **count** component is 8% of the effect and the tail component is the rest; the same decomposition on **sm4** gives $T = +0.988 = 0.629 + 0.359$, where the count component is **36%**. Any account that treats “merge the 1-D tensors” as purely a count reduction has the decomposition backwards, and any account that treats the count component as negligible has it backwards at this cell.

Practical significance, answered rather than volunteered. Scope item (iv) of §1 puts the objection in its sharpest form — the effect is ≈ 0.6 pp inside a method that is 1.8 to 4.2 pp behind a tuned cosine schedule — and it deserves an answer rather than a restatement. **The answer is that this result is addressed to whoever builds the step-size adapter, not to whoever trains the model**, and the two audiences get different verdicts.

For a practitioner choosing an optimiser, this paper recommends nothing. §7 T4 gives the deficit against a tuned schedule at three depths and it widens with depth; $+0.556$ pp does not close a 1.807 pp gap on ResNet-18 and is irrelevant to a 4.214 pp gap on ResNet-50. Worse, the swap is not even free in wallclock as we have implemented it. Over the same-contrast cells the uniform **chunk777** arm ran a median **1.185** $\times$ longer than the **nodewise** arm (range 1.086–1.360; pooled 74.9 against 64.4 minutes over 46 runs each), and the ratio holds inside each of the thirteen compute nodes on which both arms ran, so it is not node heterogeneity. On **hz3**, the only cell with the budget headroom to test it, spending that same extra wallclock on extra **epochs** of the aligned arm instead is worth between $+0.194 \pm 0.100$ (at **hz3**’s own 1.086 ratio) and $+0.448 \pm 0.084$ (at the other cells’ 1.19), against $+0.576 \pm 0.109$ for the swap; paired within run, the swap leads by $+0.129 \pm 0.101$, t 1.27, which does not resolve. **At equal wallclock, in the one cell where we can measure it, changing the partition and training 17% longer are worth about the same, and we cannot order them.**

For someone designing a step-size adapter, the verdict is different, and it is the reason we report the result. The 17% is an artefact of our indexing, not of the partition: the two arms learn 14,421 and 14,420 step sizes respectively, carry identical parameter counts and identical meta-state, and differ only in which weight index maps to which group. The intrinsic cost of the choice is **zero** — and it is a choice every method in this family makes and none of them measures. IDBD (Sutton, 1992), Autostep (Mahmood et al., 2012), hypergradient descent (Baydin et al., 2018), SwiftTD (Javed et al., 2024), CAM-HD (Jie et al., 2020), blockwise adaptivity (Zheng and Kwok, 2019) and MetaOptimize (Sharifnassab et al., 2025) itself all default to an architecture-aligned partition (per layer, per block, per channel) on the unstated premise that the architecture’s own decomposition is the right one. On this corpus that premise is measurably the wrong way round: at fixed group count the architecture-aligned partition is worse than an arbitrary uniform one in every count-matched cell we have, by $+0.556 \pm 0.045$ pp under an SGDm base, by $+1.000 \pm 0.067$ under SGD, and by up to $+1.640 \pm 0.245$ on CIFAR-100. It is also larger than the lever it is usually confused with: one decade of group count, at fixed partition family, is worth $+0.350$ to $+0.828$ pp (§4.2), so the partition is worth roughly two thirds of a decade to one and a half decades of count — a variable that method papers do tune.

What would make it a practitioner result, and why we do not claim it. Every measurement here is inside one meta-learning framework that trails a tuned schedule. Whether the same ordering holds in a competitive method is untested: the design that would test it — the four partitions applied as a fixed per-group learning-rate scale on a plain SGDm or AdamW run with no meta-learning at all — is named in §9 and was not run. Until it is, the correct reading of this paper is that a design parameter the field sets by architectural intuition has a measurable and consistent optimum in the opposite direction, on one framework, at CIFAR resolution.

4.8 Budget: the effect survives $3\times$ the budget; whether it decays is not resolved

hz3 ran the four arms for 300 epochs at 6 seeds, so D at 100, 200 and 300 epochs can be taken **within run**, per seed. That pairing cancels the seed, the run and the batch identically. It does **not** cancel the step-size clip box or the GPU class, because **hz3** is two submissions: its seed-5 **chunk777**, **nodewise1d** and **chunk2325** runs were resubmitted in the narrower box $-15:-2.3026$ and on an A100, while everything else — including their own seed-5 **nodewise** partner — ran at $-30:9.0$ on an L4 or a 2080 Ti (§7 T9). We therefore report the budget contrast **twice**: as submitted, and over the five clean seeds.

budget	D , 6 seeds as submitted	se	D , 5 clean seeds	se
100	+0.576	0.109	+0.662	0.083
200	+0.514	0.115	+0.575	0.119
300	+0.428	0.071	+0.455	0.081

The level is robust, and it is the claim we make. $D(300)$ reads $+0.428 \pm 0.086$ at 6 v 6 (t 4.94) and $+0.455 \pm 0.096$ at 5 v 5 (t 4.75) on the Welch estimator used everywhere else in this paper; all six seeds favour **chunk777** at 300 epochs (exact binomial $p = 0.0156$). **The partition gap is not an artefact of a 100-epoch budget.**

The trend is not robust, and we do not claim it. Within-run, $D(300) - D(100) = -0.149 \pm 0.105$, $t - 1.42$ over all six seeds and -0.207 ± 0.107 , $t - 1.94$ over the five clean ones. Neither resolves at $|t| \geq 2$, but the second is close enough that the difference matters, so we say where it comes from: 76% of the shift is at the **100-epoch** end (the six-seed $D(100)$ rises by 0.086 pp when seed 5 is removed, against 0.027 pp at 300 epochs), and seed 5’s low $D(100)$ is **not** explained by either defect — the clip box is provably inert before epoch 162, and the hardware term (§6.3) points the other way. It is an unexplained extreme value in an $n = 6$ cell. The honest statement is therefore the weaker one: **D does not grow with budget from 100 to 300 epochs, and we cannot resolve whether it decays.** We do not report flatness as a result. A box- and hardware-matched replacement trio (R2, §3.5) is registered and will settle it; until it lands this cell carries its sensitivity in the text.

One further reading must be disclosed rather than buried, because the batch’s own registered scorer prints it. `analysis/c87_hz3_score.py` fixes its primary window at **50 epochs**, not at `plateau5`’s 5, and on that window returns $D(300) - D(100) = +0.229 \pm 0.070$, t 3.27, verdict **GROWS**. The windows disagree because at $B = 100$ a 50-epoch trailing mean spans epochs 50–99 and therefore *contains* the mid-training trough documented below, while at $B = 300$ it contains none: the window injects the trough at exactly one budget. `plateau5` is this paper’s primary metric throughout and we do not switch it here — but the registered verdict is **GROWS**, our primary-window reading is a non-significant decline, and a reader is entitled to both. See Appendix A.2.

The trajectory is not monotone, and we publish it. Within run, D goes significantly *negative* in mid-training and recovers:

epoch	cc1 ($n=3$)	g3m, ResNet-34 ($n=9$)
25	$+0.625 \pm 0.170$	-0.081 ± 0.156
55	-0.560 ± 0.150 ($t - 3.73$)	-0.518 ± 0.128 ($t - 4.04$)
75	$+0.345 \pm 0.129$	—
100	$+0.727 \pm 0.146$	$+0.666 \pm 0.091$

Any claim about D is a claim about the end of training at these budgets.

5 Nine candidate mechanisms — four refuted — and two nulls

This section is the second half of the contribution, not an appendix. Of the nine, **four are refuted** (M2, M3, M6, M9), **one is narrowed** to a base–meta pairing rather than to a base alone (M4), **one is not**




Figure 3: D at $1\times$, $2\times$ and $3\times$ the budget, paired WITHIN run, and what one box-mismatched seed does to it. hz3 ran the four arms for 300 epochs at six seeds, so D can be read off the same run at three budgets, cancelling seed, run and batch identically. (a) Faint lines are the six per-seed trajectories; the two heavy lines are the arm-set means with 95% intervals. **Seed 5 is drawn in red because it is not box-matched:** its chunk arm ran in β -box $-15:-2.3026$ and its **nodewise** arm in $-30:9.0$, so that seed’s D is a cross-box difference and the other five are not. (b) The budget slope both ways. D is present and resolved at every budget, and the slope does not resolve either on all six seeds (-0.149 ± 0.105 , t -1.42) or on the five box-matched seeds alone (-0.207 ± 0.107 , t -1.94). Neither interval excludes zero, so the verdict does not turn on the contaminated seed — but its t does, and we print both rather than choosing.

separable from the axes it is aliased with (M7), and **three cannot be decided by this instrument or this design** (M1 untested, M5 inapplicable, M8 not identifiable). Two refutations rest on gates registered before their data existed (M2, M9); two are post-hoc contrasts on data collected for other purposes (M3, M6). One further candidate, M1, had its literature attribution withdrawn after we re-read the sources.

M9 is the one refutation in this paper that was designed, registered, run and read in that order. Its bar — **REFUTED** if $D \geq +0.55$ — was committed in `bin/c97_sm4_secondmoment.sh` before the batch was submitted, and it was cleared by a wide margin in the direction the mechanism said was impossible. We report it as a positive result, which is what it is.

M7 is not a fourth refutation, and the reason is worth stating. The contrast that would make it one is a within-batch separator that the relevant batch’s own registered scorer **refuses to compute** (§5.6, §7 T7). Under §3.4 we report the refusal and not the contrast, so the refutation is unavailable and M7 stands at *not separable*. That is the most expensive single application of §3.4’s rule in this paper, and it is the reason the rule is worth stating as a contribution rather than as a habit.

The refutations carry their own power limits and we print them with the verdicts. M3 in particular refuses a general *size law* at t 0.87 and cannot rule out a tail-specific effect; “refuted” there means the general form is refuted, not that the tail is exonerated (§5.3).

Table 6: Index of the nine candidates, the verdict on each, and where each is decided.

#	candidate mechanism	verdict	where
M1	$\sqrt{N}$ estimator-noise averaging over group members	untested, not refuted — our instrument measures the wrong quantity; its literature attribution withdrawn	§5.1, §2.5
M2	Meta-gradient correlation (an N_{eff}/m field) predicts accuracy	refuted on its own pre-registered gate: anti-concordant $t - 11.14$, dissociation $t - 23.26$	§5.2
M3	Degenerate size-1 groups are the carrier	refuted as a general size law: removing 100% of singletons buys $+0.115 \pm 0.133$ (t 0.87)	§5.3
M4	The size-1 tail carries it universally	narrowed to a base–meta pairing, not to a base: pooled $D - G = +0.514 \pm 0.056$ over the 8 CIFAR-10 SGDm cells and -0.061 ± 0.085 under AdamW + Lion (2 batches), but $+0.629 \pm 0.242$ (t 2.59) under AdamW + RMSProp . No individual G survives Holm on Welch df over the 14-test family	§5.4, §5.5
M5	Choi-style inclusion — a finer partition contains the coarser one once tuned	inapplicable: granularity is state, not hyperparameters; the only setting where the arms coincide is $\eta = 0$	§2.2, §4.5
M6	The base optimiser’s own second-moment normalisation	refuted: RMSProp and AdamW both carry one and differ by $+0.531 \pm 0.138$ (t 3.84) on two-batch level pools	§5.5
M9	A second-moment normaliser anywhere in the loop — base or meta — shrinks D	refuted on a pre-registered bar: at the corner where both carry one, $D = +0.889 \pm 0.228$ (t 3.89), the largest D in the AdamW family, against a registered refutation threshold of $D \geq +0.55$	§5.5
M7	D tracks the aligned arm’s accuracy level	not separable: the registered within-batch separator issued no verdict; the surviving slope is aliased with network and base optimiser (-0.119 ± 0.022 holding base, -0.386 ± 0.105 holding network, -0.187 ± 0.125 holding both)	§5.6
M8	Some summary statistic of the group-size distribution	not identifiable from this design (rank 3, one contrast type)	§5.7

The two nulls: architecture alignment (§4.6, §5.10) — a *bounded* null, $A = -0.009 \pm 0.157$, 95% CI $[-55\%, +51\%]$ of D , from one batch at one design point — and out-of-sample predictability of D (§5.8). §5.9 adds the failure of the classical remedy — hierarchical partial pooling — which is a dead *fix* rather than a dead mechanism, and is included because CAM-HD makes it the obvious question to ask.

5.1 Untested, not refuted: $\sqrt{N}$ estimator-noise averaging

The natural story is that a group of size N averages N meta-gradient samples, so a coarse group’s step-size estimate has noise $\propto 1/\sqrt{N}$ and fine partitions are noisier. We built a per-group drift instrument to test it and got the wrong sign in every fit — 6 of 6 — where a $1/\sqrt{N}$ model requires a log-log slope of -0.500 . (*That count is carried from the project record and was **not** re-derived at write time; see Appendix A.9. Nothing below rests on it.*)

We do not present that as a refutation of the variance claim, because it is not one. The instrument measures *drift* — the per-step magnitude $|\Delta\bar{\beta}|$ — which is net systematic movement, not an estimator variance; and under Lion’s sign update $|\Delta\beta| = \eta$ per step exactly, so the statistic is **censored at the meta-stepsize** (two published fits ran through arms pinned at exactly 1.000×10^{-3}). The surviving fits are 4-point log-log OLS with R^2 as low as 0.08. The only sentence this instrument supports is: *the net drift of a group’s step-size exponent does not fall like $1/\sqrt{\text{group size}}$; we did not measure the estimator variance.* The variance mechanism is **untested**, not refuted, and we say so.

Separately, and this is a correction to our own earlier writing: the attribution of a $\sqrt{N}$ noise-averaging argument to Adam-mini, Adalayer and SGG is **withdrawn** (§2.5).

5.2 Refuted: the meta-gradient correlation field predicts accuracy

A natural design principle for grouping is “put weights with correlated meta-gradients together”, which suggests an effective-sample-size statistic N_{eff}/m as a predictor. We instrumented it and registered a five-way concordance test in advance. The cc1 scorer, verbatim:

```
C2 aligned vs uniform at m=14,420
d_acc = +0.727 pp (t +3.63, RESOLVED)   d_N_eff/m = -0.0237 (t -11.14, RESOLVED)
-> **ANTI-CONCORDANT**
C3 aligned vs uniform at m=4,851, size-1 tail already gone
d_acc = +0.011 pp (t +0.08, unresolved) d_N_eff/m = -0.0529 (t -23.26, RESOLVED)
-> **DISSOCIATION**
-> **MIXED: THE FIELD IS NOT A SUFFICIENT STATISTIC.**
```

The negative is strong, not weak. The field channel resolves at $t = -11$ and $t = -23$ — it is not too noisy to read. It reads decisively, and decisively *against* accuracy on one contrast and silently on the other. This is the sharpest thing we can say to the meta-gradient-correlation line of argument: the correlation statistic that line reaches for does not track the outcome it is meant to justify.

5.3 Refuted as a size law: size-1 groups are the carrier

The most attractive explanation of §4.3 is that D is caused by the 9,610 degenerate size-1 groups in **nodewise**. Inside one batch (**ck1**) we removed **100% of a network’s singletons** by going from one step size per weight (**chunk1**, $m = 11,173,962$, 100% singletons) to one per pair (**chunk2**, $m = 5,586,981$, **0%** singletons):

chunk1 $90.979 \pm 0.131 \rightarrow$ chunk2 91.095 ± 0.024 . **Step** = $+0.115$, se **0.133**, t 0.87 — **unresolved**, against a count-only prediction of $+0.13$ to $+0.16$ from the family’s own local secants.

The ladder is smooth in $\log m$ with **no kink** at the one step where singleton fraction collapses from 100% to 0%. Reported against our own interest, and **with its power limit**: that step re-groups 100% of 11.17M coordinates, whereas the tail in **nodewise** is 0.086% of the weights. It refuses a general *size law*; it cannot rule out a tail-specific effect. The unit is the **tensor**, not the group size.

5.4 Narrowed to a base-meta pairing: the tail as a universal carrier

$D-G$ (Eqs. 1–2) isolates the size-1 tail’s contribution — G is the same contrast with the tail already removed from both arms, count-matched exactly.

Table 7: The tail’s contribution, cell by cell.

cell	D	G	$D - G$	t
cc1 (SGDm)	$+0.727 \pm 0.200$	$+0.011 \pm 0.147$	$+0.715 \pm 0.248$	2.88
r13 @ 10^{-4} (SGDm)	$+0.681 \pm 0.173$	-0.011 ± 0.087	$+0.692 \pm 0.194$	3.57
fa1 (SGDm)	$+0.629 \pm 0.123$	-0.001 ± 0.076	$+0.630 \pm 0.145$	4.35
hz3 (SGDm, 300 ep)	$+0.428 \pm 0.086$	-0.057 ± 0.061	$+0.484 \pm 0.106$	4.57
g3m (SGDm, R34)	$+0.666 \pm 0.094$	$+0.171 \pm 0.073$	$+0.495 \pm 0.119$	4.15
r50 (SGDm, R50)	$+0.881 \pm 0.261$	$+0.153 \pm 0.314$	$+0.729 \pm 0.409$	1.78
gm2 (SGDm, C100)	$+1.485 \pm 0.238$	$+0.068 \pm 0.069$	$+1.417 \pm 0.248$	5.72
n11 (RMSPProp)	$+0.973 \pm 0.251$	$+0.020 \pm 0.049$	$+0.953 \pm 0.256$	3.72
m12 (SGDm)	$+0.456 \pm 0.195$	$+0.173 \pm 0.073$	$+0.282 \pm 0.208$	1.36
r13 @ 3×10^{-4} (SGDm)	$+0.591 \pm 0.096$	$+0.217 \pm 0.128$	$+0.375 \pm 0.160$	2.34
n11 (SGD)	$+1.035 \pm 0.109$	$+0.525 \pm 0.294$	$+0.510 \pm 0.314$	1.63
aw1 (AdamW + Lion)	$+0.279 \pm 0.087$	$+0.232 \pm 0.089$	$+0.047 \pm 0.124$	0.38
sm3 (AdamW + Lion)	$+0.141 \pm 0.064$	$+0.296 \pm 0.096$	-0.155 ± 0.116	-1.34
sm4 (AdamW + RMSPProp)	$+0.889 \pm 0.228$	$+0.261 \pm 0.081$	$+0.629 \pm 0.242$	2.59

m12’s three standard errors are computed on **three** independent seeds, not six runs: its two nominal halves are the same command line including `--seed`, so the batch is three seeds run twice rather than six seeds (§6.1). Averaging within seed before differencing gives $D + 0.456 \pm 0.195$ (t 2.34), $G + 0.173 \pm 0.073$ (t 2.38), $D - G + 0.282 \pm 0.208$ (t 1.36). **Every point estimate is unchanged**; only the standard errors move. hz3 matched at 5 v 5 for the seed-5 mismatch of §7 T9 reads $D + 0.455 \pm 0.096$, $G - 0.050 \pm 0.075$, $D - G + 0.506 \pm 0.121$, t 4.16 — same verdict.

The G tests, corrected for multiplicity. G is measured once per count-matched cell, so the values above are a family and the ones that reach nominal significance should not be read one at a time. We declared the family as the G leg of every Table 4 cell that has one — twelve tests, fixed before the correction was computed — and applied Holm–Bonferroni step-down (Holm, 1979) at $\alpha = 0.05$. **Table 4 has since grown by four cells, two of which have a G** (sm3 and sm4; bm2 ran only `nodewise` and `chunk777`, and its registered scorer forbids any G or tail statement from it), so the same rule now yields **fourteen** tests. We report the pre-specified twelve first, unchanged, and then the fourteen, so that the effect of enlarging the family is visible rather than absorbed. Each p is the two-sided Welch test on the two arm means; $p(z)$ is the normal approximation, which is the quantity the paper’s t columns report and is anti-conservative at 2–4 degrees of freedom. Each Holm column is computed on its own ordering.

At most three of the pre-specified twelve reach nominal $\alpha = 0.05$ and none survives Holm. On the normal approximation the three are aw1 (t 2.62), m12 (t 2.38) and g3m (t 2.34), with smallest adjusted $p = 0.11$; on the Welch degrees of freedom only g3m reaches nominal α , with smallest adjusted $p = 0.39$.

On the family of fourteen, the Welch verdict is unchanged and the normal-approximation verdict is not. Adding sm3 ($G = +0.296 \pm 0.096$, t 3.07, Welch df 3.5) and sm4 ($G = +0.261 \pm 0.081$, t 3.23, Welch df 3.7): on Welch degrees of freedom **nothing survives Holm** (smallest adjusted $p = 0.46$, g3m); on the anti-conservative normal approximation **sm4 (adjusted $p = 0.017$) and sm3 (adjusted $p = 0.028$) both survive**, which no tabled cell did at twelve. We record that without resting anything on it, because the normal approximation is the wrong reference at 3.5 degrees of freedom and we say so elsewhere in this paper. The enlarged family is also heterogeneous where the pre-specified one was not — fourteen-cell G pool $+0.093 \pm 0.022$, $Q = 28.25$ on 13 df ($p = 0.0084$), against the twelve-cell $+0.067 \pm 0.023$, $Q = 18.21$ on 11 df ($p = 0.077$) — and the two cells that make it so are the two new AdamW-base ones.

We therefore make no claim that G is resolved in any individual cell, and an earlier version of this section, which rested on G being resolved under AdamW, is withdrawn as a per-cell claim.

What survives the correction is not a cell but a **contrast between base optimisers**, which is a single pre-specified comparison and needs no family correction: $D - G$ is $+0.047 \pm 0.124$ (t 0.38) under AdamW

Table 8: The twelve-test G family under Holm–Bonferroni. At most three reach nominal $\alpha = 0.05$ and none survives Holm in either convention.

cell	G	se	t	df	p (Welch)	p Holm	$p(z)$	p Holm (z)
g3m (SGDm, R34)	+0.171	0.073	2.34	16.0	0.033	0.39	0.0195	0.20
aw1 (AdamW)	+0.232	0.089	2.62	3.5	0.067	0.74	0.0089	0.11
m12 (SGDm)	+0.173	0.073	2.38	4.0	0.076	0.76	0.0171	0.19
r13 @ 3×10^{-4}	+0.217	0.128	1.69	3.9	0.168	1.00	0.091	0.73
n11 (SGD)	+0.525	0.294	1.79	2.8	0.177	1.00	0.074	0.67
gm2 (SGDm, C100)	+0.068	0.069	0.99	3.5	0.386	1.00	0.322	1.00
hz3 (SGDm, 300 ep)	−0.057	0.061	−0.93	6.2	0.389	1.00	0.354	1.00
r50 (SGDm, R50)	+0.153	0.314	0.49	3.9	0.653	1.00	0.627	1.00
n11 (RMSProp)	+0.020	0.049	0.41	3.4	0.710	1.00	0.686	1.00
r13 @ 10^{-4}	−0.011	0.087	−0.12	2.1	0.913	1.00	0.902	1.00
cc1 (SGDm)	+0.011	0.147	0.08	2.4	0.944	1.00	0.938	1.00
fa1 (SGDm)	−0.001	0.076	−0.01	9.8	0.993	1.00	0.993	1.00

against a fixed-effect pool of $+0.514 \pm 0.056$ (t 9.2) over the eight CIFAR-10 SGDm cells (Q 4.52 on 7 df; restricted to the six ResNet-18/CIFAR-10 cells it is $+0.514 \pm 0.064$, the same point estimate). Adding the CIFAR-100 cell would give $+0.558 \pm 0.055$ with Q 17.17 on 8 df, which §4.3’s commensurability rule forbids and which is that rule doing measurable work rather than being a stylistic preference. The tail accounts for essentially none of D under AdamW and for essentially all of it under SGDm, and that difference — not any single G — is the result.

The independent AdamW measurement is **sm3**, and its runs **are** now in the deposited run table, so it enters the table above rather than sitting beside it: $G = +0.296 \pm 0.096$, $D - G = -0.155 \pm 0.116$ at the same cell as **aw1** (§6.1). The two AdamW+Lion batches pool to $D - G = -0.061 \pm 0.085$ (Q 1.41 on 1 df); on the 6 v 6 concatenation of their seeds the same quantity reads -0.054 ± 0.082 (t −0.66). Under an AdamW base with a Lion meta-optimiser, the size-1 tail contributes nothing measurable to D , and that now rests on two independent submissions rather than one.

One G contrast exists in the corpus outside this table, and we name it so the family is auditable rather than convenient. **bn1** ran **nodewise1d** and **chunk2325** ($m = 4,851$ in both arms) but no **chunk777** arm, so it yields a G and no D : $G = +0.295 \pm 0.048$, the largest G in the corpus and the only large positive G under an SGDm base — which cuts against “under SGDm, removing the tail removes the gap”. (**sm3**’s G , which an earlier version of this section also listed here, is now a tabled cell.) Enlarging the family to all fifteen and re-running Holm changes the count but not the verdict on Welch degrees of freedom, where **nothing survives**; on the normal approximation **bn1** (adjusted $p = 1.2 \times 10^{-8}$), **sm4** (0.017) and **sm3** (0.028) survive. The fifteen-test family is heterogeneous (G pool $+0.128 \pm 0.020$, $Q = 42.98$ on 14 df, $p = 8.6 \times 10^{-5}$), driven by **bn1**; **bn1** and **cc1** are the same configuration in different batches and their G values differ by 0.284 ± 0.155 (t 1.83), i.e. not resolved, which is what §6.3’s cross-batch bound predicts. We disclose it rather than leave a referee to find it.

Under SGDm, removing the tail removes the gap. **Under AdamW with a Lion meta it does not**: $D - G$ is $+0.047 \pm 0.124$ in **aw1** and -0.155 ± 0.116 in **sm3**, pooling to -0.061 ± 0.085 , against $+0.514 \pm 0.056$ pooled over the eight CIFAR-10 SGDm cells. G itself is positive under AdamW ($+0.232 \pm 0.089$ and $+0.296 \pm 0.096$) but survives Holm over the G family on Welch degrees of freedom in neither cell, so the dependence is carried by the $D - G$ contrast, which is one pre-specified comparison, not by a per-cell G verdict.

The dependence is not on the base optimiser alone, and this is new. Holding the base at AdamW and everything else at **aw1**’s cell, and changing only the meta-optimiser from Lion to RMSProp:

statistic	AdamW + Lion (aw1, sm3)	AdamW + RMSProp (sm4)	difference
D	$+0.189 \pm 0.052$	$+0.889 \pm 0.228$	$+0.700 \pm 0.234$
G	$+0.261 \pm 0.065$	$+0.261 \pm 0.081$	-0.001 ± 0.104
$D - G$	-0.061 ± 0.085	$+0.629 \pm 0.242$	$+0.690 \pm 0.257$

G does not move at all — the two readings agree to a thousandth of a percentage point — and the whole of the meta-optimiser’s effect on D lands in the size-1 tail component. The sm4 scorer reaches the same two conclusions independently and states them as its own registered verdicts: **G PERSISTS**. The exactly count-matched contrast stays resolved when the meta also carries a second moment, as it is under AdamW+Lion (aw1 +0.232, sm3 +0.296). G is base-dependent, not meta-dependent, at this corner. And on the mechanism contrast: **MECHANISM CONTRAST RESOLVED** at $|t| \geq 2$ (interval $[+0.154, +1.104]$). Under AdamW+Lion it is **NOT** resolved in either batch (aw1 +0.047, sm3 -0.155), so a resolved $D-G$ here is a finding about the meta-optimiser.

What that costs the tail story. M4 was narrowed in the previous draft to “an SGDm base with a Lion meta-optimiser”, with AdamW as the counter-example. The counter-example does not survive as a statement about the base: at an AdamW base the tail carries nothing under Lion and carries $+0.629 \pm 0.242$ under RMSProp. **The tail’s contribution is a property of the base–meta pairing, not of the base**, and this corpus has measured only two of the four pairings it would take to say which side dominates — it has no SGDm+RMSProp cell at all. The two right-hand columns of the table above are **cross-batch** comparisons and carry the batch floor this paper measures elsewhere ($F(62, 85) = 5.47$, $p \ 6.9 \times 10^{-13}$), which neither pairing nor Welch removes; they are labelled comparisons, not registered tests, and the sm4 registration explicitly forbids pooling sm4 with aw1 or sm3. What is a registered test is sm4’s own within-batch $D - G = +0.629 \pm 0.242$ ($t \ 2.59$), and it is resolved. The pre-registration for that batch said, before the data existed, that “a *RESOLVED* non-null G would say the tail story is base-dependent and must be re-scoped.” The re-scoping is warranted on the contrast; the word “resolved” is not, and we do not use it.

“The gap lives in the degenerate size-1 tail” may not be written as a general claim, and may not be scoped to a base optimiser either. It holds under an SGDm base with a Lion meta-optimiser (eight cells, seven submissions) and under an AdamW base with an RMSProp meta (one cell); it fails under an AdamW base with a Lion meta (two cells). The scoping variable is the pairing, and three of its four cells are one batch each.

The same batch’s D itself is **UNRESOLVED** on its own registered rule: `analysis/c88_scorers.py`
`--score aw1` prints

```
{'D_adamw': 0.279, 'se': 0.087, 't': 3.19,
'rule': "UNRESOLVED at n=6. Report the interval. Do NOT re-cut the
data, and do NOT describe it as 'partially transferring'."}
```

The registered bar was a 95% lower bound above +0.30; the realised bound is **+0.108** (normal approximation; **+0.038** on Welch df 4), so the bar is not cleared. The registration also specified 6 seeds where only 3 reached the contrast, so the primary ran at half its registered power — a registration deviation disclosed again, with its cause, in §8. D under AdamW is reported as an interval, +0.279 [0.108, 0.450], **not as a transfer**.

5.5 Refuted twice: second-moment normalisation, in the base (M6) and in the meta (M9)

M6, the base side, post-hoc. If the mechanism were “a base optimiser that already normalises per coordinate does not need the partition to do it”, then RMSProp and AdamW — both carrying a second moment — should behave alike. They do not. On the two-batch level pools of §4.4,




Figure 4: D splits into a tail-free part G and a size-1-tail part $D - G$, and the split depends on the base-meta pairing. (a) For every batch that ran all four arms, D is drawn as the sum of G (hatched: the same uniform-versus-aligned contrast with the size-1 tail *already removed from both arms*, count-matched exactly at $m = 4,851$) and $D - G$ (solid: what the tail contributes). The black tick and whisker are D itself with its 95% interval. Under an SGDm base with a Lion meta the hatched part is ≈ 0 and the tail carries essentially all of D . Under **AdamW with a Lion meta** it does not: G is $+0.232 \pm 0.089$ (aw1) and $+0.296 \pm 0.096$ (sm3) while $D - G$ is $+0.047 \pm 0.124$ and -0.155 ± 0.116 , pooling to -0.061 ± 0.085 . Under **AdamW with an RMSProp meta** (sm4, the only non-Lion cell in the corpus) G is unchanged at $+0.261 \pm 0.081$ while $D - G$ returns to $+0.629 \pm 0.242$ (t 2.59) — G does not move with the meta-optimiser and the tail component does. (b) $D - G$ alone, with 95% intervals and each cell’s t . CIFAR-100 is placed below the rule for the reason given in Figure 1. The pooled SGDm/CIFAR-10 value is $+0.514 \pm 0.056$ over eight cells (Q 4.52 on 7 df); we quote that pool rather than cc1’s $+0.715$, which is the maximum.

$D(\text{RMSProp}) - D(\text{AdamW}) = +0.531 \pm 0.138$, t 3.84; on the single-batch reading the previous draft quoted, the independent `n11-versus-aw1` contrast, it was $+0.694 \pm 0.266$, t 2.61. The refutation is the same and it is now carried by four batches instead of two. It remains a **post-hoc contrast** on data collected for other purposes, and the AdamW half of it comes entirely from `aw1` and `sm3`: `bm2`'s registered scorer forbids `bm2` any statement about AdamW, and none is made.

M9, the meta side, pre-registered. M6 leaves the strongest form of the normalisation story standing. Every base-side comparison in this corpus varies the base while a **Lion** meta-optimiser sits above it, and Lion's sign update carries no second moment at all. So the claim " D shrinks wherever a second-moment normaliser sits, base **or** meta" had never been tested at the one corner that decides it: an AdamW base with an RMSProp meta, where both components carry one. That corner predicts the *smallest* D in the corpus. `sm3` was built to run it and did not — its own `ARGS`: line carries `--alg-meta` twice and `argparse` kept the last occurrence, so it ran Lion (§6.1). The corner was therefore untested, not refuted, and it stayed that way for a full cycle.

`sm4` runs it. Four granularity arms $\times$ three seeds at `aw1`'s own cell — ResNet-18/CIFAR-10, $\alpha_0 = 10^{-3}$, $\eta = 10^{-4}$, box $-15:-2.3026$, 100 epochs, `AUGMENT=1` — changing exactly one flag, `--alg-meta Lion` $\rightarrow$ `--alg-meta RMSProp`, with RMSProp's own attribute set live and Lion's two attributes set to the code's -1 sentinel so they are dropped before the optimiser is built. The decision rule was frozen in the batch script before submission, and we quote it from the script rather than paraphrasing it:

```
Mechanism candidate 8 says D shrinks wherever a second-moment normaliser sits. [...] This corner
carries a second moment on BOTH sides, so the pattern predicts the SMALLEST D in the corpus.
REFUTED if D >= +0.55.
CONSISTENT if D <= +0.279 (aw1's own D) with a 95% interval below the bar.
UNDECIDED otherwise -- report the interval, add no seeds to this batch, and re-register a fresh
replicate instead. All three outcomes are publishable and the scorer prints the one that fires;
nothing here is chosen after the fact.
```

Three gates fire before the contrast. The `ARGS` gate (RULE 20) read all twelve runs' own `ARGS`: lines: `CLEAN --- every run's own ARGS line matches the declared design, no repeated flag, no Lion attribute live` — which is the check `sm3` would have failed. The box gate, measured on `sm4`'s own probe records, reports `coord_lo` and `coord_hi` of exactly 0.00000 on all four arms. And an operating-point gate, blocking and registered in advance because no RMSProp meta-optimiser had ever run anywhere in this corpus and its meta-stepsize was matched to `aw1` rather than tuned: `ON ANCHOR. 12-run mean plateau5 91.567 against the AdamW+Lion anchor 93.161 (-1.594 pp, bar 3.0)`. The cell is comparable. Had it fired, every contrast in the batch would have been demoted to descriptive and labelled off-anchor — reported, not discarded.

The verdict, from `analysis/c97_sm4_score.py` run unedited:

```
D = chunk777 - nodewise = +0.8893 +- 0.2285    t +3.89 (df 2.3, p 0.0494)    n 3v3
95% CI [+0.4415, +1.3371]
VERDICT: REFUTED. D >= +0.55 at the corner where BOTH the base and the meta carry a second-moment
normaliser. Mechanism candidate 8 -- 'D shrinks wherever a second moment sits' -- is false, and
the AdamW/RMSProp pattern in the corpus needs another explanation. This is a POSITIVE result and
must be reported as one, not buried.
```

The corner that predicted the smallest D in the corpus produced the largest D in the AdamW family: $+0.889$ against `aw1`'s $+0.279$, `sm3`'s $+0.141$, and their pool of $+0.189 \pm 0.052$ — a between-batch difference of $+0.700 \pm 0.234$. Whatever the AdamW-versus-RMSProp pattern in §4.4 is, it is not a count of second-moment normalisers in the loop.

Scope, carried with the verdict because the scorer carries it. The batch runs at `aw1`'s meta-stepsize of 10^{-4} , **not at an RMSProp-meta optimum** — no meta-stepsize ladder exists under an RMSProp meta, because this corpus contains no other RMSProp-meta run of any kind. It is ResNet-18 / CIFAR-10 / 100 epochs / box $-15:-2.3026$ only. And it is **not pooled with `aw1` or `sm3`**: different meta-optimiser, different

batch. One cell refutes a claim that was stated over the whole grid; it does not establish anything about the RMSProp-meta cell that a ladder would.

What this does to §4.4’s 2×2 . The previous draft read the four-level base decomposition as agreeing with M6 about which component of the base optimiser is not the axis, on the strength of an unresolved second-moment main effect (-0.169 ± 0.145 , $z = -1.16$). On the fourteen-cell pools that main effect resolves (-0.323 ± 0.080 , $z = -4.03$) and the agreement is gone. **We withdraw the claim that the two analyses agree.** What both still support is the weaker and more useful statement: neither momentum nor second-moment normalisation, taken alone, is the moderator — the momentum collapse leaves a residual Q of 34.64 on 9 df — and M9 shows that adding a second moment to the *meta*-optimiser moves D in the direction opposite to the one the normalisation story requires.

5.6 Not dead, not separable: D and the aligned arm’s accuracy level

Over all 10 design points, regressing D on the aligned arm’s `plateau5` gives slope -0.044 ± 0.011 , $t = -4.06$, $r = -0.821$ — which looks like a law and is not one, because level and dataset are the same column: the single CIFAR-100 design point sits at level 70.44 and the nine CIFAR-10 design points at 89.63–92.98.

Inside CIFAR-10 the slope is -0.161 ± 0.067 , $t = -2.38$, $r = -0.669$ over 9 design points (exact permutation over all $9! = 362,880$ orderings, $p = 0.0454$). We report that it is nominally resolved, and then report the four things that stop it being a mechanism.

- **It is at the resolution floor by construction.** At 9 design points the two-sided 5% critical correlation is $|r| = 0.666$. The realised $|r|$ is 0.669. A result that clears its own critical value in the third decimal place is a coin landing on its edge, not a law.
- **It is aliased with the network and with the base optimiser, and the alias is the whole effect.** The two lowest-level CIFAR-10 design points are `r50` (a different network) and `n11/SGD` (a different base optimiser). Holding the base fixed at `SGDm` (6 points, spanning ResNet-18/34/50) gives -0.119 ± 0.022 ; holding the network fixed at ResNet-18 (7 points, spanning four bases) gives -0.386 ± 0.105 ; holding **both** fixed — the only four points where “level” varies with nothing else structural — gives -0.187 ± 0.125 , $t = -1.49$, exact permutation $p = 0.333$, **unresolved**. A slope whose magnitude moves by $3.2\times$ depending on which confound you hold, and which vanishes when you hold both, is measuring the confounds.
- **An instrument that does not share an arm with D does not resolve it.** D and level share the `nodewise` arm. The mechanical slope this induces is -0.017 (the mean sampling variance of a `nodewise` arm mean, 0.01808, over the variance of level across the nine points, 1.07971), one tenth of -0.161 , so the artefact is not the explanation. But replacing level by an independent instrument — the mean of the `chunk2325` and `nodewise1d` arms, neither of which enters D — gives -0.167 ± 0.101 , $t = -1.65$ within CIFAR-10 and does not clear the bar.
- **The one within-batch contrast left in the corpus does not resolve it either.** `r13` ran both meta-stepsizes in one batch: level moves $+0.599$ pp ($91.908 \rightarrow 92.507$) and D moves -0.090 ($+0.681 \rightarrow +0.591$), an implied within-batch slope of -0.150 ± 0.331 ($t = -0.45$). Same sign as the between-point slope, and an interval that contains zero and all three sub-slopes above. It is also confounded with η , which is why it is a consistency check and not a test.

The registered within-batch separator issued no verdict. `gn1` was built to answer exactly this question — one batch, one variable changed (BatchNorm $\rightarrow$ GroupNorm), batch cancels — and its pre-registered commensurability gate fired: the two halves sit on error budgets of 7.707 pp and 10.569 pp, a $1.37\times$ ratio against a registered bar of 2.0 pp, so a percentage-point contrast between them is not defined. Its scorer prints **“NO TRANSFER VERDICT IS ISSUED ... THIS IS NOT A NULL”** and halts before computing either contrast. An earlier version of this paper quoted that batch’s face values as a falsification of the level model — a within-batch dD of -0.385 ± 0.205 , described as wrong-signed by 2.2 to 7.0 se. We

withdraw that, and we withdraw the falsification with it. **Both legs of the earlier argument came from the withdrawn cell:** it was the within-batch test, and it was also the low-level point that flattened the within-CIFAR-10 slope, which without it moves from -0.021 ± 0.077 ($t = -0.27$) to the -0.161 ± 0.067 above. The paper is weaker here than the previous draft claimed, and in the direction the level model predicts.

What is left is an honest negative-space statement, not a dead mechanism. The accuracy level of the aligned arm is not separable from the network and the base optimiser in this corpus, at 9 CIFAR-10 design points and a critical $|r|$ of 0.666. We do not claim it is a carrier; we do not claim it is not; and §5.8 shows that as a *predictor* it is the worst of the four models we tried, out of sample. The design needed to separate it — a level contrast at fixed network, fixed base and commensurable error budget — is stated in §9.

5.7 Not identifiable (and un-killable): a summary statistic of the size distribution

We constructed exact group-size multisets for every (network, granularity) cell we could instantiate — 150 of 163 attempted, the rest being block partitions applied to the wrong network depth, which the optimiser rejects — by instantiating the optimiser on CPU and reading sizes off its own allocated state, with two receipts asserted on every cell ($\sum \text{size} \times \text{count} = \sum p.\text{numel}()$; $\sum \text{count} = \text{beta}$ elements allocated) passing on all 150, and reproducing every group count in this paper exactly. We then scored twenty candidate statistics — mean/median/sd/CV/skew/min/max, harmonic and geometric mean, sd of log size, categorical and weighted entropy, inverse participation ratio, Gini, singleton fractions by group and by weight, threshold fractions, and a heterogeneity index.

An earlier analysis nominated the degeneracy indicator with an apparently decisive fit. **That positive claim is deleted, and the reason it is deleted is the finding.** `frac_groups_size1` takes exactly **two values** across the entire count-matched design — 0.6664 on **nodewise**, 0.0000 on all three other arms, on every network. Re-running the analyst’s own shape test with (a) that statistic, (b) a bare 1/0 “is this arm **nodewise**” dummy, (c) an arbitrary two-level regressor, and (d) the plain test that the three non-**nodewise** arms share a mean gives χ^2 values **identical to five figures**. The fit *is* the three-arm equality test, relabelled; a regressor whose fit is invariant to its own values is not a predictor.

The structural limit governs everything above. At fixed count the corpus contains exactly one contrast type — architecture-aligned versus uniform chunk — plus one zero-dose control (**permnode**). Over the four count-matched arms the candidates are collinear at $|r| \geq 0.93$ and the design matrix has **rank 3**, so at most two statistics are jointly identifiable and all twenty collapse into about two equivalence classes. Cross-validation cannot catch this: leave-one-batch-out leaves the arm structure intact in every fold. The licensed sentence is therefore a **limit**, not a search result:

At fixed group count the corpus contains a single contrast type, so every candidate summary of the size distribution reduces to an indicator for the architecture-aligned arm and none is identifiable. The size distribution remains the surviving carrier; **which** property of it is an open question requiring new runs of a kind this corpus does not contain.

§9 states the experiment that would break the degeneracy.

5.8 Null: D is not predictable from configuration properties

The sixteen D cells of Table 4 collapse to **10 distinct design points** (the six identical ResNet-18/SGDm/ $\eta = 10^{-4}$ /100-epoch batches are one point; the two CIFAR-100 batches are one). Leave-one-design-point-out, refitting each single-predictor model on the held-in 9:

model	LOO RMSE	vs “predict the corpus mean”
mean (baseline)	0.3859	—
$D \propto k \cdot \log(\text{headroom})$	0.2693	−30%
CIFAR-100 dummy	0.3765	−2%
$D \propto k \cdot \text{headroom}$	0.3796	−2%
$D \propto \text{level (OLS)}$	0.8379	+117% WORSE

None of this is a result, and here is why. The margin is dominated by the single CIFAR-100 fold: the mean errs by -0.888 there and $k \cdot \log(\text{headroom})$ by -0.457 , and restricted to the nine CIFAR-10 folds the best model wins by 0.040 RMSE ($0.2791 \rightarrow 0.2395$, -14% ; the level model, 0.2358 , -16%), with a sign test of $7/9$, two-sided $p = 0.180$. Over all ten folds the sign test is $8/10$, $p = 0.109$. The functional form is itself the winner of about ten candidates scored on the same points, so any apparent improvement is a best-of-ten selection statistic before it is anything else. And the power bound is decisive: **at 10 design points a predictor needs $|r| \geq 0.632$ — it must explain $\geq 40\%$ of the between-design-point variance — to be visible at $p < 0.05$; seeing $|r| = 0.4$ would need ≈ 25 design points, which is not reachable by brute force.**

The level model earns its own sentence. Fitted on the nine CIFAR-10 points it predicts $D(\text{CIFAR-100}) = +4.11$ against $+1.56$ observed, an error of $+2.55$ pp — nearly three times the error of simply predicting the corpus mean. A covariate that is nominally resolved *within* CIFAR-10 (§5.6) and catastrophic the moment it is asked to leave it is a within-regime association, not a law.

Three aliasings make the corpus look richer than it is. (i) `frac_groups_size1`, `log m`, `log params`, `log(n classes)` and the CIFAR-100 dummy are **not five predictors**: the singleton fraction is pinned near $2/3$ by a structural identity (every convolution has `bias=False` and is followed by one affine norm), taking 0.66644 on every CIFAR-10 cell and 0.66438 on the CIFAR-100 cells, so regressing D on it is regressing D on the dataset dummy. Anyone who writes “*D scales with the singleton fraction*” from this corpus has written “*CIFAR-100 is different*”. (ii) m and `params` are 1-vs-9 leverage contrasts resting on one architecture point each. (iii) G shares a batch with D and has no predictive content anyway.

The draft sentence: *D is real, replicated across 20 count-matched within-batch cells. Its variation across configurations is not unattributed: a four-way split on the base optimiser accounts for 92.8% of it (§4.4, where the rival label is measured too), and within a fixed base D is homogeneous ($\tau = 0.000$, 95% upper limit 0.109 pp). What remains unpredictable is the **continuous** part — no measurable scalar property of a configuration predicts D out of sample better than the corpus mean by a margin this design, at 10 design points, can resolve.*

The two statements are not in tension, and the estimands are worth separating explicitly. §4.4 is a **within-corpus decomposition** over a categorical factor with three singleton levels; §5.8 is **out-of-sample prediction** over continuous covariates at the design-point level. A factor can account for most of the observed heterogeneity and still support no prediction rule, and here it does exactly that.

5.9 Dead fix: hierarchical partial pooling

Covered in §2.4 with numbers. Summarised here so the section is complete: three pooling operators were implemented. The `zpool` figures in §2.4 were re-derived at write time on `plateau5`; the `M0/M1` statements below are carried from the project record and were **not** (Appendix A.9). `M0 shrink` is withdrawn because the operator saturates (it applies every step, so any $\lambda \geq 0.001$ has a half-life ≤ 693 steps against $\approx 50,000$ — the sweep measured full pooling three times over); `M1 additive r` shows an interior optimum that is largely an α_0 escape-rate artefact and transfers to no other architecture or dataset; and the confound-free `zpool` rescues an over-fine partition by $+10.04$ pp and still lands 2.18 pp below plain six-block. The r -dial is a meta-learning-rate knob wearing a pooling costume: under Lion every increment is exactly $\pm\eta$, so pooling divides the effective meta-stepsize by $|2p - 1|$.

We therefore do not propose a new pooling design. CAM-HD’s remedy is the right classical instrument and we could not make it beat a good partition.

5.10 Bounded null: alignment (repeated here because it is a deliverable)

§4.6. $A = -0.009 \pm 0.157$, $t = 0.06$, scored NULL against a band registered in advance, with a 95% interval of $[-0.317, +0.298]$ pp = $[-55\%, +51\%]$ of D and an MDE of 0.440 pp = 76% of D . It excludes alignment as the principal carrier and leaves a moderate contribution open, from one batch at one design point, with permutation variance confounded with seed variance and with a registration defect we print (band half-width $0.15 < \text{realised se } 0.157$). Read at that strength it is still what redirected this programme from architecture alignment to group-size homogeneity, and it is still worth as much as the positive — a negative reported with its resolution is a deliverable; a negative reported without one is a claim we cannot support. R1 (rp1, 24 jobs, in flight, §3.5) is designed to halve the interval and to decompose the two variance sources.

6 Measurement discipline: what our own instrument did to us

This section reports four defects in our own process. Each cost a claim; each is cheap for a reader to avoid.

6.1 Two batches that failed silently on their own axis

m12 (24 runs) was built to fill the corpus’s largest external-validity hole: **every count-matched partition run uses a Lion meta-optimiser**. It was to add meta = Adam and meta = RMSProp at a fixed SGDm base. It did not. The submission script composes the base/meta flags into a variable and then appends a second `--alg-meta Lion` after it, and `argparse` takes the last occurrence. The runs’ own logged argument line reads:

```
--alg-meta Adam --normalizer-param-meta 0.999 --momentum-param-meta 0.9 --weight-decay-meta 0
--alg-meta Lion --momentum-param-meta 0.99 --Lion-beta2-meta 0.9 --weight-decay-meta 0
```

and the runtime warns `args_meta` includes unnecessary attributes `{'normalizer_param'}` — Lion does not use it. **All 24 m12 runs are meta = Lion**. Worse, after last-wins its two nominal halves resolve to **the same command line including --seed**, differing only in an output path. m12 is therefore three seeds run twice, not six seeds, and its standard errors are computed on three groups (§3.3, *Duplicate runs*). The At the time m12 was audited the corpus census confirmed the hole was total: every partition-programme run in the corpus was meta = Lion, across four base optimisers. sm4 has since opened it by exactly twelve runs — see §7 T1 for the census under an explicit rule and for how little one cell closes.

One cycle later, the same defect voided a second batch. sm3 (12 runs) was built to test the one untested corner of the base $\times$ meta grid — an AdamW base with an **RMSProp** meta-optimiser, where both components carry a second moment. Its own ARGs: line reads `--alg-meta RMSProp ... --alg-meta Lion`, so it ran Lion. The second-moment corner was therefore **untested, not refuted**, and the batch is salvageable only as an independent replicate of aw1: $D = +0.141 \pm 0.064$, $G = +0.296 \pm 0.096$, at the same box, α_0 , η and budget. Its twelve rows are now in the deposited run table, so a reader can re-derive both numbers; it is Table 4 row 19, it is the AdamW level’s second batch in §4.4, and it is a cell of the G family in §5.4. The properly composed replacement, R4 (sm4), has since run and **refuted the mechanism** the corner was built to test (§5.5) — which is the useful sense in which this failure was recoverable, and it cost a full cycle.

Reported as such, m12 is still informative, in a way we would not have bought deliberately: it is **two nondeterministic reruns of one 3 v 3 measurement** in one batch. They read $+0.554 \pm 0.217$ and $+0.357 \pm 0.211$. A 0.197 pp spread between reruns of one measurement, at fixed seeds inside one batch, is a direct **nondeterminism** bound — not a seed bound and not a batch bound — and it is of the same order as $\tau = 0.203$ pp (§4.4). Over all twelve same-seed pairs the bound is mean $|\Delta_{\text{plateau5}}|$ **0.163 pp**, max **0.472 pp**. Collapsed to its three groups, m12 gives $D = +0.456 \pm 0.195$, which is row 5 of Table 4.

What we now do about it, mechanically. `analysis/argsline_guard.py` reads a run’s own ARGs: line, fails on any repeated flag, and checks a declared design; `bin/_lib_guards.sh` refuses to sbatch a command line that does not match the declared design and re-reads the first launched job’s actual ARGs:

line afterwards. Swept over the 2,237 runs carrying an `ARGS:` line on both clusters, exactly **36 carry a repeated flag, and they are exactly m12’s 24 and sm3’s 12**. No other batch on either cluster is affected: `--alg-base`, `--stepsize-groups`, `--meta-stepsize`, `--alpha0`, `--num-epochs`, `--seed` and `--normalizer-param-*` are single-occurrence in every run in the corpus, so **the partition axis, the budget axis and the step-size axis were never overwritten anywhere**. Ten further submission scripts had been flagged by a text search for carrying two or more `--alg-meta` occurrences; all ten are clean, verified two ways — no run they produced carries a repeated flag, and every extra occurrence in their text is inside a comment or inside a guard’s quoted assertion array. The text search that produced that list was a false-positive generator, which is itself the argument for STANDING RULE 20: audit the runs, not the scripts.

6.2 The metric column, and the registered-scorer rule

Two headlines were withdrawn in this project for the same two habits: quoting the 20-epoch plateau column instead of the 5-epoch primary, and scoring a batch with a reduction written after the data arrived instead of the scorer committed before the runs existed. On `aw1`, the two columns give $G = +0.226$ (t 3.23) and $G = +0.232$ (t 2.62); the difference is small but the *verdict* moved, because the registered rule was a 95% lower bound and the hand-rolled prose restated it as a point estimate with $t \geq 2$. On `hz3`, a 50-epoch trailing window turned a flat budget effect into a growing one (§4.8). Both are one-line mistakes that survived internal review.

The rule we now follow, and recommend: **if a batch has a registered scorer, it is scored by that scorer, unedited, and its output is quoted, not paraphrased**. A hand-rolled reduction may be a cross-check; disagreements resolve in favour of the committed code. The most expensive application of that rule in this paper is `gn1`: its scorer halts before computing the contrast a whole subsection had been written around, and the subsection was rewritten rather than the scorer overridden (§5.6, §7 T7).

The rule’s companion — **no batch is submitted without a scorer registered first** — was broken more often than the record said. Five batches behind material reported here have no pre-registered scorer: `m12`, `n11`, `r50`, `gm2` and `sm3`. Two of them supply live Table 4 rows and one, `n11`, supplies both single-batch levels of the base-optimiser moderator in §4.4. Their values reconcile field-by-field against their own runs’ `ARGS:` lines; the gap is procedural, and §8’s provenance table marks each one rather than letting §3.4’s scorer list read as coverage.

6.3 Batch and seed as sources of variance — and a claim of ours that did not reproduce

Every primary in this paper is within-batch, which cancels any offset shared by the two arms regardless of how large it is. That design choice is safe and we keep it. What we can no longer support is the *reason* we previously gave for it.

Re-derived at write time: taking every configuration key (network, dataset, batch size, granularity, base, meta, η , α_0 , γ , augmentation, β -box, hierarchical mode, λ , r , epoch budget) that appears in **two or more batches at $n \geq 3$ per batch**, with fixed-LR baselines excluded and non-converged regimes removed by a dataset-relative floor, gives **26 configurations over 237 runs**. Within-config, across-batch:

- $F(39, 172) = 0.71$ — batch is **not** resolvable as a variance component;
- random-effects $sd_{\text{batch}} = 0.000$ pp against a pooled within-arm **sd of 0.193 pp**;
- observed batch-mean spread per configuration: **median 0.100 pp, mean 0.126 pp, max 0.319 pp**.

The largest single replication series in the corpus is the **nodewise** cell at $\eta = 10^{-4}$, present in **eight** batches: 91.890 / 91.961 / 92.000 / 92.012 / 92.044 / 92.064 / 92.066 / 92.184 — a 0.294 pp spread over 31 runs. (The eighth is `rp1`’s **nodewise** arm, at 92.066, which sits inside the range the seven already spanned; it

is the only **rp1** number that appears anywhere in this paper and it changes nothing. Under §3.3’s duplicate rule the 31 runs are 28 distinct experiments, **m12** contributing three groups rather than six runs.)

A two-way batch $\times$ seed ANOVA on the 14 configuration cells carrying a complete rectangle gives **SEED** $F(30, 30) = 1.50$, $p = 0.138$ — seed is null, which agrees with the internal record. But the same reduction gives **BATCH** $F(14, 30) = 0.54$, $p = 0.89$, and no cell definition we could construct reproduced the record’s $F(62, 85) = 5.47$, $sd_{batch} \approx 0.21$ **pp**. Including non-converged runs sends batch F to 722, which is a statement about accuracy regimes, not batches. **We therefore withdraw the claim that batch is a large random effect on this cluster** and replace it with what we can measure:

At matched science — same network, dataset, granularity, base, meta, η , α_0 , γ , augmentation, clip box, budget — and $n \geq 3$ per batch, cross-batch offsets on this cluster are bounded by about 0.32 pp and are not separable from within-arm noise; seed is null (F 1.50, p 0.138). We keep every primary within-batch because the offset, while small on average, is not bounded *a priori* and is half the size of the effect at its observed maximum — and because at least one batch in this corpus is known to have run with unintended state (a modified optimiser file timestamped 33 s before submission), which no variance model would have caught.

One consequence of the withdrawal, spelled out because §4.4 used to depend on the opposite. If batch carried a resolvable variance component, then a second submission of a configuration would be a *replicate* even at the same seeds, and the agreement of several such submissions would be evidence about that component. With the component withdrawn, a submission at the same seeds is a **rerun**: its agreement with the first bounds run-to-run nondeterminism and nothing else. That is exactly the situation of §4.4’s four repeated cells, three of which share seeds 0–2 and one of which shares seed 3 with a fourth, and §4.4 now says so. It is also the reason the eight-batch **nodewise** series quoted above is a *tight* bound rather than a loose one: those batches share their seeds, so the 0.294 pp spread has the seed effect divided out of it by construction, and it is the same set of runs §4.4’s $Q = 0.88$ is computed on. The two readings are one measurement, and neither of them is a seed-level replication.

Hardware was excluded by direct measurement: within (configuration $\times$ family), 2080Ti – L4 = +0.035 ($n = 45$), A100 – L4 = –0.021 ($n = 21$), A100 – 2080Ti = +0.112 ($n = 5$), over 29 distinct nodes. That last figure is used in §7 T9 to show that the one hardware mismatch in the corpus points the *wrong* way to explain the outlier it sits on.

6.4 A guard with a blind spot

`analysis/argsline_guard.py` is the corpus’s only automated check against the defect of §6.1, and its file collector globs `<dir>/*.out` without descending, so it silently skips the twelve local `.out` files under `runs/failed_hier_v1/`. All twelve are clean and none is in the run table, so nothing in this paper is affected — but a guard with an unstated blind spot is worse than a guard with a stated limit, and we record it here rather than in a commit message. The corpus sweep quoted in §6.1 was run with a recursive collector for this reason.

7 Threats to validity

T1 — Almost one meta-optimiser. Lion’s sign update makes the per-group α the *only* thing setting per-coordinate update magnitude, which is precisely the regime where the partition should matter most. Our headline is measured at the most favourable point of the axis we barely varied, and this remains the paper’s largest hole. **sm4** has narrowed it, and it is worth being exact about by how much.

The census is stated under a rule a reader can execute against the deposited run table: a **partition-family run** is an admissible row whose **granularity** is **nodewise**, **nodewise1d**, **chunk*** or **permnode***. There are **420** of them, of which **408 are meta = Lion and 12 are meta = RMSProp**; the twelve are **sm4**. Before **sm4** the count was 408 of 408. (An earlier version of this paper quoted “367 of 367” here. That figure

does not re-derive under any definition of “partition-programme run” we can reconstruct from the run table, so we drop it rather than restate it — the same treatment Appendix A.7 gives every number of that kind.)

So the axis is no longer empty, and it is not an axis either:

- **Twelve runs, one cell, under 3% of the partition programme.** One base optimiser (AdamW), one network, one dataset, one budget, one β -box, one meta-stepsize, three seeds.
- **No ladder.** The `sm4` scorer attaches the limit as part of its verdict: the batch runs at `aw1`’s meta-stepsize of 10^{-4} , **not at an RMSProp-meta optimum**, because no meta-stepsize ladder exists under an RMSProp meta — this corpus contains no other RMSProp-meta run of any kind. §4.1 shows the meta-stepsize moves the granularity gap by more than the gap, so an untuned η is not a minor caveat on this axis specifically.
- **The meta axis now has exactly the defect §4.4 spent a batch removing from the base axis.** Of its two levels, one has fourteen cells and seven-plus submissions and the other has one cell and one submission. A single-cell level cannot be tested for within-level homogeneity, so nothing in this corpus can distinguish a meta-optimiser effect from a batch effect peculiar to `sm4` — the same argument that made `n11`’s two levels inadequate before `bm2` ran.

What the one cell does buy. It is enough to refute a claim that was stated over the whole grid: the corner where base *and* meta both carry a second-moment normaliser produces the **largest** D in the AdamW family, not the smallest, against a bar registered before the runs existed (§5.5, M9). And it produces a structural result the Lion-only corpus could not have produced: at a fixed AdamW base, switching the meta-optimiser leaves G untouched (-0.001 ± 0.104) and moves $D - G$ by $+0.690 \pm 0.257$, so the tail story is scoped to a base-meta **pairing** rather than to a base (§5.4). Refuting a universal claim takes one counter-example. Establishing the axis takes a ladder, and there is none.

The experiment this hole still needs, stated so that it is not confused with what was run: a meta-stepsize ladder under an RMSProp meta at a fixed base, plus at least one non-Lion meta cell at a base other than AdamW — SGDm is the obvious one, since it is where seven of the paper’s submissions sit and where the tail story is strongest. Neither exists, and neither is in flight.

T2 — The mechanism is missing, one candidate is not identifiable from this design, and the alignment null is underpowered. §5.7 and §4.6. The size distribution is the surviving carrier and takes $+0.590 \pm 0.107$ of the decomposition, but we cannot say which property of it and this corpus provably cannot tell us. Separately, the leg that redirected us there is measured once, at $n = 3 \vee 3$, with an interval spanning $[-55\%, +51\%]$ of D and with permutation variance confounded with seed variance: **a moderate alignment contribution is not excluded**, and the batch that would exclude it (R1, `rp1`, 24 jobs) is running rather than reported.

T3 — Scale and domain. CIFAR-resolution vision, ResNets, one framework. ImageNet-1k (Deng et al., 2009) is not available to us: the copy on our allocation has 489 of 1000 training classes and an unlabelled validation set. No transformer, no language model, no reinforcement learning — and the parent method’s own home is reinforcement learning.

T4 — Competitiveness. We state the grouping key, because the answer depends on it. Among the cells in this corpus that run **plain** MetaOptimize (no hierarchical pooling), the highest is 93.317 ± 0.083 ($n = 3$; ResNet-18, six blocks, AdamW base + Adam meta, $\eta = 10^{-3}$, $\alpha_0 = 3 \times 10^{-4}$, the interior maximum of a 7-rung α_0 ladder scored on the test set). Tuned SGD + cosine (Loshchilov and Hutter, 2017) on the same network and budget is 95.124 ± 0.047 ($n = 5$; lr 0.1, the interior maximum of a bracketed 4-point grid: $94.172 / 94.844 / \mathbf{95.124} / 94.181$). The deficit is -1.807 ± 0.095 pp.

That is the *smallest* deficit in the corpus, not a typical one, and the network is doing the work. On ResNet-34 the best plain cell is 92.265 ± 0.053 ($n = 9$) against 94.823 ± 0.035 for AdamW + cosine, a deficit

of -2.558 ± 0.063 ; on ResNet-50 it is 90.833 ± 0.236 against 95.047 ± 0.110 , a deficit of -4.214 ± 0.260 . **The honest range is -1.8 to -4.2 pp behind a tuned schedule, and it widens with depth.**

The single highest MetaOptimize cell anywhere in the corpus is not the one above: it is 93.967 ± 0.074 ($n = 3$; ResNet-34, layerwise, SGDm + Lion, $\eta = 10^{-3}$, $\alpha_0 = 10^{-6}$, with this project’s own additive hierarchical pooling at η -ratio 0.025), -0.857 ± 0.081 behind the ResNet-34 baseline. We do not lead with it: it is the last rung of a four-rung η -ratio ladder within its own batch, so its argmax is at the ladder boundary and is interior only when a second batch’s rungs are appended, and §5.9 reports that this pooling operator does not generalise. It is stated here so that no reader can find a higher cell in the released table than the paper admits to.

MetaOptimize does beat a constant-LR AdamW tuned over four rungs (91.849 ± 0.092), which is the comparison the parent paper makes. We make no competitiveness claim beyond that, and all six numbers in this item are test-set argmaxes with no held-out estimate behind them (T12).

This threat is not answered by the size of D , and we do not attempt to answer it that way. The practical-significance argument for the partition result is given in §4.7 and it is designer-facing: the choice of partition is free of learned parameters and free of memory, it is the one lever in this family that no method measures, and the deficit reported here is the reason we make no recommendation to adopt the method.

T5 — Rule 11 is closed on one cell only. §4.5 closes the “you compared at one arm’s favourable η ” objection on ResNet-18/CIFAR-10. No meta-stepsize ladder exists on ResNet-34, ResNet-50, CIFAR-100 or under AdamW.

T6 — The 100-epoch horizon, and non-monotonicity within it. D goes significantly negative at epoch 55 on two independent batches and recovers by epoch 100 (§4.8). Our claims are about the end of training at 100 and 300 epochs.

T7 — BatchNorm and “size-1 tail” are under-identified. On every network we ran, the only 1-D tensors are normalisation parameters and one bias, so “the groups are degenerate” and “the groups are on the normalisation parameters” coincide exactly. **gn1** was the designated separator and **issued no verdict**. Its pre-registered commensurability gate T0.6 tests the **difference in level** between the two halves against a bar of 2.0 pp: the BatchNorm arms sit at 92.293 (error budget 7.707 pp) and the GroupNorm arms at 89.431 (budget 10.569 pp), a difference of -2.862 pp, which exceeds the bar by 0.862 pp. (The scorer also prints the budget *ratio*, $1.37\times$, next to that line; the ratio is descriptive and is not the gated quantity.) The gate fires before either contrast is computed and the scorer stops there.

Two invocations of that scorer must be distinguished, because §8’s deposit does not carry the input the second one needs. Run with `--root` pointed at the batch’s own run directory — probes present, the invocation whose output we quote — `c84_gn1_score.py` passes T0.5 with `rec_lo = rec_hi = 0.0000` on all four arms and then fails T0.6 as above. Run against the deposited tree, where `probe*.jsonl` is excluded, it never reaches T0.6: T0.5 reports `0 probe dirs ... NO PROBE -- cannot be gated` on all four arms, `VOID_ARMS = 2` is met, and the scorer exits at `T0 VERDICT: 4 of 4 arms void or dropped -> THE WHOLE BATCH IS VOID`. Both halts refuse a GroupNorm contrast, but they are different refusals for different reasons, and a reader reproducing from the deposit will see the second. §8 says which scorers this affects. **We therefore report its GroupNorm arm means (Appendix B) and no GroupNorm contrast: no D from those arms enters Table 4, the heterogeneity pool of §4.4, the design-point set of §5.8, or any count in the abstract.** Consequently every cell in this paper uses BatchNorm, and this limitation is therefore *un-probed* rather than probed-and-null; throughout we write “*normalisation scalars or, more generally, one-dimensional tensors*” rather than committing. One consequence must be stated plainly: because the scorer halts before T1, the retained **gn1** BatchNorm cell (Table 4 row 4, $+0.587 \pm 0.153$) is a **re-derivation from the four arm means the scorer does print**, not a scorer verdict. Its T0 gate passes on 4/4 seeds per arm and T0.6 gates the percentage-point comparison of the two halves rather than the BatchNorm contrast itself; had T1 been reached, $+0.587$ at t 3.83 would have cleared its registered bar

of $+0.30$ and $t \geq 2$. A maximally conservative reader may drop that row too: the same-contrast pool then loses one of its fourteen cells, and on the eleven cells the previous draft pooled it moved the estimate from $+0.571 \pm 0.037$ to $+0.570 \pm 0.038$ with Q 36.39 on 9 df — a move of 0.001 pp, which is why we keep it and label it. The design that would separate the normaliser question — a level contrast at fixed network, fixed base and commensurable error budget — is `gn2a/gn2b`, specified in advance in `bin/c84_normaliser_transfer.sh` and not run.

T8 — $n = 3$ in fifteen of twenty cells. Seed is statistically null on this cluster (§6.3), and a k -of- k per-seed agreement at $n = 3$ has exact $p = 0.25$ and carries no evidence, so we report no sign tests. But fifteen cells rest on $3 \vee 3$, and §6.1 shows two reruns of one such cell differing by 0.197 pp.

T9 — Excluded data, and one batch that is two. One batch (`ar1`, $D = +0.697 \pm 0.118$) is excluded as box-void throughout: it bound on the step-size guards asymmetrically, in the direction that inflates D .

`hz3` carries a second, different defect, and it is a **configuration** defect, not a metadata one. The batch is two submissions: twenty-one runs at job ids 4782007–4782027 in the clip box $-30:9.0$, and three seed-5 runs — `chunk777`, `nodewise1d`, `chunk2325` — resubmitted 32,288 job ids later at 4814293–4814295 in the narrower box $-15:-2.3026$. Their ENV lines say so and `results/all_runs.csv` records it faithfully; nothing here is repairable by fixing a column. `BETA_CLIP` is the only effective difference: every `ARGS` flag matches the corresponding seed-4 run exactly.

Two things follow, and they are not the same thing.

First, the box. Sweeping all 30,000 probe records of all 24 runs, no coordinate of any run reaches either rail before **step 81,000 = epoch 162**; from there the three narrow-box runs sit on their floor to the end (peak 138/14,421, 100/4,851 and 46/4,851 coordinates), while all twenty-one wide-box runs stay at 0/N throughout and the ceiling is never touched by anything. So the box difference is **inert at the 100-epoch reading and live at the 200- and 300-epoch readings**, one-sidedly, on the arm D favours. It is small — 0.95% of coordinates at its peak, against the registered 5% occupancy gate, which the batch’s scorer duly passes — but it is real and it is asymmetric.

Second, the hardware. The resubmitted trio also ran on an A100 80 GB while their seed-5 `nodewise` partner ran, with the rest of the original submission, on an RTX 2080 Ti. The batch was designed to assign GPU class **by seed** precisely so that a seed’s four arms share it; the resubmission broke that, and the batch’s registered scorer still labels all four seed-5 runs `gpu-2080ti-11g`. This mismatch applies at **every** budget, the 100-epoch reading included. It is a defect in the registered scorer’s metadata, not in its statistics; under our own file-freeze rule we do not edit a registered scorer after its data exist, so the correction is recorded here, and R2 makes the label true again.

Consequently §4.8’s within-run pairing cancels seed, run and batch, but **not** the clip box (from epoch 162) and **not** the GPU class (throughout). Both exceptions are limits on the budget reading and must be carried with it; neither is cancelled by the pairing, and §4.8 is to be read with them attached. The same applies to $T = \text{nodewise1d} - \text{nodewise}$; G and U are matched inside their pairs at every seed and are unaffected.

Seed 5 is also the batch’s per-seed outlier at 100 epochs — $D(100) = +0.148$ against **+0.548, +0.564, +0.586, +0.622, +0.990** — and unremarkable at 300 ($+0.290$, against a seed-2 low of $+0.168$). We note explicitly that **neither defect explains that outlier**: the box is inert at 100 epochs, and the hardware term (§6.3: A100 – 2080 Ti = $+0.112$ pp) sits on the arm that would *inflate* D , not deflate it. The 100-epoch outlier is unexplained. §4.8 reports the budget contrast both with and without seed 5 for this reason, and a replacement seed-5 trio in the original box and on matched hardware is R2 (§3.5).

Finally, for reuse: a scorer that groups on `beta_clip` drops exactly those three runs, keeps `nodewise` seed 5, and reads an **unbalanced** $D(300) = +0.464 \pm 0.084$ in place of the balanced $+0.428 \pm 0.086$. That is an artefact of the grouping, not a second estimate.

Neither `ar1`’s exclusion nor `hz3`’s split changes a **sign**. The `hz3` split does move one verdict’s **margin**, which is why §4.8 carries it inline rather than leaving it here.

T10 — Filter sensitivity. Several quantities in this corpus move by 0.5–0.7 pp between two defensible row filters. Every number here is re-derived at write time under the single stated gate (Eq. 11), and we recommend the same discipline to anyone reusing the data.

T11 — Novelty scoping. We claim no absolute priority. CAM-HD (Jie et al., 2020) built the granularity ladder, named the small-sample mechanism, and reported an interior optimum in 2020–2022; Zheng and Kwok (2019) studied blockwise adaptivity in 2019; Choi et al. (2019) established tuning-protocol sensitivity in 2019. What we add is count matching, the alignment null, the 1-D-tensor object, and the negative results in §5.

T12 — Selection on the test set; no validation split. No held-out validation split exists anywhere in this project. `plateau5` is a test-set quantity, the operating point (η , α_0 , the β -box) was chosen on it, and all the absolute numbers in T4 are argmaxes of ladders scored on it; 418 distinct configuration cells were scored on that same test set over the campaign (§3.3). Every absolute accuracy in this paper is therefore an optimistic estimate and none is offered as a benchmark result. The within-batch contrasts are much less exposed — no arm and no cell was selected, and §4.3 shows the reported set of count-matched contrasts is complete — but they are measured at a test-selected operating point, and D moves -0.189 pp per decade of η (§4.5). We cannot report a selection-free estimate of anything and we do not claim one.

8 Reproducibility

One command. The deposit re-derives and asserts the numbers that carry a claim in this paper; §3.4 states that scope exactly and `python3 code/c98_reproduce.py --census` measures it:

```
make reproduce      # every headline, re-derived and checked against the paper
make reproduce-table2 # Table 2 and Figure 1 alone
make figures        # regenerate all four figures from the CSV
make verify         # md5 every file against the manifest
```

(The target name `reproduce-table2` is historical: the primary contrast table is Table 4 in this manuscript and Figure 1 is the forest plot.)

`make reproduce` runs in seconds on a laptop, needs `python3` and `matplotlib` and nothing else — no GPU, no PyTorch, no cluster — and **exits non-zero if any headline fails to reproduce**. It prints one line per number: *derived value | paper value | PASS/FAIL | where it appears in this paper*. Its output at the time of writing is included in the deposit as `REPRODUCTION-AUDIT.txt`.

What the deposit does not let you re-run, stated as a table rather than as a caveat. Ten registered scorers are quoted or relied on in this paper. Run unedited against the deposit alone, two reach their quoted verdict in full, two reach part of it, and six halt or print `NO DATA`, because their box-occupancy and meta-gradient-field gates read the excluded `probe*.jsonl`. `python3 code/c98_reproduce.py --deposit` prints which is which and why, and `python3 code/c77_pp1_score.py` on the deposit tree prints `0 probe dirs` and `P2 ... cannot be scored` — we would rather a referee find that in this paragraph than at a terminal. The two verdicts §3.5 adds, `bm2`’s R0.5 box gate and `sm4`’s box gate, are in the same class and for the same reason. The rule we adopt going forward, and recommend: **a scorer registered from now on ships a --summary input path at registration time**, so that its deposit-reproducibility is part of what was committed before the runs existed. Retrofitting one onto the scorers already quoted here would re-mint every md5 in the provenance table of this section and destroy the very property — committed-before-the-data, run unedited — that makes those quotes worth anything.

Artefact and identifier. The deposit is **6.2 MB in 137 files**, with `MANIFEST.md5` covering every one of them and `make verify` checking all 137 against it. **It has no DOI, and this paper prints none.** The artefact is identified by the repository commit stamped at the top of the deposit’s `README.md`, which

MANIFEST.md5 pins byte-for-byte; a DOI is attached when the archive of record issues one, and the deposit carries the four-step procedure for doing that and for writing the resulting string into the three places that must agree. We print the commit rather than a promised identifier because a promised identifier does not resolve.

Data. `data/all_runs.csv`, 2,173 rows, one per run, with the full configuration (network, dataset, batch size, granularity, base, meta, η , α_0 , γ , augmentation, β -box, hierarchical mode, λ , r , seed), the outcome columns (`best_test`, `final_test`, `plateau5`, `plateau`, `auc`, `epochs-to-threshold`), the provenance columns (`job_id`, `account`, `node`, `wallclock_min`), the two admissibility flags (`window_ok`, `complete`), and a `dup_group` column carrying twenty-one groups — the eighteen pairs of differently-named runs that resolve to the same experiment plus the three `a0` reruns, where one member of each carries `superseded = 1` (§3.3). The column is exhaustive within a batch for every run that logs an `ENV:` line; it does not carry the nine cross-submission pairs of §4.1’s parent cell, which §4.1 handles in text. Raw per-epoch series are the Slurm `.out` files in `logs/raw_out.tar.gz`; each carries its own `ARGS:` and `ENV:` line, which is the authority on what that run actually did. The shipped log set is **2,241 files, which is every .out file on either cluster:** 2,237 with an `ARGS:` line and the four infrastructure jobs that have none. There is no shortfall to state, and the `sm3` runs that an earlier version of this ledger carried as un-ingested are in the run table.

Attrition — every job we launched, and where each one went. Table 9 reconciles the Slurm output directories on both clusters against the run table, so that a reader can confirm nothing was dropped silently. It is exhaustive: the two categories of loss are “crashed before epoch 1” and “budget too short for a plateau window”, and neither touches a count-matched arm.

Table 9: Attrition ledger. Job counts; GPU-hours where a wallclock was recorded. The exclusions are applied in the order shown, which matters: 425 rows fail `window_ok`, but 25 of them have no readable `plateau5` and are removed first, so the `window_ok` line reads 400.

stage	<i>n</i>	GPU-h	note
Slurm <code>.out</code> files on the two clusters	2,241	—	<code>runs/</code> + <code>runs_alice2/</code> + the on-cluster copies
– infrastructure jobs, no training	–4	—	<code>gtest</code> , <code>gtest2</code> , <code>mo-smoke</code> , <code>ts-pretok</code> ; no <code>ARGS</code> line, no CSV row
jobs that entered the training script	2,237	—	each logs one <code>ARGS</code> line
– crashed or cancelled before epoch 1	–64	≈0	itemised below; not one logged a single epoch
rows in <code>results/all_runs.csv</code>	2,173	1,624.8	2,150 carry a wallclock
– no readable <code>plateau5</code>	–25		2–5-epoch smoke tests
– <code>window_ok = 0</code> , <code>plateau5</code> present	–400	66.3	budget ≤ 20 epochs; <code>window_ok</code> is <code>epochs_done > 20</code>
– <code>window_ok = 1</code> , <code>complete = 0</code>	–24		truncated runs, 7 of them in-flight <code>rp1</code> snapshots; <code>complete</code> is <code>epochs_done ≥ 0.95 × requested</code>
admissible	1,724	1,558.4	the gate of Eq. 11

The 64 that never produced an epoch, by cause, read off their own tracebacks:

<i>n</i>	cause	batches
36	LAM='na' reached <code>float()</code> in the patched <code>HF.py</code>	am4, amx
12	BETA_CLIP without the colon separator reached <code>str.split(':')</code>	hs, ha
6	a block partition applied to the wrong network depth	sc
6	CUDA device-side assert (CIFAR-100 label range)	uc5
4	Slurm CANCELLED	a0, g1

All 64 died in `build_optimizer` or in the first minibatch, all are configuration errors that a later pre-submission guard now catches, and none consumed measurable GPU time. **None is in a count-matched arm**, so no contrast in Table 4, §4.6 or §5.4 is affected. Two of the five causes do touch material reported elsewhere and we name the contact rather than leave it to be found: the four Slurm cancellations are all `scalar-arm` runs (`g1-adamw-scal-s{0,1,2}` at $\alpha_0 = 10^{-6}$, `a0-scal-1e3-s0` at $\alpha_0 = 10^{-3}$), so the `scalar` column of §4.1’s α_0 table is short by four runs relative to what was submitted — which is one of the reasons that table is labelled descriptive and cross-batch there, and why the in-batch η pair above it carries the claim; and the 48 `am4/amx/hs/ha` crashes are all hierarchical-pooling runs, in the region of §2.4 and §5.9 that Appendix A.9 already marks as carried from the project record rather than re-derived.

The 449 inadmissible rows are not failures; they are short-budget probes. All 425 `window_ok = 0` rows requested ≤ 20 epochs (392 at 20, 33 at 2–5) and so could never carry a five-epoch plateau at their requested budget. Only 24 of the 449 are full-budget runs, and those are the truncated ones. By granularity, the 449 are: `layerwise` 136, `weightwise` 111, `nodewise` 108, `resnet18_blocks` 56, `scalar` 30, `permnode` 7 (the in-flight `rp1` snapshots), no-partition baseline 1 — i.e. attrition is concentrated in the coarse and the maximally fine arms of the exploratory ladders, which is where the 20-epoch probes were run.

Attrition inside the primary contrasts is exactly zero. The sixteen batches that carry a count-matched contrast contribute **272 runs, of which 272 are admissible**. More strongly: **every** uniform-chunk, `nodewise1d` and `permnode` run in the corpus outside the in-flight `rp1` batch — 238 rows — is admissible, so no count-matched cell could have been lost to the gate even in principle. (Counting `rp1` the figure is 249 of 256; the seven exceptions are its mid-flight snapshots, and they enter no contrast here.) Submitted n equals admissible n in all twenty cells of Table 4 and in the excluded `ar1` cell.

One registration deviation, disclosed here as well as in §5.4. `aw1`’s registered scorer (`analysis/c88_scorers.py`, committed before the runs) specifies “12 jobs: `nodewise` and `chunk777` $\times$ 6 seeds”. The submission script `bin/c90_awbase.sh` ran 4 arms $\times$ 3 seeds = 12 jobs instead: the same job count, four arms instead of two, and half the registered per-arm power on the primary. Every other batch’s realised seed set matches its submission script’s `SEEDS` line.

Compute. 2,150 runs carry a wallclock; they total **1,625 GPU-hours** over 29 distinct nodes and two accounts, on NVIDIA L4 24 GB, RTX 2080 Ti 11 GB, A100 80 GB and A100 MIG 40 GB partitions. GPU class is recorded per run and was measured to carry no systematic offset between arms (§6.3); every primary is within batch, so a per-node offset shared by the two arms cancels identically. The one exception, where a resubmission broke arm-level hardware matching inside a single seed, is §7 T9.

Environment. One virtual environment, both accounts, verified identical at write time: Python 3.10.4 (Python/3.10.4-GCCcore-11.3.0); PyTorch **2.0.1+cu118**; CUDA (torch build) **11.8**; torchvision 0.15.2+cu118; numpy 1.26.4 (torch 2.0.1 is incompatible with numpy ≥ 2); scheduler Slurm.

Seeds. Every run records its seed; seeds are 0–11. Batches specify their seed set in the submission script and the scorers assert the registered set is present before scoring. Table 10 gives, per batch, the seeds actually realised — which is not always the registered set, and where it is not, we say so.

Three things Table 10 is meant to stop a reader from having to guess. (i) **The β -box is not constant across cells**, so `r13` (–30:9.0) and `fa1` (–25:–2.3026) are not box-matched to the rest, and `hz3` contains one seed whose two arms sit in *different* boxes (Figure 3, §4.8, §7 T9). Nothing is pooled across boxes without the between-box Q partition of §4.4 being printed alongside. (ii) **m12’s three seeds were each run twice** under two names; the run table’s `dup_group` column records the eighteen such within-batch pairs, and any n , se or t computed over rows sharing a `dup_group` drops superseded rows, averages within the group, and counts n as the number of distinct groups. That is what makes `m12` a 3 v 3 cell with se 0.195, not a 6 v 6 cell with se 0.142 (§6.1). The reproduction audit enforces the rule, so a reader cannot accidentally recover the wrong number for `m12`. The column is **not** a census of every same-seed re-measurement in the corpus: across batches there are 154 such groups over 364 runs, none of which enters the n of a within-batch primary, and nine of which — §4.1’s `pp_/PP_` pairs — sit inside one reported cell and are handled in §4.1’s own text.

Table 10: Per-batch provenance. Every batch behind a cell in Table 4, with the script that submitted it, the scorer registered for it, that scorer’s md5, the seeds realised, the β -box, and the account. **Five batches have no scorer that was registered before their runs existed;** we mark them rather than let the reader infer coverage from §3.4.

batch	submission script	registered scorer	md5 (12)	seeds realised	n β -box	acct
cc1	c81_concordance.sh	c81_cc1_score.py	c138eab327c1	3,4,5	12 −15:−2.3026	A
mm1	c76_matched_m_partition.sh	c76_mm1_score.py	57e48c32e066	0,1,2	6 −15:−2.3026	A
pp1	c77_permuted_partition.sh	c77_pp1_score.py	3c6aa6f97627	0,1,2	9 −15:−2.3026	A
bn1 (<i>G</i> only)	c78_degenerate_tail.sh	c78_bn1_score.py	285632b3b8a6	0,1,2	9 −15:−2.3026	A
gn1	c84_normaliser_transfer.sh	c84_gn1_score.py	82c515d1ad49	0–7	24 −15:−2.3026	B
m12	c94_meta_ladder.sh	none registered	—	0,1,2 (each run twice)	24 −15:−2.3026	B
rl3	c87_rule11_ladder.sh	c87_rl3_score.py	b008216753a0	0,1,2	24 −30:9.0	B
fa1	c82_field_wideclip.sh	c82_fa1_score.py	b96ab2080cfa	0–5	24 −25:−2.3026	A
hz3	c87_horizon_300ep.sh	c87_hz3_score.py	16769a161630	0–5	24 −30:9.0 and −15:−2.3026	A
aw1	c90_awbase.sh	c88_scorers.py	0363bccb4d3	0,1,2	12 −15:−2.3026	B
nl1	c92_norm_ladder.sh	none registered	—	0,1,2	24 −15:−2.3026	A
g3m	c83_gen_r34_merged.sh	c83_gen_score.py	ffdabe0a1790	0–8	36 −15:−2.3026	A
r50	c93_resnet50.sh	none registered	—	0,1,2	12 −15:−2.3026	A
gc1	c83_gen_c100_screen.sh	c83_gc1_score.py	5e775f417031	0,1,2,3	8 −15:−2.3026	A
gm2	c91_c100_mech.sh	none registered	—	0,1,2	12 −15:−2.3026	B
sm3	c96_secondmoment.sh	none registered	—	0,1,2	12 −15:−2.3026	A
bm2	c97_bm2_basemod.sh	c97_bm2_score.py	—	3,4,5	12 −15:−2.3026	A
sm4	c97_sm4_secondmoment.sh	c97_sm4_score.py	—	0,1,2	12 −15:−2.3026	A
ar1 (excluded)	c79_argmax_robustness.sh	c79_ar1_score.py	2396e5cf2372	0,1,2	12 −15:−2.3026	A

(iii) **Five batches have no pre-registered scorer** — m12, nl1, r50, gm2 and sm3. nl1 alone supplies two of Table 4’s rows and both of the single-batch levels of the base-optimiser moderator in §4.4. Their CSV values were checked field-by-field against their own runs’ **ARGS**: lines and are correct; what is missing is not the data but the commitment-before-the-fact, and we mark it rather than claim coverage we do not have. The replication that repairs the most exposed of them, R3 (§3.5, bm2), has run: it was registered before any of its runs could exist and it supplies a second, independent batch at both of nl1’s levels (§4.4). nl1 still supplies the other half of each.

Code. The optimiser is the released MetaOptimize HF.py plus the patches in patches/ (patch_chunkwise.py, patch_nodebn.py, patch_permnode.py, patch_zpool.py, the probe patches, and the augmentation patch), each of which carries an identity test against the authors’ own working blockwise path. Analysis is 60+ scorers under analysis/; the ones that gate a verdict are named in §3.4 and quoted verbatim. The figures are analysis/c98_figures.py and contain no typed value.

Registration. Batches carry a submission script under bin/ with numbered pre-submission guards: the run table must be present; the required patches must be present in the live tree; group counts are **measured** on the instantiated optimiser and asserted equal to the registered values; there must be no name collision; the comparator batch must exist; and there must be disk and queue headroom. Since the failures of §6.1 the guards also enforce STANDING RULES 20 and 21 mechanically: the scorer must already exist and pass its own selftest before a job is composed, the composed command line is rejected if any flag repeats or if it does not match the declared design, and the first launched job’s actual **ARGS**: line is re-read after sbatch. Decision rules and bands are committed before submission. analysis/c88_scorers.py --selftest runs 31 assertions in both directions; it currently reports one failure, S6e, which asserts that the count-matched partition rows are all one base optimiser — an assertion made true at cycle 88 and made false, deliberately, by the aw1 and nl1 batches.

An integrity check a reader can run on our own logs. analysis/argsline_guard.py sweeps every run’s own **ARGS**: line for a repeated flag, because argparse silently takes the last occurrence and a submission script can therefore run a different experiment from the one it declares. Over the 2,237 runs carrying an **ARGS**: line on both clusters it finds exactly 36 with a repeated flag, in two batches, and those two batches are reported as void-as-designed in §6.1. The partition axis, the budget axis and the step-size axis (--alg-base,

--stepsize-groups, --meta-stepsize, --alpha0, --num-epochs, --seed, --normalizer-param-*) are single-occurrence in every run in the corpus. We recommend the check to anyone running batch experiments through a shell wrapper; it cost us two batches to learn, and §6.4 records the blind spot the tool itself still had when we found it.

Deviations from the parent’s configuration, stated so absolute numbers are not misread. (i) We enable RandomCrop + horizontal flip; the parent’s text and released code appear not to. Without augmentation ResNet-18 memorises CIFAR-10 within an epoch and there is no optimisation headroom for a step-size method to exploit, so we regard augmentation as scientifically necessary — but it means our absolute accuracies are not comparable to the parent’s. (ii) Our headline partition programme uses a Lion meta-optimiser where the parent’s SGDm arm uses Adam. (iii) We clip β to a box; the parent mentions no clipping. Box occupancy is measured per run, **per coordinate** from the `n_at_lo/n_at_hi` rails and never from the 62-element per-tensor summary — reading the summary instead reported 0.000000 occupancy on 24 runs that were clipped in every record and inverted an arm ranking, which voided the `ar1` batch (§3.3).

What a reader can falsify on their own hardware in an afternoon. Merge your one-dimensional tensors into one step-size group each and measure the difference at fixed group count; and re-tune both arms of any granularity comparison you have and see how much of the gap survives.

9 Conclusion

“The number of step sizes” is not one experimental variable. Within one batch, moving the shared meta-stepsize by a decade moves the layerwise-minus-scalar gap from +3.291 to +0.655 pp; the initial step size flips the sign of the ordering at the parent’s own configuration; and the group count, at fixed partition family, is worth at most +0.34 pp across 0.47 decades with an inconsistent sign.

What remains after all three are controlled is a real effect. At **matched group count**, replacing an architecture-aligned partition by a uniform one is worth a positive amount of accuracy in **every one of 20 within-batch cells** across three networks, two datasets, four base optimisers, two meta-optimisers, two meta-stepsizes and two budgets; it survives tuning each arm to its own optimum (-0.090 ± 0.178 , itself an upper bound in magnitude); it is present and resolved at $3\times$ the budget; and its variation across configurations has a candidate moderator that we can decompose but not identify — splitting the fourteen same-contrast ResNet-18 cells by base optimiser leaves them homogeneous inside every level (Q 7.36 / 10 df, p 0.69) and accounts for 92.8% of the between-cell heterogeneity (Q 95.12 / 3 df), with all four levels now resting on two or more independent submissions and the SGDm level homogeneous at Q 4.21 / 7 df, p 0.76, τ 0.000, pool $+0.556 \pm 0.045$. That share is a property of the grouping rather than of the label: batch identity accounts for 95.8% of the same Q , and the base optimiser survives conditioning on it only at ΔQ 4.11 on 2 df, p 0.13. A registered replication (`bm2`) gave the SGD and RMSProp levels a second independent batch each and both replicated, which moved that conditional test from ΔQ 0.05 on 1 df to ΔQ 4.11 on 2 df; it did not finish it. Architecture **alignment** is bounded out as the principal carrier, in one batch at one design point: holding the count *and* the per-tensor size multiset and permuting only membership is worth -0.009 ± 0.157 pp, 95% CI $[-0.317, +0.298] = [-55\%, +51\%]$ of D , against an MDE of 0.440 pp. A moderate alignment contribution is not excluded and a further batch is running to bound it more tightly. The surviving carrier is the group-size distribution, which takes $+0.590 \pm 0.107$ (t 5.53) of the same three-arm decomposition.

We cannot say which property of it, and the design cannot tell us. At fixed count this corpus contains one contrast type; twenty candidate statistics collapse to two equivalence classes and the best-fitting one is an arm indicator in disguise. Nine further mechanisms were examined and four are refuted, with the other five narrowed, unseparable or undecidable: $\sqrt{N}$ averaging (untested, not refuted, and its literature attribution withdrawn), the meta-gradient correlation field (anti-concordant at $t - 11.14$), size-1 groups as such ($+0.115 \pm 0.133$), the tail as a universal carrier (scoped to a base-meta pairing: pooled $D - G + 0.514 \pm 0.056$ under SGDm + Lion, -0.061 ± 0.085 under AdamW + Lion, $+0.629 \pm 0.242$ under AdamW + RMSProp), Choi-style inclusion (granularity is state, not hyperparameters), base-optimiser normalisation (RMSProp and

AdamW differ by $+0.531 \pm 0.138$), **second-moment normalisation anywhere in the loop — refuted on a bar registered before the data existed: the corner where base and meta both carry one gives $D = +0.889 \pm 0.228$, the largest in the AdamW family, against a predicted smallest —** and any size-distribution summary statistic (not identifiable). A ninth candidate, the aligned arm’s accuracy level, is neither confirmed nor separable: it is aliased with the network and the base optimiser, and the within-batch separator we registered for it issued no verdict (§5.6, §7 T7). And D is not predictable out of sample from any configuration property this design can resolve.

The honest description of this paper is therefore: **a robust, replicated, count-matched measurement in twenty cells; a candidate moderator for its heterogeneity, replicated at all four of its levels and homogeneous inside each, and still not separated from the batch identity it co-varies with; a single-batch bounded null that excludes the mechanism most people would guess as the principal carrier without excluding it as a contributor; a prescription that costs nothing in learned parameters and works everywhere except one base–meta pairing; nine dead or undecidable mechanisms, one of them killed on a bar we wrote down first; and no mechanism.** We would rather publish that than a mechanism that does not survive its own controls — this project produced one of those too, and killed it (§5.7) — or than a refutation that does not survive its own registered scorer, which this project also produced and withdrew (§5.6).

The experiments that would break the impasse, in the order we would run them. Two of the four registered in §3.5 have been read: R3, the base-moderator replication at fresh seeds, which replicated both levels and is in §4.4; and R4, the second-moment corner, which refuted its mechanism and is in §5.5. Two remain: the alignment replication with the permutation seed decoupled (R1, complete on disk and deliberately unscored) and the box- and hardware-matched budget trio (R2, queued). To those we add two experiments the last cycle created rather than closed. **The first would convert the base optimiser from a candidate moderator into an identified one:** a second independent batch at each level *within one submission*, or failing that a third batch at the level where base and batch are most tightly aliased, run at fresh seeds under the same registration discipline as **bm2**. Until something of that shape exists, the base optimiser and the submission it arrived in cannot be told apart at $p < 0.05$. **The second is a meta-stepsize ladder under a non-Lion meta-optimiser, and a non-Lion cell at a base other than AdamW:** **sm4** shows the meta-optimiser moves $D - G$ by more than the base does at one point in the space, and a single cell with no ladder under it cannot say whether that is a meta effect or a batch.

Beyond those, the decisive mechanism design is additive rather than subtractive. Every count-matched arm we have ever run *removes* the size-1 tail, so “tail” is perfectly confounded with “architecture-aligned”. The decisive design creates, for the first time, an arm that is **neither aligned nor tail-free**: arm A1 = plain **chunk2325** ($m = 4,851$, min size 10, size CV 0.088); arm A2 = **chunk2325** with one large convolution tensor split into 9,610 size-1 groups and merging elsewhere to hold $m = 4,851$ **exactly**. Two arbitrary partitions, identical count, differing only in whether a size-1 tail exists. If the tail is the carrier, A2 loses ≈ 0.6 pp to A1. If A2 – A1 is null, the degeneracy statistic is dead as a carrier and the size-1 tail is only a proxy. Pair it with a heterogeneity sweep at **zero** degeneracy — uniform-count partitions at $m = 4,851$, minimum size ≥ 10 , prescribed size CV in $\{0.09, 0.40, 0.76, 1.90\}$ — which dissociates CV, Gini, sd-log and the heterogeneity index from the singleton fraction for the first time. Run both under AdamW and SGDm, because that is precisely where the two candidate carriers disagree. Roughly 40 jobs. It is registered and unfunded.

Two further designs are named because this paper states limits it cannot close without them. For the level covariate, a normaliser or width contrast run at a **commensurable error budget** — **gn1** failed its own commensurability gate on a 2.862 pp difference in level against its 2.0 pp bar, and **gn2a/gn2b** were specified for exactly this and not run. And for external validity, the same four partitions applied as a fixed per-group learning-rate scale on a plain SGDm or AdamW run with **no meta-learning at all**, which is the only design in the queue that would convert an internal audit of one framework into a claim about step-size granularity.

A Discrepancy register

Numbers carried in this project’s internal record that **did not reproduce** at write time, and what replaced them.

A.1 — The meta-optimiser axis does not exist. The record carried “meta axis (ResNet-18/CIFAR-10, base = SGDm): Lion +0.727 · Adam +0.554 · RMSProp +0.357”. The arithmetic is right and the labels are wrong: the m12 batch’s two halves both ran meta = Lion (§6.1), so +0.554 and +0.357 are two nondeterministic reruns of one configuration at the same seeds, not two meta-optimisers. Every partition-family run behind the cells of Table 4 is meta = Lion, and the one batch that varies the meta-optimiser is a single cell: of the 420 admissible runs in the partition families, 408 are meta = Lion and 12 are meta = RMSProp. **One cell is not an axis, and no meta-optimiser axis may be reported.** (The record’s own count, “367 of 367”, does not re-derive under any reconstructible definition and is dropped rather than restated, per A.7’s precedent.) Consequently the record’s “ G is not null under an RMSProp meta (+0.276, t 3.21)” is also mislabelled: that is m12’s second Lion rerun, and its companion reads +0.071 ± 0.103.

A.2 — Two values from the banned metric column. The record carried $G(\text{AdamW}) = +0.226$, t 3.23 and hz3 “the gap grows with budget, +0.229, t 3.27”. Both come from 20-epoch and 50-epoch trailing windows respectively. On the primary 5-epoch window: $G(\text{AdamW}) = +0.232 \pm 0.089$, t 2.62; and the budget effect is -0.149 ± 0.105 , t -1.42, i.e. a decline that does not resolve — -0.207 ± 0.107 , t -1.94 once hz3’s mismatched seed-5 pair is set aside (§7 T9, §4.8). Both replacements match the project’s own later corrections; we record the originals so the supersession is visible. Note that the 50-epoch reading is not merely a discarded internal note: it is the **registered primary** of `analysis/c87_hz3_score.py`, whose verdict is GROWS. We keep `plateau5` as this paper’s primary metric and print both, rather than quietly preferring the window that suits the narrative.

A.3 — Batch as a large random effect: not reproduced. The record carried BATCH $F(62, 85) = 5.47$, $p = 6.9 \times 10^{-13}$, $sd_{batch} \approx 0.21$ pp. Under every cell definition we could construct, restricted to converged runs, batch is not resolvable ($F(39, 172) = 0.71$; random-effects $sd_{batch} = 0.000$ against within-arm sd 0.193). The seed half of the claim reproduces ($F(30, 30) = 1.50$, $p = 0.138$). See §6.3 for the replacement statement.

A.4 — Heterogeneity τ , and the two Q ’s. The record carried $\tau = 0.285$ pp against 0.137 pp measurement noise. DerSimonian–Laird with Welch standard errors gives $\tau = 0.215$ pp against an rms se of **0.151 pp** on the twelve-cell pool as first published, and $\tau = 0.203$ pp against **0.152 pp** on the eleven-cell pool after gn1-GroupNorm is removed (§4.4). The ratio is 1.3–1.4 $\times$, not 2 $\times$. **Two Q values are in circulation and both are correct.** Computed from full-precision arm means in `results/all_runs.csv` — the convention this paper adopts, because it is what the deposited code computes — $Q = 43.19$ on 11 df (36.40 on 10 dropping GroupNorm). Computed from the three-decimal (D , se) pairs *as printed in Table 4*, $Q = 43.01$ (36.29), with pool +0.545 (+0.570), τ 0.214 (0.203). The gap is rounding of the printed table and nothing else; an earlier appendix described it as “reproduces exactly”, which a 43.2-versus-43.0 mismatch does not support. The between-base share of §4.4 is 88.4% under both conventions. The conclusion is unchanged and is now stronger: the excess over measurement error is real. The between-base share of §4.4 was 88.4% under both conventions on that eleven-cell pool; on the fourteen cells §4.4 now uses it is 92.8% of a Q of 102.47. The conclusion is unchanged and is stronger: §4.4 decomposes almost all of it — without claiming to have identified its cause.

A.5 — The prediction null has weakened slightly with two new batches, and again with a removal. The record stated that inside CIFAR-10 the corpus mean *wins* out of sample (0.2791 vs 0.2863) with a sign test of 8/11, $p = 0.227$. With r50 and m12 added the headroom models were marginally ahead inside CIFAR-10 and the sign test was 9/11, $p = 0.065$. After the gn1-GroupNorm removal the design-point set is 10, the best model wins inside CIFAR-10 by 0.040 RMSE (0.2791 $\rightarrow$ 0.2395, -14%) and the sign tests

are 8/10 (p 0.109) overall and 7/9 (p 0.180) inside CIFAR-10. The verdict is unchanged — nothing reaches significance, the margin is dominated by one CIFAR-100 fold, the functional form is a best-of-ten selection, and the power bound $|r| \geq 0.632$ is not approached — but the null now holds by a narrower margin than the record implied, and we report the margin that re-derives rather than the record’s phrasing.

A.6 — fa1’s ceiling caveat, stated precisely. The record variously described **fa1**’s **nodewise** arm as grazing the ceiling in “5 of 6 seeds”. The scorer’s own occupancy table: all 24 arms read **rec_lo** 0.0000 exactly (the floor is free), and on the **nodewise** arm five of six seeds have non-zero **rec_hi**, of which **three** exceed the 5% gate (0.1673, 0.2193, 0.0601). Both statements are true of different thresholds; we quote both.

A.7 — Adaptation versus a frozen β . The record carried “+2.551 pp over frozen β ”. The frozen- β runs on disk are 20-epoch probes and fail this paper’s admissibility gate (**window_ok** = 0), so the claim is **not re-derivable** under our own stated rule and is dropped rather than restated.

A.8 — Corpus size. The record variously says 1,960 / 2,077 / 2,113 runs and $\sim$ 1,200–1,400 GPU-hours. At write time: **2,173 rows, 1,724 admissible, 1,625 GPU-hours** summed over the 2,150 runs carrying a wallclock, with 2,237 jobs having entered the training script and nothing awaiting ingest (§8, Table 9). The correction register runs to entry 128.

A.9 — Two blocks carried from the record without re-derivation. Two statements are **not** re-derived at write time and are marked where they appear: (i) the “6 of 6 fits with the wrong sign” count for the drift-versus-group-size instrument (§5.1) — the instrument’s own interpretation is disowned in the same paragraph, so nothing rests on the count; and (ii) the M0 **shrink** and M1 **additive r** pooling failures (§2.4, §5.9), whose replacement is the **zpool** sweep, which **was** re-derived on **plateau5** and carries that section’s argument on its own. No claim in §4 depends on either.

A.10 — Two headline conventions, superseded by measurement. Two phrases the record carried are withdrawn on evidence rather than on taste. “*Byte-identical contrast*”, used of the twelve-cell pool, is false three times over: the cells span three step-size clip boxes, **hz3** spans two boxes inside a single cell, and a registered scorer’s own header forbids pooling across boxes. The replacement is measurement, not silence: over the eight SGDm cells, which span all three boxes and are not confounded with the base axis, between-box $Q = 0.76$ on 2 df (p 0.68); §4.4 gives that number, and the three whole-pool readings it has taken at eleven, thirteen and fourteen cells, rather than only the one that flatters us. And “*a CSV metadata defect*”, used of **hz3**’s three seed-5 rows, is wrong in the direction that matters: the run table is faithful to those runs’ own ENV lines and the defect is in the experiment (§7 T9).

B The full arm table for the primary contrast

plateau5 arm means with sem, ResNet-18/CIFAR-10 unless stated, all within batch. The β -box column is carried so a reader can redo §4.4’s between-box partition of Q .

[¶] **hz3**’s seed-5 **chunk777**, **nodewise1d** and **chunk2325** runs sit in $-15:-2.3026$ instead (§7 T9).

[§] The **gn1** GroupNorm arms are reported here for completeness because the batch’s registration requires its four-arm table to be reported whole. **No contrast is formed from them anywhere in this paper:** the registered scorer’s commensurability gate fired on a 2.862 pp difference in level against its 2.0 pp bar, and issued no verdict (§7 T7). The two **gn1** lines above share one batch and 24 runs.

sm3’s, **sm4**’s and **bm2**’s arm means are re-derived from the deposited run table like every other row of this table. **ar1**’s arm means are omitted because the batch is box-void and quoting them invites the contrast we exclude; its D ($+0.697 \pm 0.118$) is quoted in §4.3 solely so the exclusion is visible.

Table 11: Arm means for every batch behind the primary contrast.

batch	β -box	nodewise	chunk777	nodewise1d	chunk2325
cc1	-15:-2.3026	91.890 $\pm$ 0.077	92.617 $\pm$ 0.185	92.706 $\pm$ 0.139	92.717 $\pm$ 0.046
mm1	-15:-2.3026	92.044 $\pm$ 0.058	92.529 $\pm$ 0.150	—	—
pp1	-15:-2.3026	92.012 $\pm$ 0.129	92.593 $\pm$ 0.058	—	—
pp1 permnode ($m = 14,420$)	-15:-2.3026	92.003 $\pm$ 0.090	—	—	—
bn1	-15:-2.3026	92.184 $\pm$ 0.033	—	92.611 $\pm$ 0.025	92.906 $\pm$ 0.041
gn1	-15:-2.3026	92.000	92.587	—	—
gn1 GroupNorm arms [§]	-15:-2.3026	89.330	89.532	—	—
ml2 (3 seeds, each run twice)	-15:-2.3026	92.064	92.519	92.683	92.856
r13 @10 ⁻⁴	-30:9.0	91.908	92.589	92.664	92.653
r13 @3 $\times$ 10 ⁻⁴	-30:9.0	92.507	93.098	92.898	93.115
fa1 @3 $\times$ 10 ⁻⁴	-25:-2.3026	92.327	92.957	92.976	92.975
hz3 @300 ep	-30:9.0 ¶	92.816	93.244	93.153	93.096
aw1 (AdamW)	-15:-2.3026	92.978	93.257	93.069	93.301
nl1 (SGD)	-15:-2.3026	91.156	92.191	91.848	92.373
nl1 (RMSProp)	-15:-2.3026	92.155	93.129	93.071	93.091
g3m (R34)	-15:-2.3026	91.336	92.002 (chunk835)	92.094	92.265 (chunk2500)
r50 (R50)	-15:-2.3026	89.631	90.513 (chunk295)	90.681	90.833 (chunk884)
gc1 (C100)	-15:-2.3026	70.311 $\pm$ 0.227	71.951 $\pm$ 0.092 (chunk771)	—	—
gm2 (C100)	-15:-2.3026	70.569	72.054 (chunk771)	71.932	72.000 (chunk2293)
ar1 (excluded, box-void)	-15:-2.3026	—	—	—	—
sm3	-15:-2.3026	93.103	93.244	93.019	93.315
sm4 (RMSProp meta)	-15:-2.3026	90.785	91.675	91.773	92.034
bm2 (SGD)	-15:-2.3026	91.203	92.181	—	—
bm2 (RMSProp)	-15:-2.3026	92.507	93.138	—	—

End matter

Data availability. The complete run table (`results/all_runs.csv`, 2,173 rows), the raw per-epoch Slurm logs (2,241 `.out` files, each carrying its own `ARGS:` and `ENV:` line), all submission scripts (`bin/`), all optimiser patches (`patches/`), all registered scorers (`analysis/`), the figure code and the reproduction audit are deposited as a single archive. **The deposit has no DOI**, because it has not been deposited; the artefact is identified by the repository commit recorded in its `README.md`, and `CITATION.cff` carries no `identifiers:` block rather than a stand-in for one. The archive is 6.2 MB, carries an md5 manifest for every file, and re-derives every number `make reproduce` checks — the list is in §3.4 — on a laptop in seconds, with no GPU and no dependency beyond `python3` and `matplotlib`. **One class of number requires data the deposit does not carry, and we state exactly which.** The per-group β trajectories (`probe*.jsonl`, ≈ 42 GB) are excluded for size and are available from the authors on request. They are not a supplement: they are the input to the box-occupancy and meta-gradient-field gates inside the registered scorers, so a reader with the deposit alone regenerates **two** of the ten scorer verdicts we quote in full (`c87_r13`, §4.5; `c87_hz3`, §4.8), **two** in part (`c81_cc1`’s D and G legs but not §5.2’s field block; `c83_gc1`’s $S1$ but not its arm-asymmetry guard), and **six** not at all — including §4.6’s $P2$ alignment block, §5.6 / §7 T7’s `gn1` halt, the box-occupancy evidence that **excludes** `ar1` in §4.3, and the two gates behind §3.5’s newly scored batches. `python3 analysis/c98_reproduce.py --deposit` prints that register and its counts. Everything else the deposit reproduces without the probes: `results/all_runs.csv` and the raw `.out` series carry every arm mean, every D , G , U , T and ρ , every pool and every Q in this paper.

Code availability. The optimiser is the released MetaOptimize `HF.py` plus the patches in `patches/`, distributed as patches rather than as a fork, each carrying an identity test against the authors’ own working `blockwise` path (`tests/`). Analysis code, figure code and the reproduction audit are released under MIT; the run table, logs and documentation under CC-BY-4.0; the patches carry the parent work’s license.

Ethics. No human or animal subjects and no personal data. CIFAR-10 and CIFAR-100 are standard public benchmarks used under their stated terms; no other data was collected. The work is a methodological audit of an optimisation method and we see no dual-use or deployment risk specific to it. The one ethical exposure we do carry is a conflict of interest, disclosed in full below rather than in a footnote, and the

mitigation for it was put in place before the results existed rather than after.

Competing interests. A supervising author of this work is a co-author of MetaOptimize (Sharifnassab et al., 2025), the method this paper audits. We state this plainly because several of this paper’s results are negative about that method: the 1.8-to-4.2 pp deficit against a tuned schedule (§7 T4), the finding that four fifths of the reported granularity benefit at the parent’s own default meta-stepsize is tolerance to an over-large meta-step rather than accuracy (§4.1), and the mechanism verdicts of §5. The mitigation is procedural and is auditable in the deposit: every decision rule and acceptance band for the negative results was committed to version control **before** the corresponding runs were submitted, in a scorer that is run unedited and whose printed verdict is quoted rather than paraphrased (§3.4, §6.2); the supervising author had no role in setting those rules; and where a batch has **no** such pre-registered scorer we say so explicitly in the provenance table in §8 rather than letting the reader assume coverage. Neither author holds any other interest in the audited method, and the authors declare no financial competing interests.

A second interest, of a different kind. The hierarchical partial-pooling design that §2.4 motivates and §5.9 evaluates is not ours. It was proposed to us by a researcher who is also a co-author of the parent work and who is not an author of this paper. §5.9 reports a **negative** result on that proposal, so we state its origin here rather than let a reader take both the design and the negative as ours. We regard origination of a design at that specificity as a substantial intellectual contribution rather than an acknowledgeable courtesy, and §5.9’s verdict was reached under the same pre-registered rules as every other verdict in §5.

Funding. This work was carried out as an MSc research project at LIACS, Leiden University. It received **no dedicated project funding and no grant**; the compute it consumed was drawn from the institutional allocation acknowledged below. The absence of a compute budget is a scope limit rather than a formality: it is the reason ImageNet-scale replication is out of reach (§7) and the reason the additive tail experiment proposed in §9 is registered but unfunded.

Acknowledgements. This work was performed using the compute resources from the Academic Leiden Interdisciplinary Cluster Environment (ALICE) provided by Leiden University. We thank the ALICE support team. One further debt is recorded under Competing interests rather than here, deliberately: origination of a design that this paper then evaluates is an intellectual contribution, and listing it as an acknowledgement would understate it.

Author contributions. Stated in CRediT terms. **M. Ahmaditeshnizi** (LIACS, Leiden University) — Conceptualization (equal), Methodology, Software (the chunkwise, 1-D-tensor, permuted-node and probe partitions and their identity tests, as patches to the released MetaOptimize implementation), Validation, Formal analysis, Investigation (all 2,173 runs), Data curation, Writing — original draft, Visualization, Project administration. **S. Salehkaleybar** (LIACS, Leiden University) — Conceptualization (equal), Supervision, Resources, Funding acquisition, Writing — review & editing, and continuity with the parent work. Explicitly **not** involved in setting the pre-registered decision rules or acceptance bands used for the negative results in §5. No one else contributed to this paper in a CRediT role. The one intellectual contribution that came from outside this list — the hierarchical partial-pooling design of §2.4 and §5.9 — is stated under Competing interests above, where its origin is also the second interest we have to disclose. Both authors accept accountability for the integrity of the work as a whole.

Use of AI assistance. Analysis scripts, batch submission scripts, the correction register, the figures and this manuscript were produced with substantial assistance from a large language model operating on the project repository under human direction, including at the analysis-design stage and not only at the writing stage. We disclose this as a first-class integrity item rather than a footnote, and we state the controls that make it checkable rather than asking to be trusted. Every number in this paper is re-derived from the run table or the raw logs at write time, and the numbers that carry a claim are additionally asserted against

what the paper prints by committed code (`analysis/c98_reproduce.py`, in the deposit; it exits non-zero if any of them fails to reproduce, and §3.4 states exactly what is and is not inside its scope); registered decision rules and bands were committed to version control before the corresponding runs were submitted; and the discrepancies this procedure surfaced between the project’s internal record and the data are recorded in Appendix A rather than corrected silently. Several claims were withdrawn by that procedure during writing — an alignment refutation downgraded to a bounded null, a mechanism refutation withdrawn because the batch’s own registered scorer refuses to compute the contrast it rested on, a budget flatness claim demoted to an unresolved trend, and a variance claim that did not reproduce — and two whole batches were found to have run a different experiment from the one they declared (§6.1). All of those outcomes are reported.

Correspondence. mohammadrezaahmaditeshnizi@gmail.com (M. Ahmaditeshnizi).

References

- Atılım Güneş Baydin, Robert Cornish, David Martínez Rubio, Mark Schmidt, and Frank Wood. Online learning rate adaptation with hypergradient descent. In *International Conference on Learning Representations (ICLR)*, 2018. arXiv:1703.04782.
- Xiangning Chen, Chen Liang, Da Huang, Esteban Real, Kaiyuan Wang, Yao Liu, Hieu Pham, Xuanyi Dong, Thang Luong, Cho-Jui Hsieh, Yifeng Lu, and Quoc V. Le. Symbolic discovery of optimization algorithms. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. the discovered optimiser is named Lion.
- Dami Choi, Christopher J. Shallue, Zachary Nado, Jaehoon Lee, Chris J. Maddison, and George E. Dahl. On empirical comparisons of optimizers for deep learning. In *arXiv preprint*, 2019. arXiv:1910.05446.
- William G. Cochran. The combination of estimates from different experiments. *Biometrics*, 10(1):101–129, 1954.
- Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. ImageNet: A large-scale hierarchical image database. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 248–255, 2009.
- Rebecca DerSimonian and Nan Laird. Meta-analysis in clinical trials. *Controlled Clinical Trials*, 7(3): 177–188, 1986.
- Tim Dettmers, Mike Lewis, Sam Shleifer, and Luke Zettlemoyer. 8-bit optimizers via block-wise quantization. In *International Conference on Learning Representations (ICLR)*, 2022. arXiv:2110.02861.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016.
- Julian P. T. Higgins and Simon G. Thompson. Quantifying heterogeneity in a meta-analysis. *Statistics in Medicine*, 21(11):1539–1558, 2002.
- Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6 (2):65–70, 1979.
- Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, pages 448–456, 2015.
- Khurram Javed, Arsalan Sharifnassab, and Richard S. Sutton. SwiftTD: A fast and robust algorithm for temporal difference learning. In *Reinforcement Learning Conference (RLC)*, 2024.
- Renlong Jie, Junbin Gao, Andrey Vasnev, and Minh-Ngoc Tran. Adaptive multi-level hyper-gradient descent. *arXiv preprint*, 2020. arXiv:2008.07277.

- Keller Jordan, Yuchen Jin, Vlado Boza, Jiacheng You, Franz Cesista, Laker Newhouse, and Jeremy Bernstein. Muon: An optimizer for hidden layers in neural networks. <https://kellerjordan.github.io/posts/muon/>, 2024.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*, 2015.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- Siyuan Li, Juanxi Tian, Zedong Wang, Xin Jin, Zicheng Liu, Wentao Zhang, and Dan Xu. Taming LLMs by scaling learning rates with gradient grouping. *arXiv preprint*, 2025. arXiv:2506.01049; the method is abbreviated SGG.
- Jingyuan Liu, Jianlin Su, Xingcheng Yao, Zhejun Jiang, Guokun Lai, et al. Muon is scalable for LLM training. *arXiv preprint*, 2025. arXiv:2502.16982; the Muon optimiser itself is due to Jordan et al. (2024).
- Ilya Loshchilov and Frank Hutter. SGDR: Stochastic gradient descent with warm restarts. In *International Conference on Learning Representations (ICLR)*, 2017.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019.
- Ashique Rupam Mahmood, Richard S. Sutton, Thomas Degris, and Patrick M. Pilarski. Tuning-free step-size adaptation. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 2121–2124, 2012.
- Arsalan Sharifnassab, Saber Salehkaleybar, and Richard S. Sutton. MetaOptimize: A framework for optimizing step sizes and other meta-parameters. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025. arXiv:2402.02342.
- Noam Shazeer and Mitchell Stern. Adafactor: Adaptive learning rates with sublinear memory cost. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2018. arXiv:1804.04235.
- Hao-Jun Michael Shi, Tsung-Hsien Lee, Shintaro Iwasaki, Jose Gallego-Posada, Zhijing Li, Kaushik Rangadurai, Dheevatsa Mudigere, and Michael Rabbat. A distributed data-parallel PyTorch implementation of the distributed Shampoo optimizer for training neural networks at-scale. *arXiv preprint*, 2023. arXiv:2309.06497.
- Richard S. Sutton. Adapting bias by gradient descent: An incremental version of delta-bar-delta. In *Proceedings of the Tenth National Conference on Artificial Intelligence (AAAI)*, pages 171–176, 1992.
- Tijmen Tieleman and Geoffrey Hinton. Lecture 6.5—RMSProp: Divide the gradient by a running average of its recent magnitude. COURSERA: Neural Networks for Machine Learning, 2012.
- B. L. Welch. The generalization of ‘student’s’ problem when several different population variances are involved. *Biometrika*, 34(1/2):28–35, 1947.
- Yuxin Wu and Kaiming He. Group normalization. In *European Conference on Computer Vision (ECCV)*, pages 3–19, 2018.
- Yang You, Igor Gitman, and Boris Ginsburg. Large batch training of convolutional networks. *arXiv preprint*, 2017. arXiv:1708.03888; the method is known as LARS.
- Yang You, Jing Li, Sashank Reddi, Jonathan Hseu, Sanjiv Kumar, Srinadh Bhojanapalli, Xiaodan Song, James Demmel, Kurt Keutzer, and Cho-Jui Hsieh. Large batch optimization for deep learning: Training BERT in 76 minutes. In *International Conference on Learning Representations (ICLR)*, 2020. arXiv:1904.00962; the method is known as LAMB.

- Yushun Zhang, Congliang Chen, Ziniu Li, Tian Ding, Chenwei Wu, Yinyu Ye, Zhi-Quan Luo, and Ruoyu Sun. Adam-mini: Use fewer learning rates to gain more. *arXiv preprint*, 2024. arXiv:2406.16793.
- Rosie Zhao, Depen Morwani, David Brandfonbrener, Nikhil Vyas, and Sham Kakade. Deconstructing what makes a good optimizer for autoregressive language models. *arXiv preprint*, 2024. arXiv:2407.07972; the layerwise variant is referred to as Adalayer.
- Shuai Zheng and James T. Kwok. Blockwise adaptivity: Faster training and better generalization in deep learning. In *arXiv preprint*, 2019. arXiv:1905.09899.